

Supplementary Materials for “Randomized Predictive P-values: A Versatile Model Diagnostic Tool with Unified Reference Distribution”

February 10, 2019

1 Proof of the Uniformity of RPPs

The RPPs are uniformly distributed, and correspondingly the NRPPs are normally distributed, under the true model. First, let us recall the well-known property of p-value for a continuous response variable written as

Theorem 1.1 *Suppose a continuous random variable Y has the CDF given by $F(y)$, then $F(Y)$ is uniformly distributed on $(0,1)$.*

Theorem 1.1 leads to the well-known fact that the p-values of a test statistic are uniformly distributed on $(0,1)$ when the null distribution is true. This uniformity is used to validate the well-calibration of computed p-values. Another equivalent way to express Theorem 1.1 is that: suppose $F(y)$ is a continuous CDF, let $F^{-1}(u)$ denote the inverse function defined as $F^{-1}(u) = \inf\{x | F(x) > u\}$, and U is uniformly distributed on $(0,1)$, then $F^{-1}(U)$ is distributed as $F(y)$. That is, we can transform uniform random numbers to random numbers with any desired continuous distribution, such as normal. When the y_i is discrete, Theorem 1.1 can be extended to:

Theorem 1.2 *Suppose the true distribution of Y_i given $X_i = x_i$ has the CDF $F_i(y_i)$ and PMF $p_i(y_i)$, where the subscript i indicates that F_i and p_i depends on a covariate x_i . The randomized predictive p-values $F^*(y_i, u_i)$ is defined as $F_i(y_i-) + u_i p_i(y_i)$. Suppose U_i is uniformly distributed on $(0,1)$. Then, we have*

$$F^*(Y_i, U_i) \sim \text{Uniform}(0,1), \text{ and } \phi^{-1}(F^*(Y_i, U_i)) \sim N(0,1). \quad (1)$$

Proof: First, we note that the normality of $\phi^{-1}(F^*(Y_i, U_i))$ can be derived from the uniformity of RPPs based on Theorem 1.1, and hence, it suffices to prove the uniformity of RPPs. Suppose all the possible values (with positive mass) for Y_i given $X_i = x_i$ are $k^{(1)}, k^{(2)}, \dots$. Let $P(Y_i = k^{(j)}) = p^{(j)}$, and $F^{(j)} = (F_i(k^{(j)}-), F_i(k^{(j)}))$; note that $\lambda(F^{(j)}) = p^{(j)}$ where $\lambda(B)$ denotes the ordinary length (ie, Lebesgue measure) of B . We note that $\cup_{j=1}^{\infty} F^{(j)} = (0,1) \setminus \{F_i(k^{(1)}), F_i(k^{(2)}), \dots\}$ and the collection of sets $\{F^{(j)} | j = 1, 2, \dots\}$

are mutually exclusive. Conditional on $Y_i = k^{(j)}$, $F^*(Y_i, U_i)$ is uniformly distributed on $F^{(j)}$ because U_i is uniformly distributed on $(0, 1)$. Therefore, for any interval (or Borel set) $B \subseteq (0, 1)$, $P(F^*(Y_i, U_i) \in B | Y_i = k^{(j)}) = \frac{\lambda(F^{(j)} \cap B)}{p^{(j)}}$. By the law of total probability, we have

$$\begin{aligned} P(F^*(Y_i, U_i) \in B) &= \sum_{j=1}^{\infty} P(F^*(Y_i, U_i) \in B | Y_i = k^{(j)}) \times P(Y_i = k^{(j)}) \\ &= \sum_{j=1}^{\infty} \frac{\lambda(F^{(j)} \cap B)}{p^{(j)}} \times p^{(j)} = \sum_{j=1}^{\infty} \lambda(F^{(j)} \cap B) \\ &= \lambda(\cup_{j=1}^{\infty} F^{(j)} \cap B) = \lambda(B) \end{aligned}$$

End of proof.

Next, we make a few remarks to clarify the applications of Theorem 1.2.

1. Since the conditional distribution of $F^*(Y_i, U_i)$ given $X_i = x_i$ is uniformly distributed on $(0, 1)$, and this distribution is free of x_i , the marginal distribution of $F^*(Y_i, U_i)$ with x_i marginalized away is still uniformly distributed on $(0, 1)$. This justifies that the overall distribution of RPPs is uniform on $(0, 1)$.
2. In frequentist paradigm, the $F_i(y_i)$ is the CDF of the true model with the true parameters that have generated the dataset. In practice, the parameters may be estimated with the sample data including y_i itself. The use of estimated parameters that have learned from y_i itself will introduce conservatism in the predictive p-values due to using y_i twice. As a result, the predictive p-values may be more concentrated around 0.5 than the uniform distribution on $(0, 1)$; correspondingly, the NRPPs tend to be more concentrated around 0 than distributed as $N(0, 1)$. This conservatism is minor when the sample size is much larger than the number of parameters. Our empirical studies (not shown in this paper) also indicated that the conservatism affects less in the following overall GOF test applied to NRPP if we use Shapiro-Wilk normality test compared to other tests, such as Kolmogorov-Smirnov test. For very complex models with a high risk of overfitting, it is necessary to eliminate this conservatism by computing cross-validatory NRPP. In this paper, we focus on discussing the necessity of using “randomization” to obtain truly uniform predictive p-values, hence, we ignore this conservatism by considering relatively simple models.
3. In Bayesian paradigm, the $F_i(y_i)$ is the CDF of the CV predictive distribution of Y_i^{rep} given y_{-i} (and covariates x_1, \dots, x_n) with model parameters θ marginalized away with respect to the posterior based on y_{-i} , as given below:

$$F_i(y_i) = P(Y_i^{\text{rep}} \leq y_i | y_{-i}) = \int P(Y_i^{\text{rep}} \leq y_i | \theta) P(\theta | y_{-i}) d\theta.$$

Therefore, the $F^*(y_i, u_i)$ is the cross-validatory randomized predictive p-values (CVRPP). Theorem 1.2 is an extension of the theorem proved by Marshall and Spiegelhalter

(2007) about the uniformity of CV predictive p-values for continuous response variable under the true model to the uniformity of CVRPP for discrete response variable. In Bayesian sense, the uniformity of CVRPP holds when both of the prior and the likelihood are correctly specified. Therefore, the non-uniformity of CVRPPs may result from mis-specification in either prior or likelihood, or both, and hence, reveals the discrepancies in both prior and likelihood. Last, we note that, although NRPPs without CV is the same as RQRs in simple regression, CVRPPs can be applied to diagnose hierarchical Bayesian models for data with complex correlation structure; see Marshall and Spiegelhalter (2007); Li et al. (2017). Theorem 1.2 provides theoretical foundation for model diagnosis with CVRPPs for models with discrete response. Further discussion of computing CVRPPs for Bayesian models is given in Section 6 of the main manuscript.

2 Power Analysis Results

2.1 Detection of Non-linearity

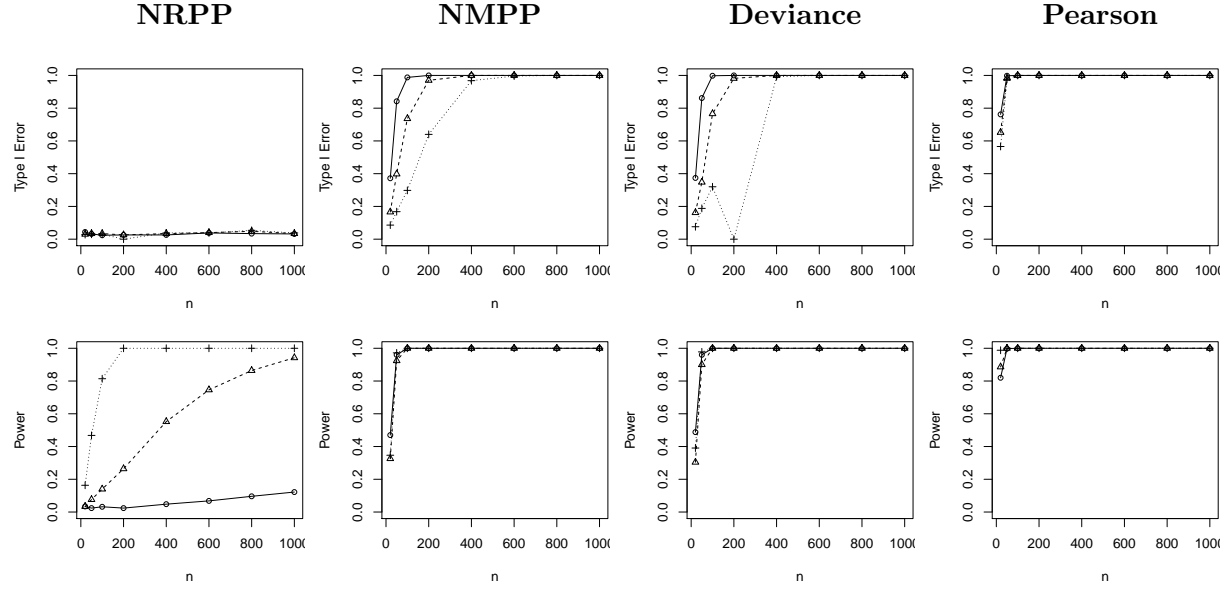


Figure S1: Comparison of the type I errors and powers of the SW tests for the NRPPs, NMPPs, and deviance and Pearson residuals. Response variable is simulated from the true model at varying sample sizes, and nonlinear covariate effects of $\beta_1 = 0.5$ (—○), 1 (---△) and 2 (.....+). True model: NB model with $\mu_i = \exp(\beta_1 x_i^2)$. Wrong model: NB model with $\mu_i = \exp(\beta_1 x_i)$.

2.2 Detection of Over-dispersion

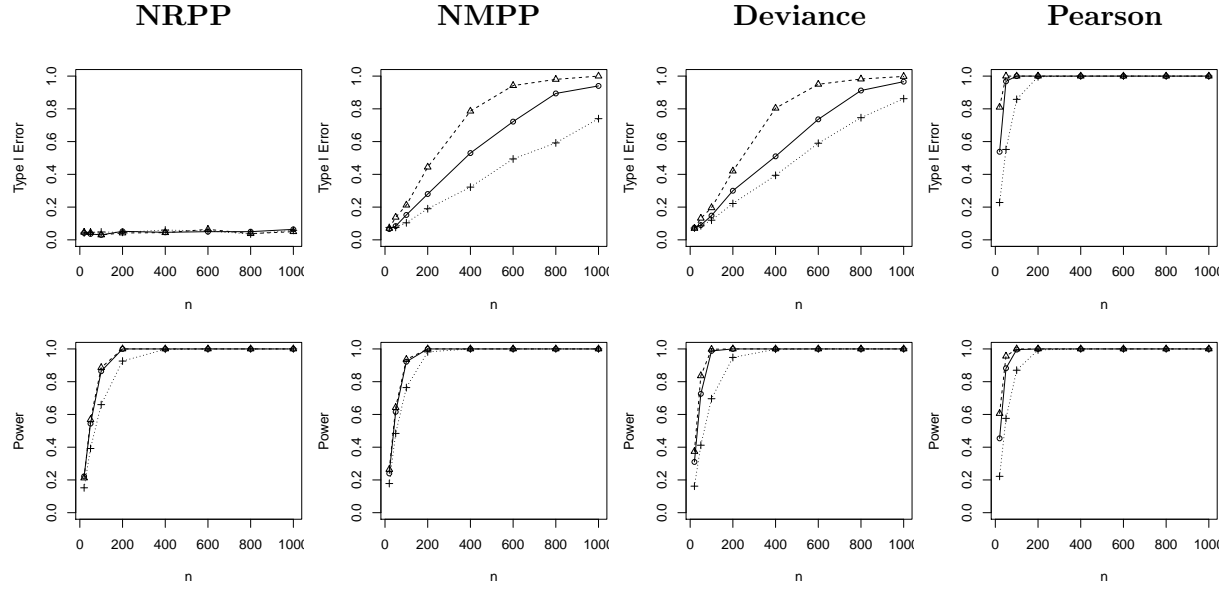


Figure S2: Comparison of the type I errors and powers of the SW tests for the NRPPs, NMPPs, and deviance and Pearson residuals. Response variable is simulated from the true model at varying sample sizes and the over-dispersion parameters of $k = 1$ (—○), 2 (---△) and 10 (.....+). True model: NB model with mean $\exp(\beta_0 + \beta_1 x)$. Wrong model: Poisson model with mean $\exp(\beta_0 + \beta_1 x)$.

2.3 Detection of Zero-Inflation

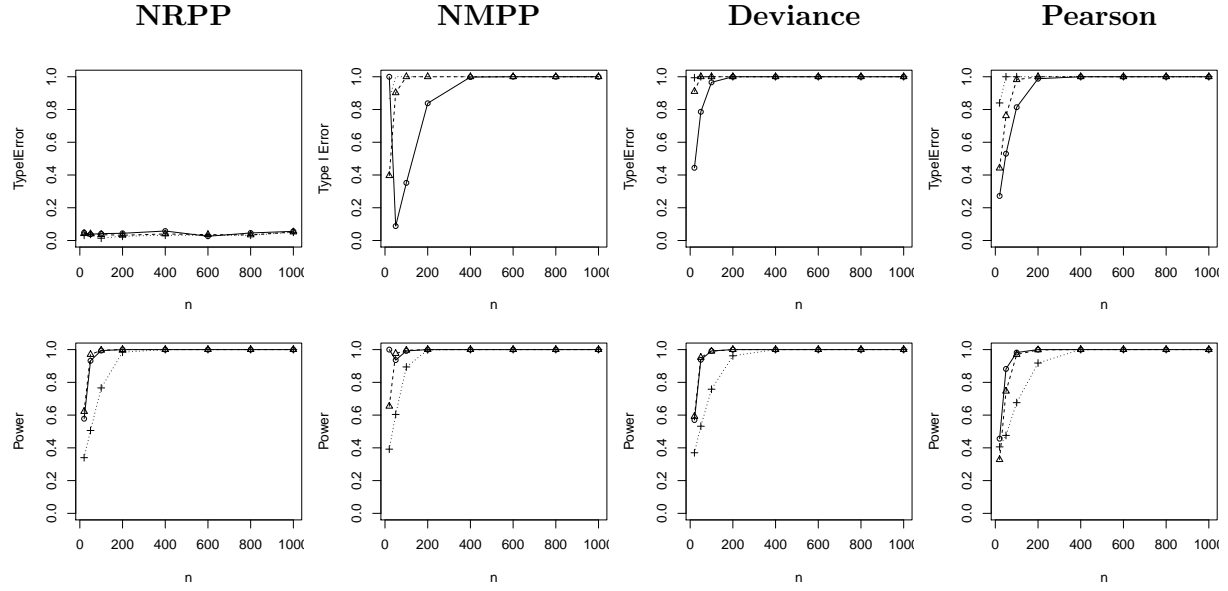


Figure S3: Comparison of the type I errors and powers of the SW tests for the NRPPs, NMPPs, and deviance and Pearson residuals. Response variable is simulated from the true model at varying sample sizes and percentages of excessive zeroes of $p = 10\%$ (—○), 30% (---△) and 50% (.....+). True model: ZIP model with mean $\exp(\beta_0 + \beta_1 x)$. Wrong model: Poisson model with mean $\exp(\beta_0 + \beta_1 x)$.

3 Results for Checking Logistic Regression

In this Section, we present the results for the performance of NRPPs and NMPPs for checking GOF of the logistic regression model.

We repeatedly simulate 500 datasets from the true model: a logistic regression model with $\text{logit}(p_i) = \beta_1 \sin(5x_i)$, with the covariate $x_i \sim \text{Uniform}(0, \pi)$ and $\beta_1 = 2$. We consider a wrong logistic regression model with linear effect $\text{logit}(p_i) = \beta_0 + \beta_1 x_i$.

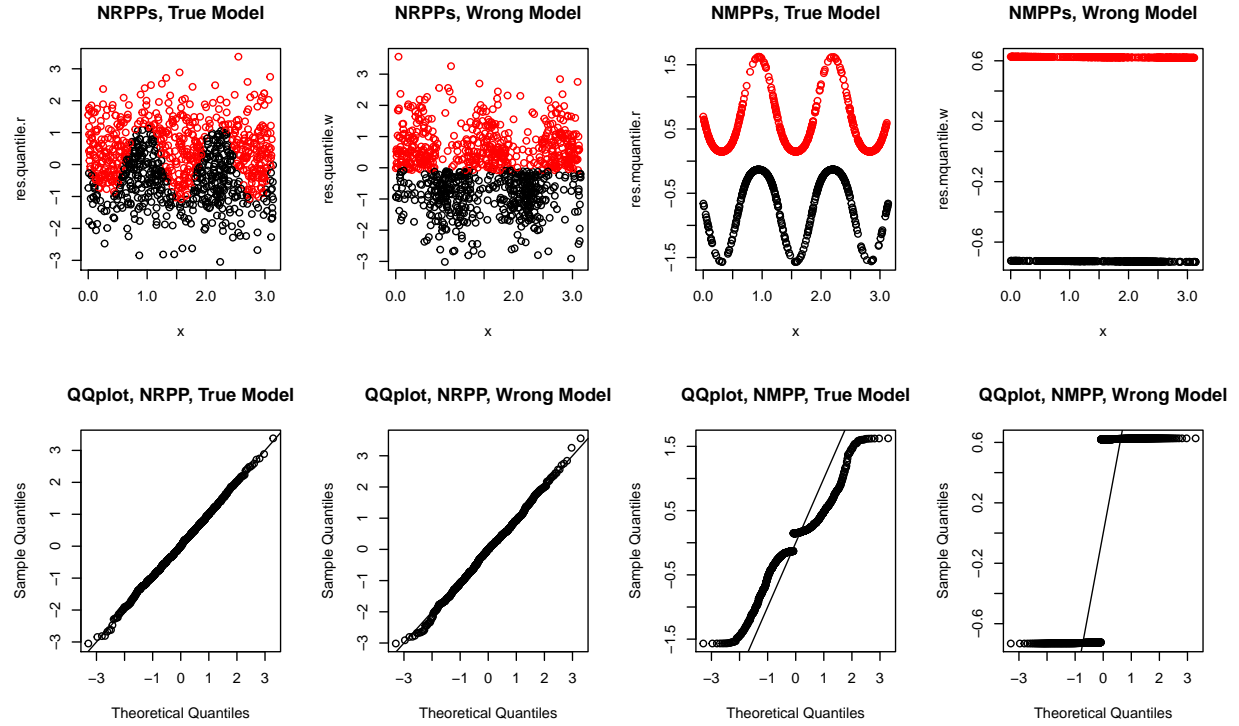


Figure S4: Residual plots of NRPPs and NMPPs versus x_i under the true vs. wrong models (top row), and the corresponding QQ plots (bottom row). The colours in the residual plot represent two values of y_i .

The residual plot of NRPPs for the wrong model clearly reveals the non-linearity. However, we see that the SW normality test applied to NRPPs has very low power in detecting the non-linearity. Another overall statistical test to measure the non-linearity pattern of NRPPs against the covariate is therefore desirable. We experiment a GOF test by testing the equal variances of NRPPs after sorted by x_i . We sort NRPPs by x_i and then divide them into 20 groups evenly from the smallest to the largest. Bartlett equal-variance test is applied to test whether the variances of the 20 groups of NRPPs are the same or not. Figure S5 shows the results of SW tests and Bartlett tests. It is clearly shown that the power of Bartlett test is very high, whereas the power of SW test is not good. Although the power of the Bartlett test is high, the results may depend on the choice of the number of groups. Hence, it remains an interesting topic to devise better GOF tests applied to NRPPs for checking model fit of logistic regression models.




Figure S5: 500 replicated GOF test p-values for NRPPs with SW normality tests and Bartlett's equal-variance tests.

4 R Code

All the R code for the simulation studies and the real data analysis, as well as the dataset used in Section 5 of the main article can be also downloaded through the corresponding author's personal website, see here [the url will be inserted upon publication].

4.1 R Codes for Illustrative Examples

4.1.1 An Example with No Covariate

```
library(ggplot2)
library(ggExtra)
library(gridExtra)
#####
# generate data
#####
#sample size
n <- 2000
smp <- sample(1:n, 100)## randomly selected sample for drawing scatterplot
#Generating data
y <- rbinom(n,2,.5)
#####
# RPP for the true model
#####
pv.r <- pbinom(y-1,2,.5)+runif(n)*dbinom(y,2,.5)
# select a small subset of pv.r for drawing scatterplot
pv.r2 <- rep (NA, n)
pv.r2[smp] <- pv.r[smp]
#####
# RPP for the wrong model
#####
# functions for calculating the pmf and the cdf of the wrong distribution(.1,.8,.1)
pmf <- function(x)
{
  return((x==0)*(.1)+(x==1)*(.8)+(x==2)*(.1))
}
cdf <- function(x)
{
  return((x==0)*(.1)+(x==1)*(.9)+(x==2)*(1))
}

#Residuals for wrong model
pv.w <- cdf(y-1)+runif(n)*pmf(y)
# select a small subset of pv.w for drawing scatterplot
```

```

pv.w2 <- rep (NA, n)
pv.w2[smp] <- pv.w[smp]

#####
# dataset for ggplot
#####
#Saving all the results in a data frame to be used in ggplot
df <- data.frame(y=y,pv.r=pv.r,pv.r=pv.r2,pv.w=pv.w,pv.w2=pv.w2)
#####
## plot settings
pch = 1
#a function for different colors of the CDF
clr <- function(x){return((x==0)*(1)+(x==1)*(2)+(x==2)*(4))}
#####

#####
# Drawing RPPs for the true model
#####
p.r <- ggplot(df, aes(as.factor(y), pv.r2)) +
  geom_point(pch=pch,colour=clr(y),size=3) +
  theme_classic() +
  ylab("RPP") +
  xlab("y") +
  theme(axis.text=element_text(size=rel(1.5))) +
  theme(axis.title = element_text(size = rel(1.5)))

#histogram for the true model
hist_right.r <- ggplot(data=df, aes(pv.r)) +
  geom_histogram(breaks=seq(0, 1, by=.1),
                 fill=I("blue"),col=I("red"),alpha=I(.2)) +
  coord_flip() +
  scale_y_reverse() +
  xlab("RPP") +
  theme(axis.text=element_text(size=rel(1.5))) +
  theme(axis.title = element_text(size = rel(1.5)))

#####
# Drawing RPPs for the wrong model
#####
p.w <- ggplot(df, aes(as.factor(y), pv.w2)) +
  geom_point(pch=pch,colour=clr(y), size=3) +
  theme_classic() +
  ylab("RPP") +
  xlab("y") +

```

```

    theme(axis.text=element_text(size=rel(1.5))) +
    theme(axis.title = element_text(size = rel(1.5)))
#histogram for the wrong model
hist_right.w <- ggplot(data=df, aes(pv.w)) +
  geom_histogram(breaks=seq(0, 1, by=.1),
                 fill=I("blue"),col=I("red"),alpha=I(.2)) +
  coord_flip() +
  scale_y_reverse() +
  xlab("RPP") +
  theme(axis.text=element_text(size=rel(1.5))) +
  theme(axis.title = element_text(size = rel(1.5)))
grid.arrange(hist_right.r,p.r, ncol=2, nrow=1)
grid.arrange(hist_right.w,p.w, ncol=2, nrow=1)

```

4.1.2 An Example with Covariate

```

# generate a dataset
n <- 1000
x <- seq (0, 2*pi, length = n)
mu <- exp (-1 + 2*sin(2*x))
y <- rpois(length (x), mu)
uy <- sort(unique(y)) ## unique values in y

# define col representation for different values of y
col <- rainbow(length (uy), s=1, start=0, end = 0.7)
# fit the true model
fit.r <- glm(y~sin(2*x), family = "poisson")
mu.r <- fit.r$fitted.values
# fit a wrong model
fit.w <- glm(y~x, family = "poisson")
mu.w <- fit.w$fitted.values
#draw plots
par(mfrow = c(4,2),mar=c(4.5, 4.5, 2, 1))
#=====
# RPP
#=====
## RPP for the true model:
RPP <- ppois(y-1, lambda = mu.r)+runif(n)*dpois(y, lambda = mu.r)
plot (x,rep (0.5, length (x)), ylim = c(0, 1), type ="n", ylab="RPP", xlab="x")
title("RPP, True Model")
k <- length (uy)
upper <- matrix (0,k, n)
for (i in 1:length (uy))
{

```

```

    upper[i,] <- ppois(uy[i], lambda = mu.r)
    lines(x, upper[i,], lty = 1, lwd=1,col = "black" )
}
points(x, RPP, col = col[match (y, uy)], pch = 20, cex=0.6)
###RPP for the wrong model
RPP <- ppois(y-1, lambda = mu.w)+runif(n)*dpois(y, lambda = mu.w)
plot (x, rep (0.5, length (x)), ylim = c(0, 1), type="n", ylab="RPP", xlab="x")
title("RPP, Wrong Model")
k <- length (uy)
upper <- matrix (0,k, n)
for (i in 1:length (uy))
{
    upper[i,] <- ppois(uy[i], lambda = mu.w)
    lines(x, upper[i,], lty = 1, col = "black" )
}
points (x, RPP, col = col[match (y, uy)], pch = 20, cex=0.6)
#=====
# MQR
#=====
### MQR for the True model:
MPP <- ppois(y-1, lambda = mu.r)+0.5*dpois(y, lambda = mu.r)
plot (x,rep (0.5, length (x)), ylim = c(0, 1), type = "n", ylab="MPP", xlab="x" )
title("MPP, True Model")
k <- length (uy)
upper <- matrix (0,k, n)
for (i in 1:length (uy))
{
    upper[i,] <- ppois(uy[i], lambda = mu.r)
    lines(x, upper[i,], lty = 1, col = "black" )
}
points(x, MPP, col = col[match (y, uy)], pch = 20, cex=0.6)

### MQR for the Wrong model:
MPP<- ppois(y-1, lambda = mu.w)+0.5*dpois(y, lambda = mu.w)
plot (x, rep (0.5, length (x)), ylim = c(0, 1), type="n", ylab="MPP", xlab="x")
title("MPP, Wrong Model")
k <- length (uy)
upper <- matrix (0,k, n)
for (i in 1:length (uy))
{
    upper[i,] <- ppois(uy[i], lambda = mu.w)
    lines(x, upper[i,], lty = 1, col = "black" )
}

```

```

points (x, MPP, col = col[match (y, uy)], pch = 20, cex=0.6)

#=====
# Deviance
#=====
res.deviance.r <- resid(fit.r,"deviance")
res.deviance.w <- resid(fit.w,"deviance")
plot (x, res.deviance.r, ylim = range (res.deviance.r),
      col = col[match (y, uy)], pch = 20, ylab="Deviance",
      main="Deviance, True Model", cex=0.6)
plot (x, res.deviance.w, ylim = range (res.deviance.w),
      col = col[match (y, uy)], pch = 20, ylab="Deviance",
      main="Deviance, Wrong Model", cex=0.6)
#=====
# Pearson
#=====
res.pearson.r <- resid(fit.r,"pearson")
res.pearson.w <- resid(fit.w,"pearson")
plot (x, res.pearson.r, ylim = range (res.pearson.r),
      col = col[match (y, uy)], pch = 20, ylab="Pearson",
      main="Pearson, True Model", cex=0.6)
plot (x, res.pearson.w, ylim = range (res.pearson.w),
      col = col[match (y, uy)], pch = 20, ylab="Pearson",
      main="Pearson, Wrong Model", cex=0.6)

```

4.2 R codes for Simulation Study

4.2.1 Simulation Setting #1: Detection of Non-Linear Covariate Effect

```

library("MASS")
n<-500 #sample size
x <- runif(n,-1.5,1.5)
lambda <- x^2
y <- rnbino(n,size = 2, mu = exp(lambda))
#####
#Model fitting
#####
fit.r <- glm.nb(y~lambda) # true model (nb: x^2)
fit.w <- glm.nb(y~x)      # wrong model(nb: x)
mu.hat.r <- fitted.values(fit.r)
mu.hat.w <- fitted.values(fit.w)
#####
#Residuals
#####

```

```

### deviance residual
res.deviance.r <- resid(fit.r,"deviance")
res.deviance.w <- resid(fit.w,"deviance")
### Pearson residual
res.pearson.r <- resid(fit.r,"pearson")
res.pearson.w <- resid(fit.w,"pearson")
### RQR (NRPP)
size.hat.r <- fit.r$theta
size.hat.w <- fit.w$theta
res.quantile.r <- qnorm(pnbinom (y-1,size = size.hat.r, mu = mu.hat.r) +
                        dnbinom (y,size = size.hat.r, mu = mu.hat.r) * runif(n))
res.quantile.w <- qnorm(pnbinom (y-1,size = size.hat.w, mu = mu.hat.w) +
                        dnbinom (y,size = size.hat.w, mu = mu.hat.w) * runif(n))
#####
### Residual plot and QQplot
#####
par (mfcol=c(2,2))
plot(x,res.quantile.r)
plot(x,res.quantile.w)
qqnorm(res.quantile.r);abline(a=0,b=1)
qqnorm(res.quantile.w);abline(a=0,b=1)
### Shapiro Wilk test
shapiro.test(res.quantile.r)
shapiro.test(res.quantile.w)

```

4.3 Simulation Setting #2: Detection of Overdispersion

```

library("MASS")
n<-100 #sample size
beta0 <- 1
beta1 <- 2
x <- runif(n,-1,2)
y <- rnbino(n,size = 2, mu = exp(beta0+beta1*x))
#####
#Model fitting
#####
fit.r <- glm.nb(y~x)
fit.w <- glm(y~x, family = "poisson")
mu.hat.r <- fitted.values(fit.r)
mu.hat.w <- fitted.values(fit.w)
size.hat.r <- fit.r$theta
#####
#Residuals
#####

```

```

### deviance residual
res.deviance.r <- resid(fit.r,"deviance")
res.deviance.w <- resid(fit.w,"deviance")
### Pearson residual
res.pearson.r <- resid(fit.r,"pearson")
res.pearson.w <- resid(fit.w,"pearson")
### RQR (NRPP)
res.quantile.r <- qnorm(pnbinom (y-1,size = size.hat.r, mu = mu.hat.r)+
dnbinom (y,size = size.hat.r, mu = mu.hat.r) * runif(n))
pvalue.w <- ppois (y-1,mu.hat.w) + dpois (y, mu.hat.w) * runif(n)
pvalue.w[pvalue.w==1] <- .999999999
res.quantile.w <- qnorm(pvalue.w)
#####
### Residual plot and QQplot
#####
par (mfcol=c(2,2))
plot(x,res.quantile.r)
plot(x,res.quantile.w)
qqnorm(res.quantile.r);abline(a=0,b=1)
qqnorm(res.quantile.w);abline(a=0,b=1)
### Shapiro Wilk test
shapiro.test(res.quantile.r)
shapiro.test(res.quantile.w)

```

4.4 Simulation Setting #3: Detecting Zero-Inflation

```

library("pscl")
dzpois <- function(x,lambda,p){return((1-p)*dpois(x,lambda)+p*(x==0))}
pzpois <- function(x,lambda,p){return((1-p)*ppois(x,lambda)+p*(x>=0))}
n<-100 #sample size
beta0 <- 1
beta1 <- 2
p <- .3
x <- runif(n,-1,2)
mu <- exp(beta0+beta1*x)
y <- rpois(n,mu)*rbinom(n,1,1-p)
#####
#Model fitting
#####
fit.r <- zeroinfl(y~x|1,dist="poisson")
fit.w1 <- glm(y~x,family="poisson")
mu.hat.r <- exp(coef(fit.r)[1]+coef(fit.r)[2]*x)
mu.hat.w1 <- fitted.values(fit.w1)
p.hat.r <- exp(coef(fit.r)[3])/(1+exp(coef(fit.r)[3]))

```

```
#####
#Residuals
#####
#deviance Residuals
res.deviance.r <- sign(y-mu.hat.r*(1-p.hat.r))*sqrt(2)*sqrt(log(dpois(y,y))
-log(dzpois(y,mu.hat.r,p.hat.r)))
res.deviance.w1 <- resid(fit.w1,"deviance")
#Pearson residuals
res.pearson.r <- resid(fit.r,"pearson")
res.pearson.w1 <- resid(fit.w1,"pearson")
### RQR (NRPP)
pvalue.r <- pzois (y-1,mu.hat.r,p.hat.r) + dzpois (y, mu.hat.r,p.hat.r) * runif(n)
pvalue.w1 <- ppois (y-1,mu.hat.w1) + dpois (y, mu.hat.w1) * runif(n)
pvalue.r[pvalue.r==1] <- .999999999
pvalue.w1[pvalue.w1==1] <- .999999999
res.quantile.r <- qnorm(pvalue.r)
res.quantile.w1 <- qnorm(pvalue.w1)
#####
### Residual plot and QQplot
#####
par (mfcol=c(2,2))
plot(x,res.quantile.r)
plot(x,res.quantile.w1)
qqnorm(res.quantile.r);abline(a=0,b=1)
qqnorm(res.quantile.w1);abline(a=0,b=1)
### Shapiro Wilk test
shapiro.test(res.quantile.r)
shapiro.test(res.quantile.w1)
```

4.5 Simulation Setting #4: Detecting of Non-Linear Covariate Effect in Logistic Regression

```
library("MASS")
#-----
# Test for equal variances of RQR against m groups of x
# n: sample size
# y: residuals
# x: covariate
# m: groups (split sample into m groups)
#-----
equavarance.test<-function(n, y, x, m){
  y.ordered<-y[order(x)] #order RQR by the covariate x
  test.eqvan.p<-bartlett.test(y.ordered~as.factor(rep(1:(n/m), each=m)))$p.value
  return(test.eqvan.p)
```



```

}
n<-1000 #sample size
x <- runif(n,0,pi)
xsq<- sin(5*x)
beta0 <- 0
beta1 <- 2
xbeta <- beta0+beta1*xsq
prob <- exp(xbeta)/(1+exp(xbeta))
y <- rbinom(n, 1,prob=prob)
#####
#Model fitting
#####
fit.r <- glm(y~xsq, family = "binomial")
fit.w <- glm(y~x, family ="binomial")
mu.hat.r <- fitted.values(fit.r)
mu.hat.w <- fitted.values(fit.w)
#####
#Residuals
#####
### deviance residual
res.deviance.r <- resid(fit.r,"deviance")
res.deviance.w <- resid(fit.w,"deviance")
### Pearson residual
res.pearson.r <- resid(fit.r,"pearson")
res.pearson.w <- resid(fit.w,"pearson")
### NMPP
mpvalue.r <- pbinom (y-1,1,prob=mu.hat.r)+dbinom (y,1,prob =mu.hat.r)*0.5
mpvalue.w <- pbinom (y-1,1,prob=mu.hat.w)+dbinom (y,1,prob=mu.hat.w)*0.5
res.mquantile.r<- qnorm (mpvalue.r)
res.mquantile.w<- qnorm (mpvalue.w)
### NRPP
pvalue.r <- pbinom (y-1,1,prob = mu.hat.r)+dbinom (y,1,prob=mu.hat.r)*runif(n)
pvalue.w <- pbinom (y-1,1,prob = mu.hat.w)+dbinom (y,1,prob=mu.hat.w)*runif(n)
res.quantile.r<- qnorm (pvalue.r)
res.quantile.w<- qnorm (pvalue.w)
#####
### residual analysis with RQR (NRPP)
#####
par (mfrow=c(2,4))
plot(x,res.quantile.r, main="NRPPs, True Model", col = y+1)
plot(x,res.quantile.w, main="NRPPs, Wrong Model", col = y+1)
plot(x,res.mquantile.r, main="NMPPs, True Model", col = y+1)
plot(x,res.mquantile.w, main="NMPPs, Wrong Model", col = y+1)
qqnorm(res.quantile.r, main="QQplot, NRPP, True Model");abline(a=0,b=1)

```

```

qqnorm(res.quantile.w, main="QQplot, NRPP, Wrong Model");abline(a=0,b=1)
qqnorm(res.mquantile.r, main="QQplot, NMPP, True Model");abline(a=0,b=1)
qqnorm(res.mquantile.w, main="QQplot, NMPP, Wrong Model");abline(a=0,b=1)
shapiro.test(res.quantile.r)
shapiro.test(res.quantile.w)
shapiro.test(res.mquantile.r)
shapiro.test(res.mquantile.w)
equavarance.test(n, res.quantile.r, x, m=20)
equavarance.test(n, res.quantile.w, x, m=20)
#####
## Replicate RQR for R times (power analysis)
#####
R<-500
pv.sw.r<-rep(NA, R)
pv.sw.w<-rep(NA, R)
pv.bartlett.r<-rep(NA, R)
pv.bartlett.w<-rep(NA, R)
for(r in 1:R){
  x <- runif(n,0,pi)
  xsq<- sin(5*x)
  beta0 <- 0
  beta1 <- 2
  xbeta <- beta0+beta1*xsq
  prob <- exp(xbeta)/(1+exp(xbeta))
  y <- rbinom(n, 1,prob=prob)
  #####
  #Model fitting
  #####
  fit.r <- glm(y~xsq, family = "binomial")
  fit.w <- glm(y~x, family ="binomial")
  mu.hat.r <- fitted.values(fit.r)
  mu.hat.w <- fitted.values(fit.w)
  pvalue.r <- pbinom (y-1,1,prob = mu.hat.r)+dbinom (y,1,prob=mu.hat.r)*runif(n)
  pvalue.w <- pbinom (y-1,1,prob = mu.hat.w)+dbinom (y,1,prob=mu.hat.w)*runif(n)
  res.quantile.r<- qnorm (pvalue.r)
  res.quantile.w<- qnorm (pvalue.w)
  pv.sw.r[r]<-shapiro.test(res.quantile.r)$p.value
  pv.sw.w[r]<-shapiro.test(res.quantile.w)$p.value
  pv.bartlett.r[r]<-equavarance.test(n, res.quantile.r, x, m=20)
  pv.bartlett.w[r]<-equavarance.test(n, res.quantile.w, x, m=20)
}
par(mfcol = c(2,2))
hist(pv.sw.r, main = "SW pvalues,NRPP,True Model")
hist(pv.sw.w, main = "SW pvalues,NRPP,Wrong Model")

```

```

hist(pv.bartlett.r,main = "Bartlett pvalues,NRPP,True Model")
hist(pv.bartlett.w,main = "Bartlett pvalues,NRPP,Wrong Model", nclass=20,xlim=c(0,1))
mean(pv.sw.w<=0.05)
mean(pv.sw.r<=0.05)
mean(pv.bartlett.w<=0.05)
mean(pv.bartlett.r<=0.05)

```

4.6 R code for real data analysis

```

#####
# Case study: number of hospital visits
#
#Demand for medical care by the elderly Deb and Trivedi (1997)
#analyze data on 4406 individuals, aged 66 and over, who are
#covered by Medicare, a public insurance program. Originally
#obtained from the US National Medical Expenditure Survey (NMES)
#for 1987/88, the data are available from the data archive of
#the Journal of Applied Econometrics at
#http://www.econ.queensu.ca/jae/1997-v12.3/deb-trivedi/.
#
#The objective is to model the demand for medical care-as captured
#by the number hospital stays-by the covariates available for the patients.
#
#Outcome: hosp (number of hospital stays)
#Covariates:
#age:      age in years (divided by 10)
#gender:   =1 if the person is male
#numchron: number of chronic conditions (cancer, heart attack,
#      gall bladder problems, emphysema, arthritis, diabetes,
#      other health disease)
#health:   self-perceived health status (excellent, average and poor),
#adldiff:  =1 if the person has a condition that limits activity of daily living
#####
library("pscl")
library("MASS")
#####
# CDF and PMF used in RPP
#####
#PDF of ZIP distribemervisitson;
dzpois <- function(x,xbeta,p)
{
return((1-p)*dpois(x,xbeta)+p*(x==0))
}
#CDF of ZIP distribemervisitson;

```

```

pzpois <- function(x,xbeta,p)
{
  return((1-p)*ppois(x,xbeta)+p*(x>=0))
}
#PDF of ZINB distribemervisitson;
dznbinom <- function(x,size,mu,p) #size: dispersion parameter
{
  return((1-p)*dnbinom(x,size = size, mu = mu)+p*(x==0))
}
#CDF of ZINB distribemervisitson;
pznbinom <- function(x,size,mu,p)
{
  return(
    (1-p)*pnbinom(x,size = size, mu = mu)+p*(x>=0)
  )
}
#####
# plotting function
#####
plotf <- function(a,b,c,y,z)
{
  plot(a, b, ylab = y, xlab = "x", main = z)
  lines(lowess(a, b),col="blue",lwd=3)
  plot(c,b, ylab = y, xlab = "Fitted Values", main = z)
  lines(lowess(c,b),col="blue",lwd=3)
}
qqplotf <- function(res,x)
{
  qqnorm (res,main = paste(x),cex.main=2, cex.lab=1.5)
  abline(a=0,b=1, lty=2)
}
#####
## data analysis
#####
### load and process data
load("DebTrivedi.rda")
dt <- DebTrivedi
attach(dt)
health_poor<-as.numeric(health=="poor")
health_average<-as.numeric(health=="average")
health_excellent<-as.numeric(health=="excellent")
gendermale<-as.numeric(gender=="male")
adldiffyes<-as.numeric(adldiff=="yes")
y<-emer

```

```

par (mfrow = c(1,1))
hist(y, xlab="Number of emergency room visits", main="", col="blue", breaks=30)

### fit Poission, NB, ZIP, ZINB models
fit.pois <-glm(y~ black+numchron+ health_excellent+health_average+
adldiffyes+school, data = dt, family = poisson)
summary(fit.pois)
fit.nb <- glm.nb(y~ numchron+health_excellent+health_average+
adldiffyes+school, data = dt)
summary(fit.nb)
fit.zip <- zeroinfl(y~ black+numchron+ health_excellent+health_average+
adldiffyes+school|1, data = dt, dist="poisson")
summary(fit.zip)
fit.zinb <- zeroinfl(y~ numchron+ health_excellent+health_average+
adldiffyes+school|1, data = dt, dist="negbin")
summary(fit.zinb)
### compute AICs
AIC(fit.pois)
AIC(fit.nb)
AIC(fit.zip)
AIC(fit.zinb)
### Parameter estimates and fitted values
mu.hat.pois<-fitted.values(fit.pois)
mu.hat.nb <- fitted.values(fit.nb)
mu.hat.zip<-fitted.values(fit.zip)
mu.hat.zinb<-fitted.values(fit.zinb)
p.zip <- coef(fit.zip)[8]
p.zinb <- coef(fit.zinb)[7]
p.hat.zip<-exp(p.zip)/(1+exp(p.zip))
p.hat.zinb<-exp(p.zinb)/(1+exp(p.zinb))
size.hat.nb<- fit.nb$theta
size.hat.zinb <- fit.zinb$theta
## define colorized representation of values in y
uy <- sort (unique(y))
col <- rainbow(length(uy), s=1, start=0, end = 0.7)#[rm.col]
col.y <- col[match(y, uy)]
#####
#   Pearson residuals
#####
res.pearson.pois <- resid(fit.pois,"pearson")
res.pearson.nb <- resid(fit.nb,"pearson")
res.pearson.zip <- resid(fit.zip,"pearson")
res.pearson.zinb <- resid(fit.zinb,"pearson")
#Residuals vs. fitted values

```

```

par (mfrow = c(2,2))
plot(mu.hat.pois,res.pearson.pois,xlab="Fitted values",
ylab="Pearson Residuals",main="Poisson", cex.main=2, cex.lab=1.5, log="x", col=col.y)
plot(mu.hat.nb,res.pearson.nb,xlab="Fitted values",
ylab="Pearson Residuals",main="NB", cex.main=2, cex.lab=1.5, log="x", col=col.y)
plot(mu.hat.zip,res.pearson.zip,xlab="Fitted values",
ylab="Pearson Residuals",main="ZIP", cex.main=2, cex.lab=1.5, log="x", col=col.y)
plot(mu.hat.zinb,res.pearson.zinb,xlab="Fitted values",
ylab="Pearson Residuals",main="ZINB", cex.main=2, cex.lab=1.5, log="x", col=col.y)
#QQ plot
par (mfrow = c(2,2))
qqplotf(res.pearson.pois,"Poisson")
qqplotf(res.pearson.nb,"NB")
qqplotf(res.pearson.zip,"ZIP")
qqplotf(res.pearson.zinb,"ZINB")
#Shapiro-wilk normality test p.value
pv_pearson.pois<-shapiro.test(res.pearson.pois)$p.value;pv_pearson.pois
pv_pearson.nb<-shapiro.test(res.pearson.nb)$p.value; pv_pearson.nb
pv_pearson.zip<-shapiro.test(res.pearson.zip)$p.value; pv_pearson.zip
pv_pearson.zinb<-shapiro.test(res.pearson.zinb)$p.value; pv_pearson.zinb

#####
#   Deviance residuals
#####
res.deviance.pois <- resid(fit.pois,"deviance")
res.deviance.nb <- resid(fit.nb,"deviance")
res.deviance.zip <- sign(y-mu.hat.zip*(1-p.hat.zip))*sqrt(2)*
sqrt(log(dpois(y,y))-log(dzpois(y,mu.hat.zip,p.hat.zip)))
res.deviance.zinb <- sign(y-mu.hat.zinb*(1-p.hat.zinb))*sqrt(2)*
sqrt(log(dnbinom(y,size = size.hat.zinb, mu = y))-
log(dznbinom(y,size.hat.zinb, mu.hat.zinb,p.hat.zinb)))
#Residuals vs. fitted values
par (mfrow = c(2,2))
plot(mu.hat.pois,res.deviance.pois,xlab="Fitted values",
ylab="Deviance Residuals",main="Poisson", cex.main=2, cex.lab=1.5, log="x", col=col.y)
plot(mu.hat.nb,res.deviance.nb,xlab="Fitted values",
ylab="Deviance Residuals",main="NB", cex.main=2, cex.lab=1.5, log="x", col=col.y)
plot(mu.hat.zip,res.deviance.zip,xlab="Fitted values",
ylab="Deviance Residuals",main="ZIP", cex.main=2, cex.lab=1.5, log="x", col=col.y)
plot(mu.hat.zinb,res.deviance.zinb,xlab="Fitted values",
ylab="Deviance Residuals",main="ZINB", cex.main=2, cex.lab=1.5, log="x", col=col.y)
#QQ plot
par (mfrow = c(2,2))
qqplotf(res.deviance.pois,"Poisson")

```

```

qqplotf(res.deviance.nb,"NB")
qqplotf(res.deviance.zip,"ZIP")
qqplotf(res.deviance.zinb,"ZINB")
#Shapiro-wilk normality test p.value
pv_deviance.pois<-shapiro.test(res.deviance.pois)$p.value; pv_deviance.pois
pv_deviance.nb<-shapiro.test(res.deviance.nb)$p.value; pv_deviance.nb
pv_deviance.zip<-shapiro.test(res.deviance.zip)$p.value; pv_deviance.zip
pv_deviance.zinb<-shapiro.test(res.deviance.zinb)$p.value; pv_deviance.zinb

#####
# Randomized Quantile Residuals (NRPP)
#####
set.seed(3)
#CDF
n<-length(y)
pvalue.pois <- ppois (y-1,mu.hat.pois) + dpois (y, mu.hat.pois)*runif(n)
pvalue.nb <- pnbinom (y-1,size = size.hat.nb, mu = mu.hat.nb) +
            dnbinom (y,size = size.hat.nb, mu = mu.hat.nb) * runif(n)
pvalue.zip <- pzipois (y-1,mu.hat.zip,p.hat.zip) +
            dzpois (y, mu.hat.zip, p.hat.zip) * runif(n)
pvalue.zinb <- pznbinom (y-1,size.hat.zinb, mu.hat.zinb, p.hat.zinb) +
            dznbinom (y,size.hat.zinb,mu.hat.zinb,p.hat.zinb) * runif(n)
pvalue.nb_linear <- pnbinom (y-1,size = size.hat.nb, mu = mu.hat.nb) +
            dnbinom (y,size = size.hat.nb, mu = mu.hat.nb) * runif(n)
pvalue.pois[pvalue.pois==1]<-0.999999999
pvalue.nb[pvalue.nb==1]<-0.999999999
pvalue.zip[pvalue.zip==1]<-0.999999999
pvalue.zinb[pvalue.zinb==1]<-0.999999999
#randomized quantile residuals (inverse CDF)
res.new.pois <- qnorm (pvalue.pois)
res.new.nb <- qnorm (pvalue.nb)
res.new.zip <- qnorm (pvalue.zip)
res.new.zinb <- qnorm (pvalue.zinb)
res.new.nb_linear <- qnorm (pvalue.nb_linear)
#Residuals vs. fitted values
par (mfrow = c(2,2))
plot(mu.hat.pois,res.new.pois,xlab="Fitted values", ylab="NRPP",
main="Poisson", cex.main=2, cex.lab=1.5,log="x", ylim=c(-6, 6))
abline(h=c(-3, 3), lty=2, col="blue")
plot(mu.hat.nb,res.new.nb,xlab="Fitted values",ylab="NRPP",
main="NB", cex.main=2, cex.lab=1.5,log="x", ylim=c(-6, 6))
abline(h=c(-3, 3), lty=2, col="blue")
plot(mu.hat.zip,res.new.zip,xlab="Fitted values",ylab="NRPP",
main="ZIP", cex.main=2, cex.lab=1.5,log="x", ylim=c(-6, 6))

```

```

abline(h=c(-3, 3), lty=2, col="blue")
plot(mu.hat.zinb,res.new.zinb,xlab="Fitted values",ylab="NRPP",
main="ZINB", cex.main=2, cex.lab=1.5,log="x", ylim=c(-6, 6))
abline(h=c(-3, 3), lty=2, col="blue")

#QQ plot
par (mfrow = c(2,2))
qqplotf(res.new.pois,"Poisson")
qqplotf(res.new.nb,"NB")
qqplotf(res.new.zip,"ZIP")
qqplotf(res.new.zinb,"ZINB")
#Shapiro-wilk normality test p.value
shapiro.test(res.new.pois)
shapiro.test(res.new.nb)
shapiro.test(res.new.zip)
shapiro.test(res.new.zinb)
#####
### replicate RQR (NRPP) for R times to examine the impact of randomization on RQR
#####
set.seed(1)
R <- 1000
pv.pois<-rep(NA, R)
pv.nb<-rep(NA, R)
pv.zip<-rep(NA, R)
pv.zinb<-rep(NA, R)
for (i in 1:R){
  #RPP
  pvalue.pois <- ppois (y-1,mu.hat.pois) + dpois (y, mu.hat.pois)*runif(n)
  pvalue.nb <- pnbinom (y-1,size = size.hat.nb, mu = mu.hat.nb) +
  dnbinom (y,size = size.hat.nb, mu = mu.hat.nb) * runif(n)
  pvalue.zip<-ppzpois(y-1,mu.hat.zip,p.hat.zip)+dzpois (y,mu.hat.zip,p.hat.zip)*runif(n)
  pvalue.zinb<-pznbino(y-1,size.hat.zinb,mu.hat.zinb,p.hat.zinb) +
  dznbinom (y,size.hat.zinb,mu.hat.zinb,p.hat.zinb)*runif(n)

  pvalue.pois[pvalue.pois==1]<-0.999999999
  pvalue.nb[pvalue.nb==1]<- 0.999999999
  pvalue.zip [pvalue.zip==1]<-0.999999999
  pvalue.zinb [pvalue.zinb==1]<-0.999999999

  #randomized quantile residuals (NRPP)
  res.new.pois <- qnorm (pvalue.pois)
  res.new.nb <- qnorm (pvalue.nb)
  res.new.zip <- qnorm (pvalue.zip)
  res.new.zinb <- qnorm (pvalue.zinb)

```



```

pv.pois[i]<-shapiro.test(res.new.pois)$p.value
pv.nb[i]<-shapiro.test(res.new.nb)$p.value
pv.zip[i]<-shapiro.test(res.new.zip)$p.value
pv.zinb[i]<-shapiro.test(res.new.zinb)$p.value
}
par (mfrow = c(2,2))
hist(pv.pois, main="Poisson", xlab="SW p-value", xlim=c(0, 1))
hist(pv.nb, main="NB", xlab="SW p-value", xlim=c(0, 1))
hist(pv.zip, main="ZIP", xlab="SW p-value", xlim=c(0, 1))
hist(pv.zinb, main="ZINB", xlab="SW p-value", xlim=c(0, 1))

```

References

- Li, L., Feng, C. X. and Qiu, S. (2017), ‘Estimating cross-validators predictive p-values with integrated importance sampling for disease mapping models’, *Statistics in Medicine* **36**(14), 2220–2236.
- Marshall, E. C. and Spiegelhalter, D. J. (2007), ‘Identifying outliers in Bayesian hierarchical models: a simulation-based approach’, *Bayesian Analysis* **2**(2), 409–444.