

STAT 812: Computational Statistics

Importance Sampling

Longhai Li

September 2018

Contents

1	Estimate Normal Probability $P(a < X < b)$ Using Importance Sampling	1
1.1	Test 1	2
1.2	Test 2	3
1.3	Test 3	3
1.4	Test 4: An example of computing a underflow probability	4
2	Estimate $E(a(X))$ of Truncated Normal Using Importance Sampling	5
2.1	Estimating Function	5
2.2	Test 1	7
2.3	Test 2	8
3	Computing Log Marginalized Likelihood for Normal Models with Importance Sampling	10
3.1	Functions	10
3.2	Testing	12

1 Estimate Normal Probability $P(a < X < b)$ Using Importance Sampling

```
## a function computing the sum of numbers represented with logarithm
## lx      --- a vector of numbers, which are the log of another vector x.
## the log of sum of x is returned
log_sum_exp <- function(lx)
{
  mlx <- max(lx)
  mlx + log(sum(exp(lx-mlx)))
}

## estimating the probability  $P(X \text{ in } A)$  for  $X \sim N(0,1)$ 
## by sampling from  $N(0,1)$ 
est_normprob_mc <- function(A, iters_mc)
{
  X <- rnorm(iters_mc)
  mean((X > A[1]) * (X < A[2]))
}

## estimating the probability  $P(X \text{ in } A)$  for  $X \sim N(0,1)$ 
## by sampling from  $Unif(A[1], A[2])$ 
est_normprobimps <- function(A, iters_mc)
```

```

{
  X <- runif(iters_mc,A[1],A[2])
  mean(dnorm(X))*(A[2]-A[1])
}

## estimating the probability P(X in A) for X ~ N(0,1)
## by sampling from Unif(A[1],A[2])
est_log_normprobimps <- function(A, iters_mc)
{
  X <- runif(iters_mc,A[1],A[2])
  log_sum_exp(dnorm(X,log = TRUE))-log(iters_mc)+ log(A[2]-A[1])
}

```



1.1 Test 1

```

A <- c(-2,2)
tp <- pnorm (A[2]) - pnorm (A[1])

probs_mc <- replicate(1000,est_normprob_mc(A,100))
probsimps <- replicate(1000,est_normprobimps(A,100))
par(mfrow = c(1,2))
ylim <- range (probs_mc, probsimps)
plot(probs_mc, ylim = ylim); abline (h=tp)
plot(probsimps, ylim = ylim); abline (h=tp)

```



```
mean((probs_mc-tp)^2)
```

```
## [1] 0.00043699
```


```
mean((probsimps-tp)^2)
```

```
## [1] 0.002079687
```

1.2 Test 2

```
A <- c(5,6)
tp <- pnorm (A[2]) - pnorm (A[1])

probs_mc <- replicate(1000,est_normprob_mc(A,100))
probsimps <- replicate(1000,est_normprobimps(A,100))
par(mfrow = c(1,2))
ylim <- range (probs_mc, probsimps)
plot(probs_mc, ylim = ylim); abline (h=tp)
plot(probsimps, ylim = ylim); abline (h=tp)
```



```
mean((probs_mc-tp)^2)
```

```
## [1] 8.160448e-14
```



```
mean((probsimps-tp)^2)
```

```
## [1] 1.378715e-15
```

1.3 Test 3

```
A <- c(-2,2)
tp <- pnorm (A[2]) - pnorm (A[1])

probs_mc <- replicate(1000,est_normprob_mc(A,100))
probsimps <- replicate(1000,est_normprobimps(A,100))
par(mfrow = c(1,2))
ylim <- range (probs_mc, probsimps)
plot(probs_mc, ylim = ylim); abline (h=tp)
plot(probsimps, ylim = ylim); abline (h=tp)
```



```
mean((probs_mc-tp)^2)
```

```
## [1] 0.0004562895
```

```
mean((probsimps-tp)^2)
```

```
## [1] 0.002244745
```

1.4 Test 4: An example of computing a underflow probability

```
A <- c(50,100)
```

```
#A method that cannot compute so small probability
```

```
tp1 <- pnorm(A[2]) - pnorm(A[1]); log(tp1)
```

```
## [1] -Inf
```

```
#The reason is that pnorm(A[1]) is so close to 1, or pnorm(-A[1]) underflow
```

```
pnorm(A[1])
```

```
## [1] 1
```

```
log(pnorm(A[1]))
```

```
## [1] 0
```

```
#Instead, we can use this way to compute the log probability
```

```
log_minus_exp <- function(la, lb) la + log(1-exp(lb-la))
```

```
la <- pnorm(-A[1], log=TRUE); la
```

```
## [1] -1254.831
```

```
lb <- pnorm(-A[2], log=TRUE); lb
```

```
## [1] -5005.524
```


```
log_tp2 <- log_minus_exp(la, lb); log_tp2
```

```
## [1] -1254.831
```

```
#We still use the same importance sampling procedure, but taking care of the underflow problem in dnorm
log_probsimps <- replicate(100,est_log_normprobimps(A,1000)); log_probsimps
```

```
## [1] -1257.701 -1254.564 -1258.275 -1254.933 -1258.755 -1253.772 -1254.338
## [8] -1255.530 -1254.048 -1257.340 -1259.333 -1256.555 -1253.495 -1255.619
## [15] -1253.464 -1259.737 -1254.494 -1253.780 -1258.678 -1262.472 -1254.979
## [22] -1253.998 -1255.943 -1254.194 -1254.543 -1254.743 -1263.006 -1259.066
## [29] -1255.268 -1259.054 -1256.712 -1256.951 -1262.242 -1255.913 -1254.591
## [36] -1253.613 -1258.030 -1257.995 -1254.963 -1253.434 -1254.063 -1258.861
## [43] -1254.563 -1267.488 -1257.681 -1253.927 -1253.821 -1254.851 -1254.956
## [50] -1254.310 -1255.387 -1255.949 -1263.110 -1254.487 -1253.941 -1254.775
## [57] -1258.338 -1256.559 -1257.098 -1253.420 -1256.386 -1256.942 -1257.405
## [64] -1255.353 -1254.906 -1256.141 -1260.230 -1257.334 -1260.913 -1256.597
## [71] -1256.585 -1254.066 -1255.903 -1260.663 -1269.147 -1254.911 -1257.999
## [78] -1256.473 -1254.026 -1253.922 -1257.722 -1257.249 -1254.119 -1254.058
## [85] -1262.687 -1253.976 -1256.211 -1254.977 -1253.866 -1255.375 -1260.788
## [92] -1256.886 -1256.629 -1258.077 -1261.576 -1254.250 -1254.296 -1254.967
## [99] -1258.773 -1259.492
```

```
plot(log_probsimps); abline(h = log_tp2)
```



2 Estimate $E(a(X))$ of Truncated Normal Using Importance Sampling

2.1 Estimating Function

```
## compute E(a) with importance sampling
est_tnormimps <- function(a, A, iters_mc)
{
  X <- runif(iters_mc,A[1],A[2])
  W <- dnorm(X)
  ahat <- sum(a(X) * W) / sum(W)
```

```

    attr(ahat, "effective sample size") <- 1/sum((W/sum(W))^2)
    ahat
  }

## a generic function for approximating 1-D integral with midpoint rule
## the logarithms of the function values are passed in
## the log of the integral result is returned
## log_f --- a function computing the logarithm of the integrand function
## range --- the range of integral variable, a vector of two elements
## n --- the number of points at which the integrand is evaluated
## ... --- other parameters needed by log_f
log_int_mid <- function(log_f, range, n,...)
{
  if(range[1] >= range[2])
    stop("Wrong ranges")
  h <- (range[2]-range[1]) / n
  v_log_f <- sapply(range[1] + (1:n - 0.5) * h, log_f,...)
  log_sum_exp(v_log_f) + log(h)
}

## compute E(a) with midpoint rule
est_tnorm_mid <- function(a, A, iters_mc)
{
  log_f <- function(x) dnorm(x, log = T) + log(a(x))
  exp(log_int_mid(log_f, A, iters_mc)) / (pnorm(A[2]) - pnorm(A[1]))
}

library(ars)

## a direct rejection sampling for truncated normal
sample_tnorm_drs <- function(n, lb = -Inf, ub = Inf)
{
  x <- rep(0, n)
  for(i in 1:n)
  {
    rej <- TRUE
    while(rej)
    {
      x[i] <- rnorm(1)
      if(x[i] >= lb & x[i] <= ub) rej <- FALSE
    }
  }
  x
}

## sample from truncated normal using ars package
sample_tnorm_ars <- function(n, lb, ub)
{
  logf <- function(x) dnorm(x, log = TRUE) ## define log density
  fprima <- function(x) -x ## define derivative of log density

  ars(10000, f = logf, fprima = fprima,
      x = c(lb, (lb + ub)/2, ub), # starting points

```

```

    lb = TRUE, ub = TRUE, xlb = lb, xub = ub) # boundary of log density
}

```

2.2 Test 1

```

## define the function a
a <- function (x) x^2
interval <- c(1,2)
A <- est_tnorm_mid (a, interval, 100000) ## midpoint rule

## estimate E(a) with rejection sampling
system.time(
  {
    rn_tnorm_ars <- sample_tnorm_ars (1000, interval[1],interval[2]) # draw samples from tnorm
    mean (a (rn_tnorm_ars))
  }
)

## user system elapsed
## 0.091 0.022 0.115

system.time(
  est_tnormimps (a, interval, 1000000) ## importance sampling
)

## user system elapsed
## 0.062 0.008 0.071

## simulation comparison of importance sampling and rejection sampling
timesimps <- system.time(
  EAimps <- replicate (100, est_tnormimps (a, interval, 1000000))
)

timesimps

## user system elapsed
## 5.915 0.179 6.199



times.rej <- system.time (
  EA_rej <- replicate (100,
    {
      rn_tnorm_ars <- sample_tnorm_ars (1000, interval[1],interval[2])
      mean (a (rn_tnorm_ars))
    }
  )
)

times.rej

## user system elapsed
## 7.958 0.697 8.712

par (mfrow = c(1,2))
ylim <- range (EAimps, EA_rej)
plot (EAimps, ylim = ylim, main = "Importance Sampling")
abline (h = A)
plot (EA_rej, ylim = ylim, main = "Adaptive Rejection Sampling")
abline (h = A)

```



```
mean ((EA_rej-A)^2)
```

```
## [1] 6.261365e-05
```

```
mean ((EAimps-A)^2)
```

```
## [1] 6.894636e-07
```

2.3 Test 2

```
## define the function a
a <- function (x) x^2
interval <- c(-1,1)
A <- est_tnorm_mid (a, interval, 100000) ## midpoint rule

## estimate E(a) with rejection sampling
system.time(
{
  rn_tnorm_ars <- sample_tnorm_ars (1000, interval[1],interval[2]) # draw samples from tnorm
  mean (a (rn_tnorm_ars))
}
)

## user system elapsed
## 0.076 0.006 0.083

system.time(
est_tnormimps (a, interval, 1000000) ## importance sampling
)

## user system elapsed
## 0.059 0.007 0.067
```



```
## simulation comparison of importance sampling and rejection sampling
timesimps <- system.time(
EAimps <- replicate (100, est_tnormimps (a, interval, 1000000))
)



timesimps

##      user  system elapsed
##  5.849   0.175   6.146

timesrej <- system.time (
EArej <- replicate (100,
  {
    rn_tnorm_ars <- sample_tnorm_ars (1000, interval[1],interval[2])
    mean (a (rn_tnorm_ars))
  }
)
)
timesrej

##      user  system elapsed
##  8.155   0.619   8.845

par (mfrow = c(1,2))
ylim <- range (EAimps, EArej)
plot (EAimps, ylim = ylim, main = "Importance Sampling")
abline (h = A)
plot (EArej, ylim = ylim, main = "Adaptive Rejection Sampling")
abline (h = A)
```



```
mean ((EArej-A)^2)
```

```
## [1] 6.258295e-06
```

```
mean ((EAimps-A)^2)
```

```
## [1] 5.611324e-08
```

3 Computing Log Marginalized Likelihood for Normal Models with Importance Sampling

3.1 Functions

```
## a function computing the sum of numbers represented with logarithm
## lx      --- a vector of numbers, which are the log of another vector x.
## the log of sum of x is returned
log_sum_exp <- function(lx)
{  mlx <- max(lx)
  mlx + log(sum(exp(lx-mlx)))
}

## computing the log probability density function of multivariate normal
## x      --- a vector, the p.d.f at x will be computed
## mu     --- the mean vector of multivariate normal distribution
## A      --- the inverse covariance matrix of multivariate normal distribution
log_pdf_mnormal <- function(x, mu, A)
{  0.5 * ( -length(mu)*log(2*pi) + sum(log(svd(A)$d)) - t(x-mu) %*% A %*% (x-mu) )
}

## the function for computing log likelihood of normal data
log_lik <- function(x,mu,w)
{  sum(dnorm(x,mu,exp(w),log=TRUE))
}

## the function for computing log prior
log_prior <- function(mu,w, mu_0,sigma_mu,w_0,sigma_w)
{  dnorm(mu,mu_0,sigma_mu,log=TRUE) + dnorm(w,w_0,sigma_w,log=TRUE)
}

## the function for computing the negative log of likelihood * prior
neg_log_post <- function(x, theta, mu_0,sigma_mu,w_0,sigma_w)
{  - log_lik(x,theta[1], theta[2]) -
  log_prior(theta[1],theta[2],mu_0,sigma_mu,w_0,sigma_w)
}

## computing the log marginal likelihood using importance sampling with
## the posterior distribution approximated by the Gaussian distribution at
## its mode
log_mar_gaussianimps <- function(x,mu_0,sigma_mu,w_0,sigma_w,itors_mc)
{  result_min <- nlm(f=neg_log_post,p=c(mean(x),log(sqrt(var(x)))),
  hessian=TRUE,
  x=x,mu_0=mu_0,sigma_mu=sigma_mu,w_0=w_0,sigma_w=sigma_w)
  hessian <- result_min$hessian
  mu <- result_min$estimate

  ## finding the multiplier for sampling from multivariate normal
```

```

Sigma <- t( chol(solve(hessian)) )
## draw samples from N(mu, Sigma %%% Sigma')
thetas <- Sigma %%% matrix(rnorm(2*iters_mc),2,iters_mc) + mu

## values of log approximate p.d.f. at samples
log_pdf_mnormal_thetas <- apply(thetas,2,log_pdf_mnormal,mu=mu,A=hessian)
## values of log true p.d.f. at samples
log_post_thetas <- - apply(thetas,2,neg_log_post,x=x, mu_0=mu_0,
                           sigma_mu=sigma_mu,w_0=w_0,sigma_w=sigma_w)

## averaging the weights, returning its log
log_sum_exp(log_post_thetas-log_pdf_mnormal_thetas) - log(iters_mc)
}

## we use Monte Carlo method to debug the above function
log_mar_gaussian_mc <- function(x,mu_0,sigma_mu,w_0,sigma_w,iters_mc)
{
  ## draw samples from the priors
  mus <- rnorm(iters_mc,mu_0,sigma_mu)
  ws <- rnorm(iters_mc,w_0,sigma_w)
  one_log_lik <- function(i)
  { log_lik(x,mus[i],ws[i])
  }
  v_log_lik <- sapply(1:iters_mc,one_log_lik)
  log_sum_exp(v_log_lik) - log(iters_mc)
}

## the generic function for finding laplace approximation of integral of 'f'
## neg_log_f      --- the negative log of the intergrand function
## p0            --- initial value in searching mode
## ...           --- other arguments needed by neg_log_f
bayes_inference_lap <- function(neg_log_f,p0,...)
{
  ## looking for the mode and hessian of the log likelihood function
  result_min <- nlm(f=neg_log_f,p=p0, hessian=TRUE,...)
  hessian <- result_min$hessian
  neg_log_like_mode <- result_min$minimum

  estimates <- result_min$estimate ## posterior mode
  SIGMA <- solve(result_min$hessian) ## covariance matrix of posterior mode
  sds <- sqrt(diag(SIGMA)) ## standard errors of each estimate
  log_mar_lik <- ## log marginalized likelihood
    - neg_log_like_mode + 0.5 * ( sum(log(2*pi) - log(svd(hessian)$d) ))

  list(estimates = estimates, sds = sds, SIGMA = SIGMA, log_mar_lik = log_mar_lik)
}

## approximating the log of integral of likelihood * prior
bayes_inference_lap_gaussian <- function(x,mu_0,sigma_mu,w_0,sigma_w)
{
  bayes_inference_lap(
    neg_log_post,p0=c(mean(x),log(sqrt(var(x)))),
    x=x,mu_0=mu_0,sigma_mu=sigma_mu,w_0=w_0,sigma_w=sigma_w
  )
}

```

```
)  
}
```

3.2 Testing

```
## debugging the program  
x <- rnorm(50)  
log_mar_gaussianimps(x,0,1,0,5,100)  
  
## [1] -73.35034  
log_mar_gaussian_mc(x,0,1,0,5,10000)  
  
## [1] -73.5719  
bayes_inference_lap_gaussian(x,0,1,0,5)  
  
## $estimates  
## [1] 0.02905778 -0.06963118  
##  
## $sds  
## [1] 0.13077623 0.09999276  
##  
## $SIGMA  
##           [,1]      [,2]  
## [1,] 1.710242e-02 -8.963959e-06  
## [2,] -8.963959e-06 9.998553e-03  
##  
## $log_mar_lik  
## [1] -73.41088  
  
x <- rnorm(10) # another debug  
log_mar_gaussianimps(x,0,1,0,5,100)  
  
## [1] -18.96407  
log_mar_gaussian_mc(x,0,1,0,5,10000)  
  
## [1] -19.01825  
bayes_inference_lap_gaussian(x,0,1,0,5)  
  
## $estimates  
## [1] -0.07190741 0.05481965  
##  
## $sds  
## [1] 0.3168537 0.2233930  
##  
## $SIGMA  
##           [,1]      [,2]  
## [1,] 0.1003962818 0.0007248575  
## [2,] 0.0007248575 0.0499044160  
##  
## $log_mar_lik  
## [1] -18.99896
```

```

## comparing importance sampling with Gaussian approximation with naive monte carlo
x <- rnorm(200)
bayes_inference_lap_gaussian(x,0,1,0,5)

## $estimates
## [1] 0.007028495 0.118606615
##
## $sds
## [1] 0.07936393 0.05000188
##
## $SIGMA
##           [,1]      [,2]
## [1,] 6.298633e-03 2.461681e-08
## [2,] 2.461681e-08 2.500188e-03
##
## $log_mar_lik
## [1] -314.6507

v_log_marimps <- replicate(1000, log_mar_gaussianimps(x,0,1,0,5,100))
v_log_mar_mc <- replicate(1000, log_mar_gaussian_mc(x,0,1,0,5,100))

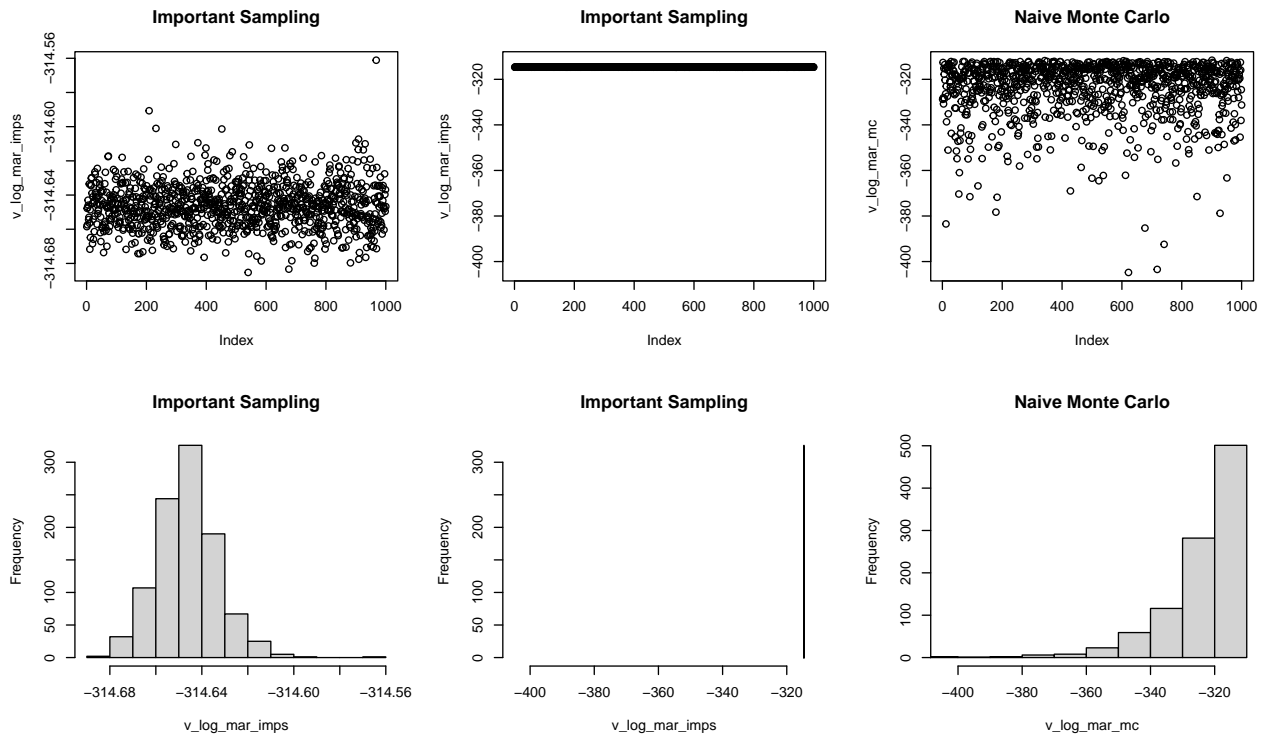
par(mfcol=c(2,3))
xlim <- c(min(c(v_log_marimps,v_log_mar_mc)),max(c(v_log_marimps,v_log_mar_mc)))

plot(v_log_marimps, main="Important Sampling")
hist(v_log_marimps,main="Important Sampling")

plot(v_log_marimps, ylim = xlim, main="Important Sampling")
hist(v_log_marimps,main="Important Sampling",xlim=xlim)

plot(v_log_mar_mc,main="Naive Monte Carlo",ylim=xlim)
hist(v_log_mar_mc,main="Naive Monte Carlo",xlim=xlim)

```



```
## comparing variance of importance sampling and naive sampling
var (v_log_marimps)
```

```
## [1] 0.0001784331
```

```
var (v_log_mar_mc)
```

```
## [1] 152.9014
```

Comparing Monte Carlo variance of Importance sampling and Laplace Approximation

```
mean (v_log_marimps)
```

```
## [1] -314.6463
```

```
bayes_inference_lap_gaussian(x,0,1,0,5)
```

```
## $estimates
```

```
## [1] 0.007028495 0.118606615
```

```
##
```

```
## $sds
```

```
## [1] 0.07936393 0.05000188
```

```
##
```

```
## $SIGMA
```

```
##           [,1]           [,2]
```

```
## [1,] 6.298633e-03 2.461681e-08
```

```
## [2,] 2.461681e-08 2.500188e-03
```

```
##
```

```
## $log_mar_lik
```

```
## [1] -314.6507
```

We see that Laplace approximation is really good for this problem, but it may not be good for other problems in which the posterior cannot be approximated by Gaussian.