

Statistical Methods for Research

Longhai Li

2025-10-31

Table of contents

1	Introduction to Statistical Methods for Research	1
	Welcome	1
	Table of Contents	1
2	A Quick Introduction to using R for Data Analysis	3
2.1	Basic R Objects and Operations	3
2.2	Import a dataset into R environment and Simple Operation	6
2.3	Create your own function	11
2.4	Include Images Saved in An External File	12
3	Simple Linear Regression	15
3.1	Overview of Simple Linear Regression	15
3.1.1	Least Squares Estimation	16
3.1.2	Residual and Sum of Squares Definitions	17
3.1.3	Coefficient of Determination (R^2)	18
3.1.4	F-test for Overall Significance	18
3.1.5	t-test for the Slope β_1	20
3.1.6	Prediction for a New Case x_0	21
3.2	Example 1: Vehicle Insurance Premium (warm-up)	21
3.2.1	Input data	21
3.2.2	Estimating regression coefficients	23
3.2.3	Residuals and fitted values (geometry picture)	23
3.2.4	SST, SSR, SSE and their meanings	25
3.2.5	Visual ANOVA on an RSS plot	26
3.2.6	R^2 , F and a compact ANOVA table	27
3.2.7	Sampling distributions via animation	28
3.3	Example 2: Oxygen Purity Data	29
3.3.1	Data	29
3.3.2	Fit and quick summary	30
3.3.3	Scatter with fitted line	31
3.3.4	Coefficient CIs and ANOVA	32
3.3.5	Mean-response and prediction bands	33
3.3.6	Residual diagnostics (assumptions check)	34
3.4	Correlation analysis (for comparison, not causation)	35
3.4.1	Data and scatter	35
3.4.2	Pearson correlation and test	36
3.5	What to report (checklist)	37

4	Multiple Linear Regression	39
4.1	An Example: Wire Bond Strength Dataset	39
4.1.1	Loading Data and Visualization	39
4.1.2	Model Fitting and Summary	42
4.1.3	Confidence Intervals and Model Components	43
4.2	RSS-based Inference: F-test, and adjusted R^2	44
4.2.1	RSS-Based Quantities	45
4.2.2	Remarks	46
4.2.3	A Simulation Study to Understand the Distributions of RSS	47
4.2.4	Example: Modelling Children Weight with Height and Age	48
4.2.5	Example: Wire bond strength	52
4.2.6	Relationship between t-test and partial F-test	53
4.3	Predictions for Mean Response and a Future Observation	54
4.3.1	Confidence Interval for Mean Response	54
4.3.2	Prediction Interval for a New Observation	54
4.4	Model Diagnostics	54
4.4.1	Residual Calculations	54
4.4.2	Residual Plots	56
4.5	Influential Observations	57
4.5.1	Plotting with the <code>olsrr</code> Package	58
4.6	Polynomial Regression	61
4.7	Handling Categorical Variables with Dummy Variables	63
4.7.1	Four Models Involving “sex”	65
4.7.2	Orders of Terms Matters in ANOVA and Warnings in Interpreting t-test Tables	78
4.8	Model Building	82
4.8.1	All Possible Regression	82
4.8.2	Automated Stepwise Procedures	83
4.9	Multicollinearity	86
4.9.1	A Simple Example	86
4.9.2	VIFs in the Wine Quality Data	88
4.9.3	VIFs in the Children Height Data	88
5	Logistic Regression	89
5.1	Odds as a Function of Probability	89
5.2	A Simulated Data	91
5.2.1	Fit a logistic model to the simulated data	92
5.3	Example of Coronary Heart Disease Data	93
5.3.1	Load a dataset	93
5.3.2	Fit Logistic Regression Model for a Single Variable	108
5.3.3	Fit Logistic Regression Model with all variables	113
5.4	Inference for Coefficients: Confidence Intervals and Covariance Matrix	114
5.5	Inference for Odds Ratios	115
5.5.1	Interpretation of Odds Ratios in Logistic Regression	115
5.5.2	Examples of Finding ORs and Their CIs for the CHD Dataset	119
5.6	Assessing Statistical Significance with Wilks’ Theorem (Analogue of F-test for OLS)	120

5.7	Assessing Predictive Effect-Size (Analogue to R_{adj}^2)	127
5.7.1	Understanding the Confusion Matrix and Metrics	127
5.7.2	Illustration with the Simulated Dataset	128
5.7.3	Application to the CHD Dataset	136
6	One-factor Design	143
6.1	Completely Randomized Design	143
6.1.1	Plasma Etching Experiment	143
6.2	Unbalanced Designs with nequal Sample Sizes	155
6.3	Randomized Complete Block Design	157
6.3.1	Vascular Graft Experiment	157
7	Two-Factor Factorial Design	165
7.1	Battery Design Experiment	165
7.2	Data Setup and Preparation	165
7.3	Exploratory Data Analysis and Visualization	166
7.4	Boxplots of Main Effects	167
7.5	Interaction Plot	167
7.6	Model Fitting and Analysis of Variance (ANOVA)	169
7.7	Model Adequacy Checks	170
7.8	Post-Hoc Analysis: Pairwise Comparisons	171
7.9	Tukey's HSD Test	171
7.10	Fisher's LSD Method	173

1 Introduction to Statistical Methods for Research

Welcome

This book contains lecture notes for **STAT 845: Statistical Methods for Research** at the **University of Saskatchewan**.

Table of Contents

2 A Quick Introduction to using R for Data Analysis

2.1 Basic R Objects and Operations

```
## create a vector
x <- 1:10
x <- seq(30,3, by = -2)
a <- c(66.32, 69.87, 70.12, 90.37, 50.08, 61.20, 65.00, 57.65)
d <- a[1]
a[1] <- 85.34

mean(a)
```

```
[1] 68.70375
```

```
ma <- mean(a)
## read a vector of numbers from a file
x <- scan("numbers.txt")
x2 <- scan("number2.txt")

## one can also read number without saving to a file
y <- scan(text = "7 8 9 10 11 12 13 13 14 17 17 45")

## create a matrix
A <- matrix(0, 4, 2)

A <- matrix(1:8, 4, 2)

A
```

```
      [,1] [,2]
[1,]     1     5
[2,]     2     6
[3,]     3     7
[4,]     4     8
```

2 A Quick Introduction to using R for Data Analysis

```
D <- matrix (a, 4, 2, byrow=T)
```

```
D <- matrix(1:8, 2, 4)
```

```
D
```

```
      [,1] [,2] [,3] [,4]
[1,]     1     3     5     7
[2,]     2     4     6     8
```

```
## create another matrix with all entry 0
```

```
B <- matrix (1:5000, 100, 50)
```

```
## assign a number to B
```

```
B[2,4] <- 45
```

```
B[1,]
```

```
[1]      1  101  201  301  401  501  601  701  801  901 1001 1101 1201 1301 1401
[16] 1501 1601 1701 1801 1901 2001 2101 2201 2301 2401 2501 2601 2701 2801 2901
[31] 3001 3101 3201 3301 3401 3501 3601 3701 3801 3901 4001 4101 4201 4301 4401
[46] 4501 4601 4701 4801 4901
```

```
B[,1]
```

```
[1]      1      2      3      4      5      6      7      8      9     10     11     12     13     14     15     16     17     18
[19]     19     20     21     22     23     24     25     26     27     28     29     30     31     32     33     34     35     36
[37]     37     38     39     40     41     42     43     44     45     46     47     48     49     50     51     52     53     54
[55]     55     56     57     58     59     60     61     62     63     64     65     66     67     68     69     70     71     72
[73]     73     74     75     76     77     78     79     80     81     82     83     84     85     86     87     88     89     90
[91]     91     92     93     94     95     96     97     98     99    100
```

```
B[1,] <- 1:50
```

```
## create a list
```

```
E <- list (newa = a, newA = A)
```

```
## list the names of components
```

```
names (E)
```

```
[1] "newa" "newA"
```

```
## to look at the component of E
E$newA
```

```
      [,1] [,2]
[1,]    1    5
[2,]    2    6
[3,]    3    7
[4,]    4    8
```

```
E$newA <- 10:17

## create a dataframe
scores <- c(30, 45, 50)
names <- c("Peter", "John", "Alice")
stat245_scores <- data.frame(names, scores)
stat245_scores
```

```
  names scores
1 Peter     30
2  John     45
3 Alice     50
```

```
stat245_scores$names
```

```
[1] "Peter" "John"  "Alice"
```

```
stat245_scores$scores [1] <- 40
stat245_scores
```

```
  names scores
1 Peter     40
2  John     45
3 Alice     50
```

```
stat245_scores$perc <- stat245_scores$scores/50 * 100
stat245_scores
```

```
  names scores perc
1 Peter     40   80
2  John     45   90
3 Alice     50  100
```

2 A Quick Introduction to using R for Data Analysis

```
stat245_scores$adj <- stat245_scores$perc + 10
stat245_scores
```

```
  names scores perc adj
1 Peter     40   80  90
2  John     45   90 100
3 Alice     50  100 110
```

```
#####
```

2.2 Import a dataset into R environment and Simple Operation

```
#####
```

```
## import myagpop.csv into an R data frame called 'myagpop'
agpop <- read.csv("agpop.csv")
```

```
## Now, we can use the data:
```

```
## preview agpop
head (agpop)
```

	county	state	acres92	acres87	acres82	farms92	farms87	farms82
1	ALEUTIAN ISLANDS AREA	AK	683533	726596	764514	26	27	28
2	ANCHORAGE AREA	AK	47146	59297	256709	217	245	223
3	FAIRBANKS AREA	AK	141338	154913	204568	168	175	170
4	JUNEAU AREA	AK	210	214	127	8	8	12
5	KENAI PENINSULA AREA	AK	50810	85712	98035	93	119	137
6	AUTAUGA COUNTY	AL	107259	116050	145044	322	388	453
	largef92	largef87	largef82	smallf92	smallf87	smallf82	region	
1	14	16	20	6	4	1	W	
2	9	10	11	41	52	38	W	
3	25	28	21	12	18	25	W	
4	0	0	0	5	4	8	W	
5	9	18	17	12	18	19	W	
6	25	32	32	8	19	17	S	

```
## look at the variable name
colnames (agpop)
```

2.2 Import a dataset into R environment and Simple Operation

```
[1] "county"    "state"      "acres92"    "acres87"    "acres82"    "farms92"
[7] "farms87"    "farms82"    "largef92"   "largef87"   "largef82"   "smallf92"
[13] "smallf87"   "smallf82"   "region"
```

```
## find number of cols
ncol (agpop)
```

```
[1] 15
```

```
## find number of rows
nrow (agpop)
```

```
[1] 3078
```

```
## access a certain row
agpop [2, ]
```

```
      county state acres92 acres87 acres82 farms92 farms87 farms82 largef92
2 ANCHORAGE AREA   AK  47146  59297 256709     217     245     223         9
  largef87 largef82 smallf92 smallf87 smallf82 region
2      10      11      41      52      38      W
```

```
## access a certain column
agpop [1:20, "acres92"] ## equivalent to
```

```
[1] 683533 47146 141338    210 50810 107259 167832 177189 48022 137426
[11] 144799 96427 73841 109555 121504 99466 67950 61426 68478 47200
```

```
agpop$acres92[1:20]
```

```
[1] 683533 47146 141338    210 50810 107259 167832 177189 48022 137426
[11] 144799 96427 73841 109555 121504 99466 67950 61426 68478 47200
```

```
agpop$largef92[1:20]
```

```
[1] 14 9 25 0 9 25 24 40 6 9 29 18 4 22 24 8 9 13 4 5
```

```
## find mean of acres92
mean (agpop $acres92)
```

```
[1] 306677
```


2 A Quick Introduction to using R for Data Analysis

```
## find sd of acres92  
sd (agpop $acres92)
```

```
[1] 424686.7
```

```
agpop_AK <- agpop [agpop$state == "AK", ]  
agpop_AK <- subset (agpop, state == "AK")  
agpop_W <- subset (agpop, region == "W")  
agpop_largefarm <- subset (agpop, largef92 > 10)  
  
hist (agpop$acres92)
```


Histogram of agpop\$acres92



Produce Plots

```
#pdf ("hist_acres92.pdf") ## use this command and dev.off to save the output to a file
hist (agpop$acres92)
```


Histogram of agpop\$acres92



```
#dev.off()

#jpeg ("agpop_acres_87v92.jpg")

plot (agpop$acres87, agpop$acres92)
abline (a = 0, b = 1)
```

```
#dev.off()## this is used to close the jpeg file
```

2.3 Create your own function

```
### data is a matrix or data.frame
means_col <- function (data)
{
  n <- ncol (data)
  cmeans <- rep (NA, n)
  for (j in 1:n)
  {
    cmeans[j] <- mean (data[,j])
  }
  cmeans
}
```

```
}  
  
### apply function  
means_col (agpop[, 3:13])
```

```
[1] 306676.97141 313016.37817 320193.69298 625.50357 678.28428  
[6] 728.06238 56.17674 54.86160 52.62248 54.09227  
[11] 59.53769
```

```
### R built-in function  
colMeans (agpop[, 3:13])
```

acres92	acres87	acres82	farms92	farms87	farms82
306676.97141	313016.37817	320193.69298	625.50357	678.28428	728.06238
largef92	largef87	largef82	smallf92	smallf87	
56.17674	54.86160	52.62248	54.09227	59.53769	

2.4 Include Images Saved in An External File

Using the following R code to include your images saved in an external file.

```
knitr::include_graphics("handwriting.png")
```

Q1:

$$a+b=c$$

Q2:

$$c=1+2$$

You can hide the above R code by setting “echo=FALSE” for the r chunk. For example, I will include the image once again as follows:

Q1:
 $a+b=c$
Q2 :
 $c=1+2$

Figure 2.1: This is a figure inserted from the file called “handwriting.png”

3 Simple Linear Regression

A Simulation Illustration with R

```
require("knitr")
knitr::opts_chunk$set(
  comment = "#",
  fig.width = 8,
  fig.height = 6,
  cache = TRUE
)
set.seed(47)

options(sim_rebuild=FALSE)
```

3.1 Overview of Simple Linear Regression

To make the simple linear regression model concrete, let's first visualize a simulated dataset that follows

$$Y_i = \beta_0 + \beta_1 X_i + \varepsilon_i, \quad \varepsilon_i \sim \mathcal{N}(0, \sigma^2).$$

Here, β_0 is the intercept, β_1 is the slope, and ε_i represents random noise.


```
set.seed(2025)
n <- 40
beta0 <- 2; beta1 <- 1.5; sigma <- 2
x <- runif(n, 0, 10)
y <- beta0 + beta1 * x + rnorm(n, 0, sigma)
dat <- data.frame(x, y)

fit <- lm(y ~ x, data = dat)

plot(x, y, pch = 19, col = "steelblue",
     xlab = "Predictor X", ylab = "Response Y",
     main = "Simulated Data with Fitted Linear Regression Line")
```

```
abline(fit, col = "red", lwd = 2)
legend("topleft", legend = c("Observed data", "Fitted line"),
      pch = c(19, NA), lty = c(NA, 1), col = c("steelblue", "red"), bty = "n")
```

Simulated Data with Fitted Linear Regression Line



The scatterplot shows data points scattered around a line — the red line is the fitted regression model.

3.1.1 Least Squares Estimation

Goal: Find $\hat{\beta}_0$ and $\hat{\beta}_1$ that minimize

$$\text{SSE} = \sum_{i=1}^n (y_i - \hat{\beta}_0 - \hat{\beta}_1 x_i)^2.$$

Solutions:

$$\hat{\beta}_1 = \frac{\sum_i (x_i - \bar{x})(y_i - \bar{y})}{\sum_i (x_i - \bar{x})^2} = \frac{S_{xy}}{S_{xx}}, \quad \hat{\beta}_0 = \bar{y} - \hat{\beta}_1 \bar{x}.$$

Here

$$S_{xy} = \sum_i (x_i - \bar{x})(y_i - \bar{y}), \quad S_{xx} = \sum_i (x_i - \bar{x})^2.$$

Shortcut (computational) formulas:

$$S_{xy} = \sum_i x_i y_i - n \bar{x} \bar{y}, \quad S_{xx} = \sum_i x_i^2 - n \bar{x}^2.$$

Interpretation:

- $\hat{\beta}_1$ measures the estimated change in Y for each unit increase in X .
- $\hat{\beta}_0$ represents the fitted value of Y when $X = 0$.

3.1.2 Residual and Sum of Squares Definitions

Let $\hat{y}_i = \hat{\beta}_0 + \hat{\beta}_1 x_i$ and $e_i = y_i - \hat{y}_i$.

Symbol	Definition	Computing Formula (in terms of S_{xx}, S_{xy} , etc.)
SST	Total Sum of Squares	$\sum_i (y_i - \bar{y})^2 = S_{yy} = \sum_i y_i^2 - n \bar{y}^2$
SSR	Regression Sum of Squares	$\sum_i (\hat{y}_i - \bar{y})^2 = \hat{\beta}_1^2 S_{xx} = \frac{S_{xy}^2}{S_{xx}}$
SSE	Error (Residual) Sum of Squares	$\sum_i (y_i - \hat{y}_i)^2 = S_{yy} - \frac{S_{xy}^2}{S_{xx}}$

Identity:

$$\text{SST} = \text{SSR} + \text{SSE}.$$

Here,

$$S_{xx} = \sum_i (x_i - \bar{x})^2 = \sum_i x_i^2 - n \bar{x}^2, \quad S_{yy} = \sum_i (y_i - \bar{y})^2 = \sum_i y_i^2 - n \bar{y}^2, \quad S_{xy} = \sum_i (x_i - \bar{x})(y_i - \bar{y}) = \sum_i x_i y_i - n \bar{x} \bar{y}.$$

3.1.3 Coefficient of Determination (R^2)

Measures the proportion of total variation in Y explained by X :

$$R^2 = \frac{SSR}{SST} = 1 - \frac{SSE}{SST}.$$

Interpretation:

- $R^2 = 1$ means perfect linear fit;
- $R^2 = 0$ means the model explains none of the variation.

3.1.4 F-test for Overall Significance

Tests whether X is linearly related to Y .

Hypotheses:

$$H_0 : \beta_1 = 0 \quad \text{vs.} \quad H_A : \beta_1 \neq 0.$$

Test Statistic:

$$F = \frac{MSR}{MSE} = \frac{SSR/1}{SSE/(n-2)} \sim F_{1,n-2} \quad (H_0).$$

p-value approach for observe F^{obs} :

Given the observed statistic F^{obs} with $(1, n-2)$ df,

$$p\text{-value} = \Pr(F_{1,n-2} \geq F^{\text{obs}}) = \text{pf}(F^{\text{obs}}, 1, n-2, \text{lower.tail} = \text{FALSE}).$$

```
## -- Inputs (provide these from your analysis context) -----
## n    <- ...    # sample size
## SSR  <- ...    # regression sum of squares
## SSE  <- ...    # error sum of squares
n    <- 20
SSR  <- 5
SSE  <- 40

df1  <- 1
df2  <- n - 2
Fobs <- (SSR/df1) / (SSE/df2)      # observed F
pval <- pf(Fobs, df1 = df1, df2 = df2, lower.tail = FALSE)
pval
```



```
[1] 0.1509505
```

```
## -- Plot F density and shade the p-value tail (with proper annotations) -----
xmax <- max(qf(0.995, df1, df2), Fobs * 1.2) # extra space for labels
peak <- max(df(seq(0, xmax, length.out = 500), df1, df2))

## Density curve
curve(df(x, df1, df2), from = 0, to = xmax,
      xlab = "F", ylab = "Density",
      main = sprintf("F(%d, %d) density | observed F = %.3f", df1, df2, Fobs))


## Shade right tail (p-value region)
xs <- seq(Fobs, xmax, length.out = 300)
ys <- df(xs, df1, df2)
polygon(c(Fobs, xs, xmax), c(0, ys, 0),
       col = rgb(0, 0, 0, 0.18), border = NA)

## Vertical line at Fobs (optional visual aid)
abline(v = Fobs, lwd = 2)

## ---- Annotation for F^obs pointing to the x-axis value (Fobs, 0) -----
x_txt_F <- Fobs + 0.06 * xmax
y_txt_F <- 0.45 * peak
arrows(x0 = x_txt_F, y0 = y_txt_F, x1 = Fobs, y1 = 0,
      length = 0.08, lwd = 1.5)
text(x_txt_F, y_txt_F,
     labels = bquote(F^{obs} == .(format(Fobs, digits = 3))),
     pos = 4)

## ---- Annotation for p-value pointing into the shaded tail -----
x_tip_p <- (Fobs + xmax) / 1.7
y_tip_p <- df(x_tip_p, df1, df2)
x_txt_p <- Fobs + 0.08 * xmax
y_txt_p <- 0.80 * peak
arrows(x0 = x_txt_p, y0 = y_txt_p, x1 = x_tip_p, y1 = y_tip_p,
      length = 0.08, lwd = 1.5)
text(x_txt_p, y_txt_p,
     labels = bquote(p == .(format(pval, digits = 4, scientific = TRUE))),
     pos = 4)
```

F(1, 18) density | observed F = 2.250



3.1.5 t-test for the Slope β_1

Equivalent to the F -test in simple regression since $t^2 = F$.

Formula:

$$t = \frac{\hat{\beta}_1}{\text{SE}(\hat{\beta}_1)}, \quad \text{SE}(\hat{\beta}_1) = \sqrt{\frac{\hat{\sigma}^2}{\sum_i (x_i - \bar{x})^2}}, \quad \hat{\sigma}^2 = \frac{\text{SSE}}{n-2}.$$

Distribution:

$$t \sim t_{n-2} \quad (H_0 : \beta_1 = 0).$$

3.1.6 Prediction for a New Case x_0

Predicted mean response:

$$\hat{y}(x_0) = \hat{\beta}_0 + \hat{\beta}_1 x_0.$$

95% Confidence interval for mean response:

$$\hat{y}(x_0) \pm t_{1-\alpha/2, n-2} \hat{\sigma} \sqrt{\frac{1}{n} + \frac{(x_0 - \bar{x})^2}{\sum_i (x_i - \bar{x})^2}}.$$

95% Prediction interval for a new observation:

$$\hat{y}(x_0) \pm t_{1-\alpha/2, n-2} \hat{\sigma} \sqrt{1 + \frac{1}{n} + \frac{(x_0 - \bar{x})^2}{\sum_i (x_i - \bar{x})^2}}.$$

Summary Cheat Sheet

Concept	Key Formula
Model	$Y_i = \beta_0 + \beta_1 X_i + \varepsilon_i$
LS Estimates	$\hat{\beta}_1 = S_{xy}/S_{xx}, \hat{\beta}_0 = \bar{y} - \hat{\beta}_1 \bar{x}$
Decomposition	$SST = SSR + SSE$
R^2	$R^2 = 1 - SSE/SST$
F -test	$F = (SSR/1)/(SSE/(n-2))$
t -test	$t = \hat{\beta}_1 / SE(\hat{\beta}_1)$
Prediction	$\hat{y}(x_0) = \hat{\beta}_0 + \hat{\beta}_1 x_0$

3.2 Example 1: Vehicle Insurance Premium (warm-up)

We examine premiums y vs. driving amount x . The scatterplot hints at a **downward** trend.

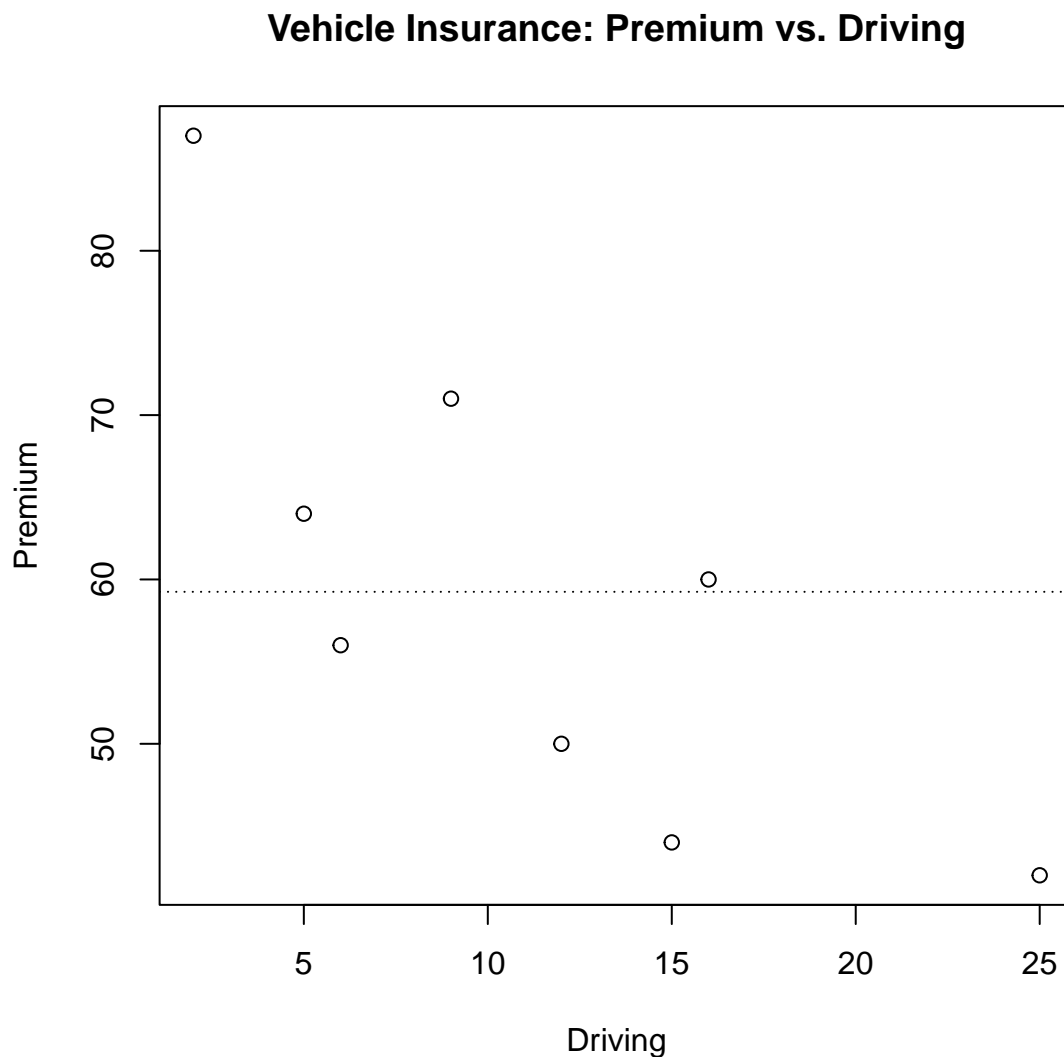
3.2.1 Input data

```
issu <- data.frame(
  driving = c(5, 2, 12, 9, 15, 6, 25, 16),
  premium = c(64, 87, 50, 71, 44, 56, 42, 60)
)
```

3 Simple Linear Regression

```
y <- issu$premium
x <- issu$driving
xbar <- mean(x); ybar <- mean(y); n <- length(y)

plot(x, y, xlab = "Driving", ylab = "Premium",
     main = "Vehicle Insurance: Premium vs. Driving")
abline(h = ybar, lty = 3)
```




Narrative. The horizontal line at \bar{y} represents the intercept-only model. Any fitted line that tilts away from this must earn its keep by reducing residual variation enough to offset the loss of one degree of freedom.

3.2.2 Estimating regression coefficients

```
fit.issu <- lm(y ~ x)
plot(x, y, xlab = "Driving", ylab = "Premium",
     main = "Fitted Simple Linear Regression")
abline(fit.issu, lwd = 2)
```

Fitted Simple Linear Regression



The slope estimate $\hat{\beta}_1$ captures the **marginal change in premium per unit of driving** (units of y per unit of x). Inference on β_1 tells us whether the pattern rises above noise.

3.2.3 Residuals and fitted values (geometry picture)

Let $\hat{y}_i = \hat{\beta}_0 + \hat{\beta}_1 x_i$ and $\tilde{y}_i = \bar{y}$. Residuals are $e_i = y_i - \hat{y}_i$ (model) and $y_i - \bar{y}$ (null). Visualizing all three clarifies the ANOVA identity.

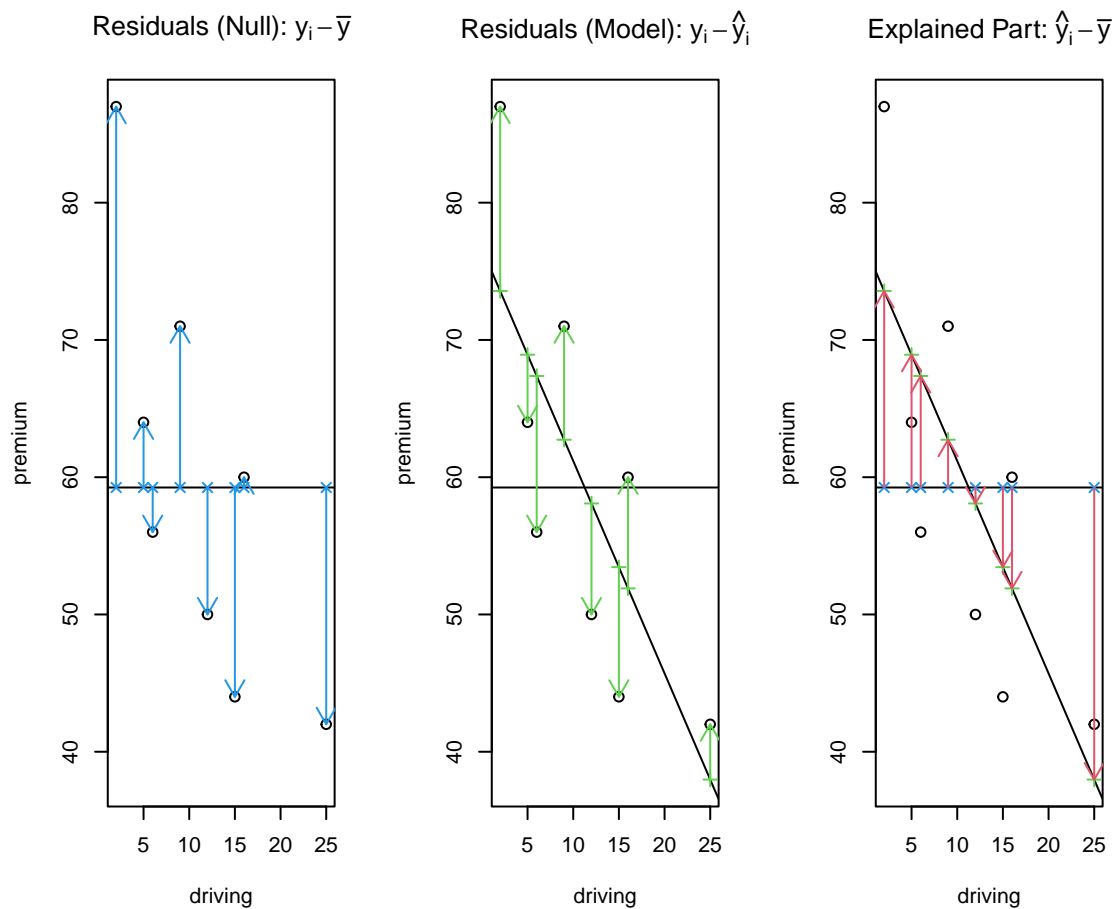
3 Simple Linear Regression

```
beta0 <- coef(fit.issu)[1]
beta1 <- coef(fit.issu)[2]
fitted1 <- beta0 + beta1 * x
fitted0 <- rep(ybar, n)
residual1 <- y - fitted1
residual0 <- y - fitted0

data.frame(y, fitted0, residual0, fitted1, residual1,
           diff.fitted = fitted1 - fitted0)
```

	y	fitted0	residual0	fitted1	residual1	diff.fitted
1	64	59.25	4.75	68.92243	-4.922425	9.672425
2	87	59.25	27.75	73.56519	13.434811	14.315189
3	50	59.25	-9.25	58.08931	-8.089309	-1.160691
4	71	59.25	11.75	62.73207	8.267927	3.482073
5	44	59.25	-15.25	53.44654	-9.446545	-5.803455
6	56	59.25	-3.25	67.37484	-11.374837	8.124837
7	42	59.25	-17.25	37.97066	4.029335	-21.279335
8	60	59.25	0.75	51.89896	8.101043	-7.351043

3.2 Example 1: Vehicle Insurance Premium (warm-up)



3.2.4 SST, SSR, SSE and their meanings

- $SST = \sum (y_i - \bar{y})^2$ quantifies **total** variability around the mean.
- $SSR = \sum (\hat{y}_i - \bar{y})^2$ is the part **explained by x** .
- $SSE = \sum (y_i - \hat{y}_i)^2$ is the **leftover** (unexplained) variability.

```
SST <- sum((y - fitted0)^2); SST
```

```
[1] 1557.5
```

```
SSE <- sum((y - fitted1)^2); SSE
```

```
[1] 639.0065
```

```
SSR <- SST - SSE; SSR
```

3 Simple Linear Regression

```
[1] 918.4935
```

Direct check: $SSR = \sum(\hat{y}_i - \bar{y})^2$.

```
sum((fitted1 - fitted0)^2)
```

```
[1] 918.4935
```

3.2.5 Visual ANOVA on an RSS plot

We place the **residual sum of squares** against model dimension to show the trade-off between fit and df.

```
## Recompute cleanly
SST <- sum((y - mean(y))^2)
SSE <- sum(resid(fit.issu)^2)
SSR <- SST - SSE
df_SSR <- 1
df_SSE <- n - 2

par(mar = c(6, 4, 4, 2) + 0.1)
plot(c(1, 2, n), c(SST, SSE, 0), type = "b", pch = 19,
     xlab = "Number of Parameters in Model",
     ylab = "Residual Sum of Squares (RSS)",
     main = "ANOVA Geometry on RSS vs. Model Size",
     xlim = c(0, 14), ylim = c(-400, SST * 1.1), xaxt = "n")
axis(1, at = c(1, 2, n), labels = c("1 (Intercept)", "2 (+Slope)", paste(n, "(Saturated)")))
abline(h = seq(0, 2000, by = 100), lty = 3, col = "grey")

par(xpd = TRUE)
arrows(9, 0, 9, SSE, col = "blue", code = 3, angle = 90, length = 0.1, lwd = 2)
text(9, SSE/2, "SSE", col = "blue", pos = 4, font = 2, cex = 1.2)

arrows(9, SSE, 9, SST, col = "red", code = 3, angle = 90, length = 0.1, lwd = 2)
text(9, (SST + SSE)/2, "SSR", col = "red", pos = 4, font = 2, cex = 1.2)

arrows(2, -200, n, -200, col = "blue", code = 3, angle = 90, length = 0.1, lwd = 2)
text((2 + n)/2, -250, paste("df_SSE =", df_SSE), col = "blue", font = 2)


arrows(1, -200, 2, -200, col = "red", code = 3, angle = 90, length = 0.1, lwd = 2)
text(1.5, -250, paste("df_SSR =", df_SSR), col = "red", font = 2)
par(xpd = FALSE)

f_value <- (SSR/df_SSR) / (SSE/df_SSE)
```


3.2 Example 1: Vehicle Insurance Premium (warm-up)

```
p_value <- pf(f_value, df1 = df_SSR, df2 = df_SSE, lower.tail = FALSE)
legend("topright",
      legend = c(sprintf("F-statistic: %.2f", f_value),
                  sprintf("p-value: %.3f", p_value)),
      title = "ANOVA Results", bty = "o", cex = 0.9)
```

ANOVA Geometry on RSS vs. Model Size



3.2.6 R^2 , F and a compact ANOVA table

```
R2 <- SSR / SST; R2
```

```
[1] 0.5897229
```

3 Simple Linear Regression

```
f <- (SSR/1) / (SSE/(n-2)); f
```

```
[1] 8.624264
```

```
pvf <- pf(f, df1 = 1, df2 = n-2, lower.tail = FALSE); pvf
```

```
[1] 0.0260588
```

```
Ftable <- data.frame(
  Source = c("Regression", "Error"),
  df      = c(1, n - 2),
  SS      = c(SSR, SSE),
  MS      = c(SSR/1, SSE/(n-2)),
  F       = c(f, NA),
  pvalue  = c(pvf, NA),
  R2part  = c(SSR, SSE) / SST
)
Ftable
```

	Source	df	SS	MS	F	pvalue	R2part
1	Regression	1	918.4935	918.4935	8.624264	0.0260588	0.5897229
2	Error	6	639.0065	106.5011	NA	NA	0.4102771

A call to `anova()` reproduces the same test:

```
anova(fit.issu)
```

Analysis of Variance Table

Response: y

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
x	1	918.49	918.49	8.6243	0.02606 *
Residuals	6	639.01	106.50		

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

3.2.7 Sampling distributions via animation

Under $H_0 : \beta_1 = 0$, F follows $F_{1,n-2}$. Under H_A , the distribution shifts right (noncentral F).




Figure 3.1: Simulation under H0: animated GIF (HTML) and static PNG (PDF).

3.2.7.1 Null world (H_0 true)

3.2.7.2 Alternative world (H_1 true)

3.3 Example 2: Oxygen Purity Data

We model oxygen purity y as a function of hydrocarbon level x and report both **mean response** and **prediction uncertainty**.

3.3.1 Data

```
x <- c(0.99, 1.02, 1.15, 1.29, 1.46, 1.36, 0.87, 1.23, 1.55, 1.40, 1.19,
       1.15, 0.98, 1.01, 1.11, 1.20, 1.26, 1.32, 1.43, 0.95)
y <- c(90.01, 89.05, 91.43, 93.74, 96.73, 94.45, 87.59, 91.77, 99.42, 93.65,
       93.54, 92.52, 90.56, 89.54, 89.85, 90.39, 93.25, 93.41, 94.98, 87.33)
n <- length(x); n
```

[1] 20

3 Simple Linear Regression




Figure 3.2: Simulation under HA (slope = -2): animated GIF for HTML, static PNG for PDF.

```
purity.data <- data.frame(x = x, y = y)
head(purity.data)
```

```
      x      y
1 0.99 90.01
2 1.02 89.05
3 1.15 91.43
4 1.29 93.74
5 1.46 96.73
6 1.36 94.45
```

3.3.2 Fit and quick summary

```
fit <- lm(y ~ x, data = purity.data)
summary(fit)
```

Call:

```
lm(formula = y ~ x, data = purity.data)
```

Residuals:

Min	1Q	Median	3Q	Max
-----	----	--------	----	-----

```
-1.83029 -0.73334 0.04497 0.69969 1.96809
```

Coefficients:

```
      Estimate Std. Error t value Pr(>|t|)
(Intercept)  74.283      1.593   46.62 < 2e-16 ***
x             14.947      1.317   11.35 1.23e-09 ***
```

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.087 on 18 degrees of freedom


Multiple R-squared: 0.8774, Adjusted R-squared: 0.8706

F-statistic: 128.9 on 1 and 18 DF, p-value: 1.227e-09

Interpretation. The slope's sign gives the direction of association; its t test (or F with 1 df) assesses evidence for a trend. Look at $\hat{\sigma}$ for noise scale and R^2 for variance explained.

3.3.3 Scatter with fitted line

```
plot(purity.data$x, purity.data$y,
     xlab = "Hydrocarbon level (x)", ylab = "Purity (y)",
     main = "Oxygen Purity vs Hydrocarbon Level")
abline(fit, col = "red", lwd = 2)
```



3.3.4 Coefficient CIs and ANOVA

```
confint(fit, level = 0.95)
```

	2.5 %	97.5 %
(Intercept)	70.93555	77.63108
x	12.18107	17.71389

```
anova(fit)
```

Analysis of Variance Table

```

Response: y
      Df Sum Sq Mean Sq F value    Pr(>F)
x         1 152.13  152.127   128.86 1.227e-09 ***
Residuals 18   21.25    1.181
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```

3.3.5 Mean-response and prediction bands

The **mean-response CI** narrows near \bar{x} and widens at the extremes; the **prediction band** is wider by the irreducible noise term.

```

x0 <- seq(min(purity.data$x), max(purity.data$x), length = 50)
newdata <- data.frame(x = x0)

est.mean <- predict(fit, newdata = newdata, interval = "confidence", level = 0.95)
pred.new <- predict(fit, newdata = newdata, interval = "prediction", level = 0.95)


```

```

plot(purity.data$x, purity.data$y,
     xlab = "Hydrocarbon level (x)", ylab = "Purity (y)",
     main = "Regression Line with Confidence and Prediction Bands")
abline(fit)
matlines(x0, est.mean[, 2:3], col = "blue", lty = 2, lwd = 2)
matlines(x0, pred.new[, 2:3], col = "red", lty = 3, lwd = 2)
legend("topleft", c("Confidence Bands (mean)", "Prediction Bands (new y)"),
     col = c("blue", "red"), lty = 2:3, bty = "n")

```

Regression Line with Confidence and Prediction Bands




3.3.6 Residual diagnostics (assumptions check)

We look for **no pattern** in residuals vs. fits and **approximate straightness** in the Q-Q plot.

```
pred <- fitted.values(fit)
e <- resid(fit)
d <- e / summary(fit)$sigma

par(mfrow = c(2,2))
plot(purity.data$x, purity.data$y, xlab = "x", ylab = "y"); abline(fit)
qqnorm(d, main = "Normal Q-Q"); qqline(d)
plot(pred, d, xlab = "Fitted", ylab = "Std. residuals", main = "Residuals vs Fits"); abline(h = 0)
plot(1:n, d, xlab = "Order", ylab = "Std. residuals", main = "Residuals vs Order"); abline(h = 0)
```


3.4 Correlation analysis (for comparison, not causation)



```
par(mfrow = c(1,1))
```

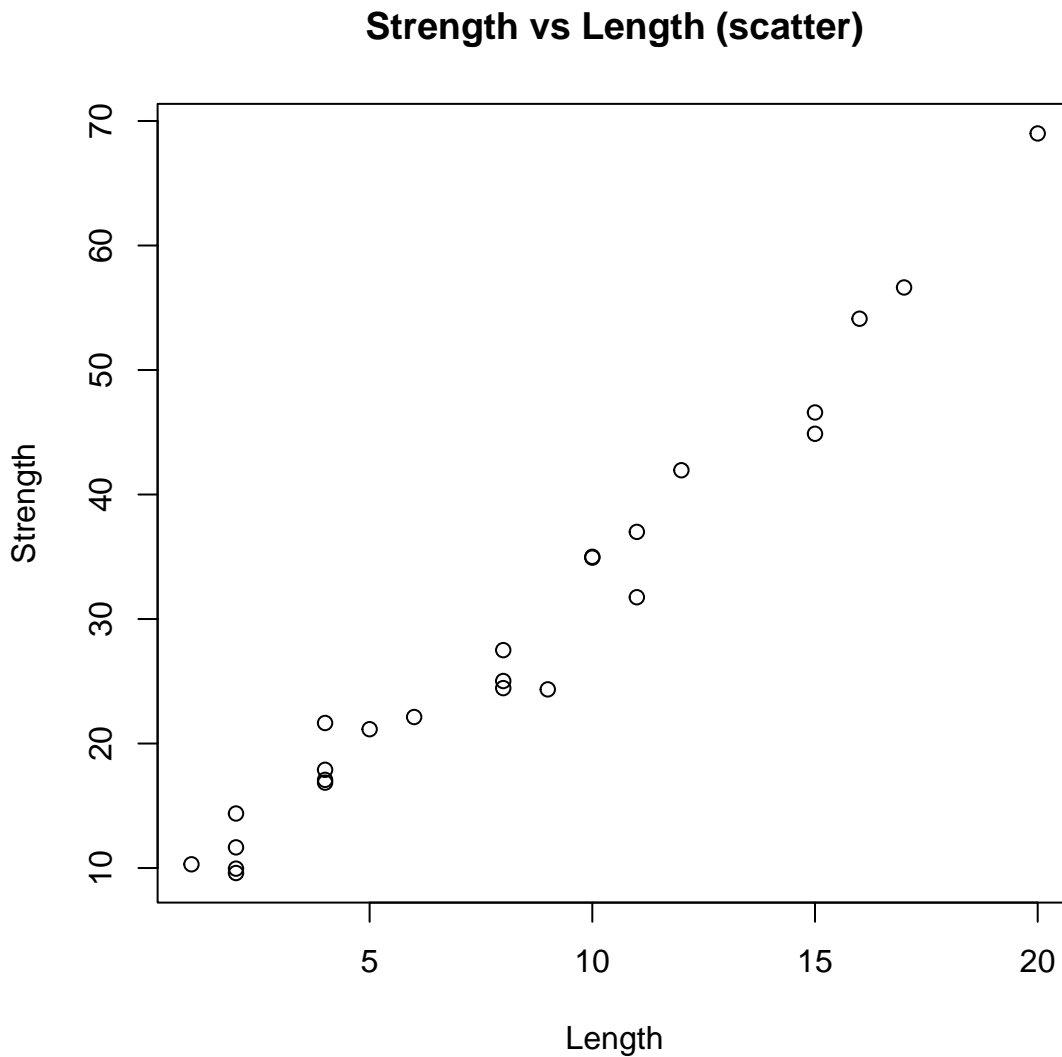
3.4 Correlation analysis (for comparison, not causation)

Correlation summarizes linear association without fitting a line or making model assumptions.

3.4.1 Data and scatter

3 Simple Linear Regression

```
strength <- c(9.95,24.45,31.75,35.00,25.02,16.86,14.38,9.60,24.35,  
             27.50,17.08,37.00,41.95,11.66,21.65,17.89,69.00,10.30,  
             34.93,46.59,44.88,54.12,56.63,22.13,21.15)  
length <- c(2,8,11,10,8,4,2,2,9,8,4,11,12,2,4,4,20,1,10,  
           15,15,16,17,6,5)  
plot(length, strength, xlab = "Length", ylab = "Strength",  
     main = "Strength vs Length (scatter)")
```



3.4.2 Pearson correlation and test

```
cor(strength, length)
```

```
[1] 0.9818118
```

```
cor.test(strength, length)
```

```
Pearson's product-moment correlation
```

```
data: strength and length
t = 24.801, df = 23, p-value < 2.2e-16
alternative hypothesis: true correlation is not equal to 0
95 percent confidence interval:
 0.9585414 0.9920735
sample estimates:
      cor
0.9818118
```

Note. A large $|r|$ and small p indicate linear association; regression further quantifies the slope and supports prediction, with diagnostics to check assumptions.

3.5 What to report (checklist)

- Estimated line $\hat{y} = \hat{\beta}_0 + \hat{\beta}_1 x$ with units.
- t/F test for slope, p -value and CI for β_1 .
- R^2 and $\hat{\sigma}$ (RMSE) for fit quality.
- Mean-response and prediction intervals at substantively relevant x_0 .
- Residual diagnostics and any remedies (transformations, robust methods) if needed.

4 Multiple Linear Regression

4.1 An Example: Wire Bond Strength Dataset

4.1.1 Loading Data and Visualization

Note: You must change the file paths in the `read.csv()` functions below to match the location of the files on your computer (for example `C:\\Users\\<YourUsername>\\Documents` on Windows).

```
## Read data. Change the path as necessary.  
## Example: bond.data <- read.csv("wire-bond.csv")  
bond.data <- read.csv("wire-bond.csv")  
  
## This will now be automatically rendered as a paged table  
bond.data
```

	strength	length	height
1	9.95	2	50
2	24.45	8	110
3	31.75	11	120
4	35.00	10	550
5	25.02	8	295
6	16.86	4	200
7	14.38	2	375
8	9.60	2	52
9	24.35	9	100
10	27.50	8	300
11	17.08	4	412
12	37.00	11	400
13	41.95	12	500
14	11.66	2	360
15	21.65	4	205
16	17.89	4	400
17	69.00	20	600
18	10.30	1	585
19	34.93	10	540
20	46.59	15	250
21	44.88	15	290

4 Multiple Linear Regression

22	54.12	16	510
23	56.63	17	590
24	22.13	6	100
25	21.15	5	400

2D Visualization


```
par(mfrow = c(1, 3), mar = c(5, 4, 2, 1))

## 1) length vs strength
i1 <- which(!is.na(bond.data$length) & !is.na(bond.data$strength))
plot(bond.data$length[i1], bond.data$strength[i1],
      xlab = "Wire Length", ylab = "Pull strength", pch = 19)
text(bond.data$length[i1], bond.data$strength[i1],
      labels = i1, pos = 1, offset = 0.4, cex = 0.75)

## 2) height vs strength
i2 <- which(!is.na(bond.data$height) & !is.na(bond.data$strength))
plot(bond.data$height[i2], bond.data$strength[i2],
      xlab = "Die height", ylab = "Pull strength", pch = 19)
text(bond.data$height[i2], bond.data$strength[i2],
      labels = i2, pos = 1, offset = 0.4, cex = 0.75)

## 3) height vs length
i3 <- which(!is.na(bond.data$height) & !is.na(bond.data$length))
plot(bond.data$height[i3], bond.data$length[i3],
      xlab = "Die height", ylab = "Length", pch = 19)
text(bond.data$height[i3], bond.data$length[i3],
      labels = i3, pos = 1, offset = 0.4, cex = 0.75)
```

4.1 An Example: Wire Bond Strength Dataset



3D Visualize


```
library(scatterplot3d)

par(mfrow = c(1,1))
s3d <- with(bond.data, scatterplot3d(
  x = length,
  y = height,
  z = strength,
  pch = 19,
  color = "steelblue",
  main = "3D Scatterplot: Strength vs. Length and Height",
  xlab = "Length",
  ylab = "Height",
```

4 Multiple Linear Regression

```
zlab = "Strength",  
angle = 60  
)  
  
fit <- lm(strength ~ length + height, data = bond.data)  
s3d$plane3d(fit, lty.box = "solid")
```

3D Scatterplot: Strength vs. Length and Height



4.1.2 Model Fitting and Summary

We fit a multiple linear regression model with `strength` as the response variable and `length` and `height` as predictors.

```
fit <- lm(strength ~ length + height, data = bond.data)  
summary(fit)
```



```

Call:
lm(formula = strength ~ length + height, data = bond.data)

Residuals:
    Min       1Q   Median       3Q      Max
-3.865 -1.542 -0.362  1.196  5.841

Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept)  2.263791    1.060066   2.136 0.044099 *
length       2.744270    0.093524  29.343 < 2e-16 ***
height       0.012528    0.002798   4.477 0.000188 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 2.288 on 22 degrees of freedom
Multiple R-squared:  0.9811,    Adjusted R-squared:  0.9794
F-statistic: 572.2 on 2 and 22 DF,  p-value: < 2.2e-16

```

The summary provides the ANOVA F-test for overall significance, R^2 , adjusted R^2 , and t-tests for individual coefficients.

4.1.3 Confidence Intervals and Model Components

```

## Confidence intervals
confint(fit)

              2.5 %      97.5 %
(Intercept) 0.065348613 4.46223426
length      2.550313061 2.93822623
height      0.006724246 0.01833138

## Fitted values and residuals
pred <- fitted.values(fit)
e <- resid(fit)
data.frame(y = bond.data$strength, y.hat = pred, e = e)

   y    y.hat      e
1  9.95  8.378721  1.57127871
2 24.45 25.596008 -1.14600783
3 31.75 33.954095 -2.20409488

```

4 Multiple Linear Regression

```
4 35.00 36.596784 -1.59678413
5 25.02 27.913653 -2.89365294
6 16.86 15.746432 1.11356772
7 14.38 12.450260 1.92974001
8 9.60 8.403777 1.19622309
9 24.35 28.214999 -3.86499936
10 27.50 27.976292 -0.47629200
11 17.08 18.402328 -1.32232830
12 37.00 37.461882 -0.46188206
13 41.95 41.458933 0.49106715
14 11.66 12.262343 -0.60234282
15 21.65 15.809071 5.84092866
16 17.89 18.251995 -0.36199456
17 69.00 64.665871 4.33412887
18 10.30 12.336831 -2.03683074
19 34.93 36.471506 -1.54150602
20 46.59 46.559789 0.03021107
21 44.88 47.060901 -2.18090138
22 54.12 52.561290 1.55871047
23 56.63 56.307784 0.32221591
24 22.13 19.982190 2.14780957
25 21.15 20.996264 0.15373580
```

```
## Covariance matrix and standard errors
cov.mat <- vcov(fit)
cov.mat
```

```
              (Intercept)          length          height
(Intercept)  1.123740429 -3.921612e-02 -1.781991e-03
length       -0.039216122  8.746709e-03 -9.903775e-05
height       -0.001781991 -9.903775e-05  7.831149e-06
```

```
data.frame(std.error = sqrt(diag(cov.mat)))
```

```
              std.error
(Intercept) 1.060066238
length      0.093523844
height      0.002798419
```

4.2 RSS-based Inference: F-test, and adjusted R^2

The General Linear Model

The general linear model is:

$$y = X\beta + \epsilon$$

- y : $n \times 1$ vector of responses
- X : $n \times p$ design matrix (first column often ones)
- β : $p \times 1$ parameter vector, where $p = k + 1$
- ϵ : $n \times 1$ error vector

4.2.1 RSS-Based Quantities

4.2.1.1 RSS-Based Quantities

Source	Sum of Squares	R^2	df	Mean Squares	F	SS_{adj}	$\hat{\sigma}^2$	R^2_{adj}
$x^\top \beta$	$SSR = \sum_{i=1}^n (\hat{y}_i - \bar{y})^2$	$\frac{SSR}{SST}$	k	$\frac{MSR}{SSR} = \frac{MSR}{MSE}$	$\frac{MSR}{MSE}$	SSR_{adj}	$\hat{\sigma}_{x^\top \beta}^2 = \frac{SSR_{\text{adj}}}{n-1}$	$\frac{SSR_{\text{adj}}}{SST} = 1 - \frac{MSE}{MST}$
ϵ	$SSE = \sum_{i=1}^n (y_i - \hat{y}_i)^2$	—	$n - p$	$\frac{MSE}{SSE} = \frac{SSE}{n-p}$	—	SSE	$\hat{\sigma}_\epsilon^2 = \frac{SSE}{n-p}$	—
y	$SST = \sum_{i=1}^n (y_i - \bar{y})^2$	—	$n - 1$	$\frac{MST}{SST} = \frac{SST}{n-1}$	—	SST	$\hat{\sigma}_y^2 = \frac{SST}{n-1}$	—

Interpretation of the $\hat{\sigma}^2$ Column

The $\hat{\sigma}^2$ column highlights how each sum of squares corresponds to an estimated variance. This view makes the adjusted coefficient of determination clear:

$$R^2_{\text{adj}} = 1 - \frac{\hat{\sigma}_\epsilon^2}{\hat{\sigma}_y^2} = \frac{\hat{\sigma}_{x^\top \beta}^2}{\hat{\sigma}_y^2}.$$

Hence, the adjusted R^2 simply expresses the **proportion of total estimated variance** attributable to the fitted model $X\beta$ rather than the residual noise ϵ .

4.2.2 Remarks

4.2.2.1 Fundamental Identities

$$\begin{aligned} \text{SST} &= \text{SSR} + \text{SSE}, \\ \text{MST} &= \text{MSE} + \frac{\text{SSR}_{\text{adj}}}{n-1}. \end{aligned}$$

where

$$\text{SSR}_{\text{adj}} = (n-1)\text{MST} - (n-p+k)\text{MSE} = \text{SST} - \text{SSE} - k\text{MSE} = \text{SSR} - k\text{MSE}.$$

4.2.2.2 Difference of $\hat{\sigma}^2$ and Mean Squares

The quantity $\hat{\sigma}^2$ represents the **estimated variance** associated with each component of the model. MSE and MST are the estimated variances of the ϵ and y itself. However, the MSR, although called **Mean Square for Regression (MSR)** is *NOT* an estimate of the variance or sample variance of $x^\top \beta$. The name of “mean” here is used to indicate a different thing. Its name “Mean Square” reflects that it is also an estimate estimate of noise variance σ^2 under $H_0: \beta = 0$:

$$E[\text{MSR} \mid H_0] = \sigma^2, \quad E[\text{MSR} \mid H_1] > \sigma^2.$$

Hence the F-statistic

$$F = \frac{\text{MSR}}{\text{MSE}}$$

is approximately equal to 1 subject to the variability as characterized with F-distribution with degree freedoms of k and $n-p$. This test is to test whether any regression coefficients are not equal to 0.

$$4.2.2.3 \quad \hat{\sigma}_{x^\top \beta}^2 = \frac{\text{SSR}_{\text{adj}}}{n-1}$$

$\hat{\sigma}_{x^\top \beta}^2$ is an unbiased estimator of the variance of linear signal when x is regarded as a random variable. This can be seen from the following equations:

$$E[\text{SSR}] = k\sigma^2 + \beta^\top X^\top (I - J/n) X \beta, \quad E[\text{MSE}] = \sigma^2.$$

Hence,

$$\begin{aligned} E[\text{SSR}_{\text{adj}}] &= E[\text{SSR}] - k E[\text{MSE}] \\ &= \beta^\top X^\top (I - J/n) X \beta \\ &= \sum_{i=1}^n (\mu_i - \bar{\mu})^2, \end{aligned}$$

where

$$\begin{aligned} \mu_i &= x_i^\top \beta \\ \bar{\mu} &= \frac{1}{n} \sum_{i=1}^n \mu_i \end{aligned}$$

For fixed X , $\text{SSR}_{\text{adj}}/(n-1)$ equals the **sample variance** of the true means $\{\mu_i\}$ over the observed design points. If the rows of X are independently sampled with covariance matrix Σ_X (the random- X model), then

$$\mathbb{E}_X \left[\frac{\text{SSR}_{\text{adj}}}{n-1} \right] = \beta^\top \Sigma_X \beta = \text{Var}(x^\top \beta),$$

4.2.2.4 Connection to Rao-Blackwell Formula

The decomposition of $\hat{\sigma}^2$ is consistent with the **Rao-Blackwell formula** for total variance:

$$\text{Var}(y) = \text{Var}(E[y | x]) + E(\text{Var}[y | x]).$$

Here,

- $\text{Var}(E[y | x])$ corresponds to the **explained variation** due to the regression component $x^\top \beta$, and
- $E(\text{Var}[y | x])$ corresponds to the **residual variation** due to ϵ .

4.2.3 A Simulation Study to Understand the Distributions of RSS

Data Generating Model

For $n = 30$ and $p_{\max} = 20$, simulate with either $H_0 : \beta = \mathbf{0}$ or H_1 where only $\beta_1 \neq 0$; $\epsilon_i \sim N(0, 1)$.

Sequence of Fitted Models

4 Multiple Linear Regression

Model Name	# of Predictors (k)	# of Parameters (p)	R Formula
Model 0	0	1	$y \sim 1$
Model 1	2	3	$y \sim x_1 + x_2$
...
Final Model	20	21	$y \sim x_1 + \dots + x_{20}$

4.2.3.1 When H_0 is true

[rss-h0.mp4](#)

4.2.3.2 When H_1 is true

[rss-h1.mp4](#)

4.2.4 Example: Modelling Children Weight with Height and Age

```
## Data: Weight, height and age of children
wgt <- c(64, 71, 53, 67, 55, 58, 77, 57, 56, 51, 76, 68)
hgt <- c(57, 59, 49, 62, 51, 50, 55, 48, 42, 42, 61, 57)
age <- c(8, 10, 6, 11, 8, 7, 10, 9, 10, 6, 12, 9)
child.data <- data.frame(wgt, hgt, age)
```

4.2.4.1 Problem 1: Height then Age

```
fit_hgt_age <- lm(wgt ~ hgt + age, data = child.data)
summary(fit_hgt_age)
```

Call:

```
lm(formula = wgt ~ hgt + age, data = child.data)
```

Residuals:

Min	1Q	Median	3Q	Max
-6.8708	-1.7004	0.3454	1.4642	10.2336

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	6.5530	10.9448	0.599	0.5641

4.2 RSS-based Inference: F-test, and adjusted R^2

```
hgt          0.7220      0.2608   2.768   0.0218 *
age          2.0501      0.9372   2.187   0.0565 .
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
Residual standard error: 4.66 on 9 degrees of freedom
Multiple R-squared:  0.78, Adjusted R-squared:  0.7311
F-statistic: 15.95 on 2 and 9 DF,  p-value: 0.001099
```

```
fit_hgt <- lm(wgt ~ hgt, data = child.data)
summary(fit_hgt)
```

```
Call:
lm(formula = wgt ~ hgt, data = child.data)
```

```
Residuals:
    Min       1Q   Median       3Q      Max
-5.8736 -3.8973 -0.4402  2.2624 11.8375
```

```
Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept)   6.1898     12.8487   0.482  0.64035
hgt           1.0722      0.2417   4.436  0.00126 **
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
Residual standard error: 5.471 on 10 degrees of freedom
Multiple R-squared:  0.663, Adjusted R-squared:  0.6293
F-statistic: 19.67 on 1 and 10 DF,  p-value: 0.001263
```

```
anova(fit_hgt, fit_hgt_age)
```

Analysis of Variance Table

```
Model 1: wgt ~ hgt
Model 2: wgt ~ hgt + age
   Res.Df  RSS Df Sum of Sq    F Pr(>F)
1     10 299.33
2      9 195.43  1    103.9 4.7849 0.05649 .
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

4 Multiple Linear Regression

```
anova(fit_hgt_age)
```

Analysis of Variance Table

Response: wgt

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
hgt	1	588.92	588.92	27.1216	0.0005582 ***
age	1	103.90	103.90	4.7849	0.0564853 .
Residuals	9	195.43	21.71		

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

4.2.4.2 Problem 2: Age then Height

```
fit_age <- lm(wgt ~ age, data = child.data)
summary(fit_age)
```

Call:

```
lm(formula = wgt ~ age, data = child.data)
```

Residuals:

Min	1Q	Median	3Q	Max
-11.000	-3.911	1.143	4.071	10.000

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	30.5714	8.6137	3.549	0.00528 **
age	3.6429	0.9551	3.814	0.00341 **

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 6.015 on 10 degrees of freedom

Multiple R-squared: 0.5926, Adjusted R-squared: 0.5519

F-statistic: 14.55 on 1 and 10 DF, p-value: 0.003407

```
fit_age_hgt <- lm(wgt ~ age + hgt, data = child.data)
summary(fit_age_hgt)
```

Call:

4.2 RSS-based Inference: F-test, and adjusted R^2

```
lm(formula = wgt ~ age + hgt, data = child.data)
```

Residuals:

Min	1Q	Median	3Q	Max
-6.8708	-1.7004	0.3454	1.4642	10.2336

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	6.5530	10.9448	0.599	0.5641
age	2.0501	0.9372	2.187	0.0565 .
hgt	0.7220	0.2608	2.768	0.0218 *

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 4.66 on 9 degrees of freedom

Multiple R-squared: 0.78, Adjusted R-squared: 0.7311

F-statistic: 15.95 on 2 and 9 DF, p-value: 0.001099

```
anova(fit_age, fit_age_hgt)
```

Analysis of Variance Table

Model 1: wgt ~ age

Model 2: wgt ~ age + hgt

	Res.Df	RSS	Df	Sum of Sq	F	Pr(>F)
1	10	361.86				
2	9	195.43	1	166.43	7.6646	0.02181 *

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```
anova(fit_age_hgt)
```

Analysis of Variance Table

Response: wgt

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
age	1	526.39	526.39	24.2419	0.0008205 ***
hgt	1	166.43	166.43	7.6646	0.0218070 *
Residuals	9	195.43	21.71		

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

4.2.5 Example: Wire bond strength

```
fit_len_hgt <- lm(strength ~ length + height, data = bond.data)
fit_hgt_len <- lm(strength ~ height+length, data = bond.data)
anova(fit_len_hgt)
```

Analysis of Variance Table

Response: strength

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
length	1	5885.9	5885.9	1124.293	< 2.2e-16 ***
height	1	104.9	104.9	20.041	0.0001883 ***
Residuals	22	115.2	5.2		

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```
anova(fit_hgt_len)
```

Analysis of Variance Table

Response: strength

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
height	1	1483.2	1483.2	283.32	4.731e-14 ***
length	1	4507.5	4507.5	861.01	< 2.2e-16 ***
Residuals	22	115.2	5.2		

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```
summary(fit_hgt_len)
```

Call:

```
lm(formula = strength ~ height + length, data = bond.data)
```

Residuals:

Min	1Q	Median	3Q	Max
-3.865	-1.542	-0.362	1.196	5.841

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	2.263791	1.060066	2.136	0.044099 *

```

height      0.012528   0.002798   4.477 0.000188 ***
length      2.744270   0.093524  29.343 < 2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 2.288 on 22 degrees of freedom
Multiple R-squared:  0.9811,    Adjusted R-squared:  0.9794
F-statistic: 572.2 on 2 and 22 DF,  p-value: < 2.2e-16

```

```
summary(fit_len_hgt)
```

```

Call:
lm(formula = strength ~ length + height, data = bond.data)

Residuals:
    Min       1Q   Median       3Q      Max
-3.865 -1.542 -0.362  1.196  5.841

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept)  2.263791   1.060066   2.136 0.044099 *
length       2.744270   0.093524  29.343 < 2e-16 ***
height       0.012528   0.002798   4.477 0.000188 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 2.288 on 22 degrees of freedom
Multiple R-squared:  0.9811,    Adjusted R-squared:  0.9794
F-statistic: 572.2 on 2 and 22 DF,  p-value: < 2.2e-16

```

4.2.6 Relationship between t-test and partial F-test

- A t-test for a single coefficient is a special case of the partial F-test; the relationship is $F = t^2$ for 1 df in the numerator.
- The p-value from t-test (output of `summary(lm())`) is the same as anova test for: $H_0 : \beta_j = 0$ vs H_1 : all covaraites have non-zero effects.

4.3 Predictions for Mean Response and a Future Observation

4.3.1 Confidence Interval for Mean Response

```
predict(fit, newdata = data.frame(length = 8, height = 275),
        interval = "confidence", level = 0.95)
```

```
      fit      lwr      upr
1 27.6631 26.66324 28.66296
```

4.3.2 Prediction Interval for a New Observation

```
predict(fit, newdata = data.frame(length = 8, height = 275),
        interval = "prediction", level = 0.95)
```

```
      fit      lwr      upr
1 27.6631 22.81378 32.51241
```

4.4 Model Diagnostics

4.4.1 Residual Calculations

```
residuals_df <- data.frame(
  hat_values = hatvalues(fit),
  ordinary_resid = resid(fit),
  standardized_resid = resid(fit) / sigma(fit),
  studentized_internal = rstandard(fit),
  studentized_external = rstudent(fit)
)
residuals_df
```

```
      hat_values ordinary_resid standardized_resid studentized_internal
1  0.15728923    1.57127871      0.68673363      0.74808172
2  0.11164598   -1.14600783     -0.50086730     -0.53140990
3  0.14191905   -2.20409488     -0.96330846     -1.03992315
4  0.10188923   -1.59678413     -0.69788088     -0.73640435
5  0.04178381   -2.89365294     -1.26468257     -1.29196212
```

6	0.07486842	1.11356772	0.48668921	0.50599936
7	0.11806106	1.92974001	0.84340057	0.89807919
8	0.15608149	1.19622309	0.52281407	0.56911105
9	0.12797685	-3.86499936	-1.68921340	-1.80892479
10	0.04131672	-0.47629200	-0.20816532	-0.21260369
11	0.09253979	-1.32232830	-0.57792886	-0.60668127
12	0.05256700	-0.46188206	-0.20186740	-0.20739197
13	0.08202675	0.49106715	0.21462286	0.22400668
14	0.11291577	-0.60234282	-0.26325633	-0.27950939
15	0.07373697	5.84092866	2.55280118	2.65246601
16	0.08794942	-0.36199456	-0.15821117	-0.16566382
17	0.25934228	4.33412887	1.89424832	2.20104100
18	0.29287870	-2.03683074	-0.89020500	-1.05862725
19	0.09617553	-1.54150602	-0.67372136	-0.70866056
20	0.14726101	0.03021107	0.01320387	0.01429859
21	0.12963943	-2.18090138	-0.95317165	-1.02169558
22	0.13580052	1.55871047	0.68124063	0.73281364
23	0.18237610	0.32221591	0.14082575	0.15574183
24	0.10908869	2.14780957	0.93870874	0.99452024
25	0.07287021	0.15373580	0.06719084	0.06978142

studentized_external

1	0.74035927
2	-0.52255660
3	-1.04194550
4	-0.72850799
5	-1.31305171
6	0.49726770
7	0.89397096
8	0.56016499
9	-1.91552083
10	-0.20792931
11	-0.59775404
12	-0.20282206
13	0.21910643
14	-0.27356920
15	3.14216850
16	-0.16195600
17	2.43521394
18	-1.06168251
19	-0.70040768
20	0.01396991
21	-1.02276424
22	0.72486668
23	0.15224503
24	0.99426154


4 Multiple Linear Regression

25 0.06818458

4.4.2 Residual Plots

```
n <- nrow(bond.data)
r <- rstudent(fit)
y.hat <- fitted.values(fit)

par(mfrow = c(2, 3))
qqnorm(r, main = "Normal Q-Q Plot"); qqline(r)
plot(y.hat, r, xlab = "Fitted values", ylab = "Studentized Residuals"); abline(h = 0)
plot(1:n, r, xlab = "Observation Number", ylab = "Studentized Residuals"); abline(h = 0)
plot(bond.data$length, r, xlab = "Wire Length", ylab = "Studentized Residuals"); abline(h = 0)
plot(bond.data$height, r, xlab = "Die Height", ylab = "Studentized Residuals"); abline(h = 0)
```



4.5 Influential Observations

```
influence_df <- data.frame(dffits = dffits(fit),
                           cook.D = cooks.distance(fit),
                           dfbetas(fit))
influence_df
```

	dffits	cook.D	X.Intercept.	length	height
1	0.319854702	3.481748e-02	0.3179493921	-0.100534181	-0.200085326
2	-0.185251548	1.183028e-02	-0.1403477437	-0.051464370	0.148315370
3	-0.423741713	5.962023e-02	-0.2219151046	-0.237135616	0.339340552


4 Multiple Linear Regression

4	-0.245376811	2.050736e-02	0.0787635526	0.022343842	-0.184260891
5	-0.274191565	2.426179e-02	-0.1572410603	-0.009662357	0.055328303
6	0.141461363	6.906752e-03	0.1301249135	-0.058073567	-0.049408295
7	0.327082639	3.598953e-02	0.1479853099	-0.261848970	0.142220906
8	0.240902557	1.996750e-02	0.2394962591	-0.076575387	-0.149787102
9	-0.733818383	1.600749e-01	-0.5011686139	-0.283749099	0.605559181
10	-0.043165927	6.493384e-04	-0.0241138520	-0.001038287	0.007460760
11	-0.190885522	1.251125e-02	-0.0602003847	0.132053762	-0.102733745
12	-0.047774650	7.954763e-04	0.0017554145	-0.016926531	-0.008495572
13	0.065496465	1.494604e-03	-0.0224255162	0.017340091	0.033756253
14	-0.097602753	3.314830e-03	-0.0483552961	0.077880727	-0.038066705
15	0.886553487	1.866936e-01	0.8097463528	-0.374290156	-0.292048214
16	-0.050292599	8.821617e-04	-0.0177647675	0.034746758	-0.025275682
17	1.441003392	5.654455e-01	-0.8513738015	1.008880052	0.413618783
18	-0.683268805	1.547244e-01	-0.0218935465	0.521608456	-0.532432956
19	-0.228476293	1.781295e-02	0.0700729860	0.018004228	-0.167581999
20	0.005805362	1.176892e-05	0.0005613509	0.004752581	-0.003094588
21	-0.394724862	5.182743e-02	-0.0084618169	-0.324109965	0.170622396
22	0.287343813	2.812893e-02	-0.1326208213	0.183002776	0.076391058
23	0.071903545	1.803448e-03	-0.0412553376	0.040164093	0.030365108
24	0.347915085	4.036930e-02	0.3084561584	0.016541769	-0.262089402
25	0.019115733	1.275757e-04	0.0062730782	-0.011614674	0.009459029


4.5.1 Plotting with the olsrr Package

```
## install.packages("olsrr") # Run once if needed
library(olsrr)

ols_plot_cooksd_chart(fit)
```

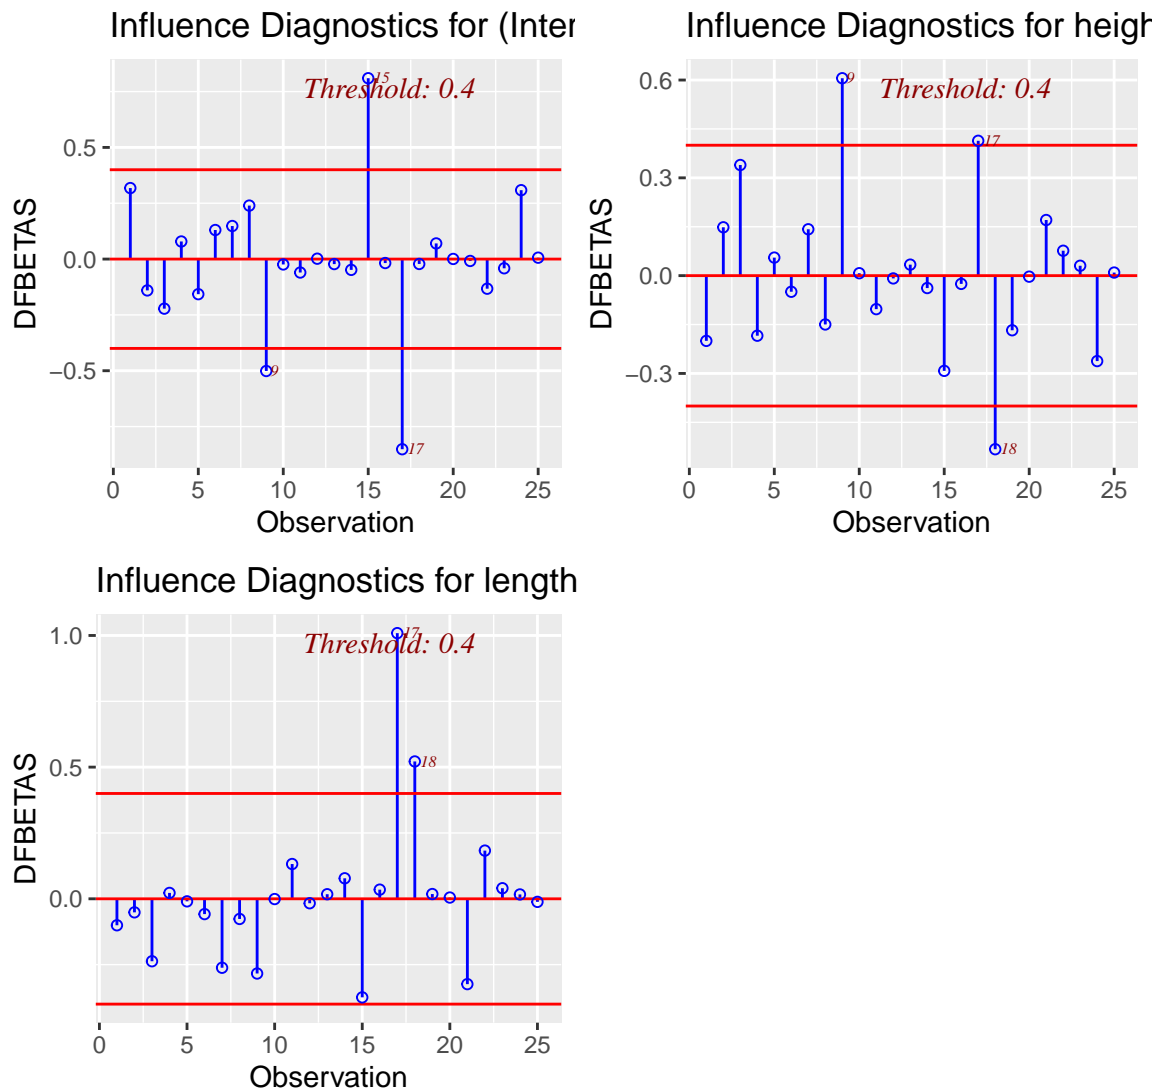



```
ols_plot_dffits(fit)
```



```
ols_plot_dfbetas(fit)
```

page 1 of 1



4.6 Polynomial Regression

```
y <- c(1.81, 1.70, 1.65, 1.55, 1.48, 1.40, 1.30, 1.26, 1.24, 1.21, 1.20, 1.18)
x <- c(20, 25, 30, 35, 40, 50, 60, 65, 70, 75, 80, 90)
fit_poly <- lm(y ~ x + I(x^2))
summary(fit_poly)
```

Call:

```
lm(formula = y ~ x + I(x^2))
```

4 Multiple Linear Regression

Residuals:

	Min	1Q	Median	3Q	Max
	-0.0174763	-0.0065087	0.0001297	0.0071482	0.0151887

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	2.198e+00	2.255e-02	97.48	6.38e-15	***
x	-2.252e-02	9.424e-04	-23.90	1.88e-09	***
I(x^2)	1.251e-04	8.658e-06	14.45	1.56e-07	***


Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.01219 on 9 degrees of freedom

Multiple R-squared: 0.9975, Adjusted R-squared: 0.9969

F-statistic: 1767 on 2 and 9 DF, p-value: 2.096e-12

```
plot(x, y, xlab = "Lot size, x", ylab = "Average cost per unit, y")
lines(x, predict(fit_poly, newdata = data.frame(x = x)), type = "l")
```



```
fit1 <- lm(y ~ x)
anova(fit1, fit_poly)
```

Analysis of Variance Table

Model 1: y ~ x

Model 2: y ~ x + I(x^2)

	Res.Df	RSS	Df	Sum of Sq	F	Pr(>F)
1	10	0.032340				
2	9	0.001337	1	0.031002	208.67	1.564e-07 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

4.7 Handling Categorical Variables with Dummy Variables

Investigate the common observation that males tend to have higher blood pressure than females of similar age.

4 Multiple Linear Regression

```
## Note: Update this path to your local file location  
sbpdata <- read.csv("sbpdata.csv")  
sbpdata
```

	sex	sbp	age
1	0	144	39
2	0	138	45
3	0	145	47
4	0	162	65
5	0	142	46
6	0	170	67
7	0	124	42
8	0	158	67
9	0	154	56
10	0	162	64
11	0	150	56
12	0	140	59
13	0	110	34
14	0	128	42
15	0	130	48
16	0	135	45
17	0	114	17
18	0	116	20
19	0	124	19
20	0	136	36
21	0	142	50
22	0	120	39
23	0	120	21
24	0	160	44
25	0	158	53
26	0	144	63
27	0	130	29
28	0	125	25
29	0	175	69
30	1	158	41
31	1	185	60
32	1	152	41
33	1	159	47
34	1	176	66
35	1	156	47
36	1	184	68
37	1	138	43
38	1	172	68
39	1	168	57

40	1	176	65
41	1	164	57
42	1	154	61
43	1	124	36
44	1	142	44
45	1	144	50
46	1	149	47
47	1	128	19
48	1	130	22
49	1	138	21
50	1	150	38
51	1	156	52
52	1	134	41
53	1	134	18
54	1	174	51
55	1	174	55
56	1	158	65
57	1	144	33
58	1	139	23
59	1	180	70
60	1	165	56
61	1	172	62
62	1	160	51
63	1	157	48
64	1	170	59
65	1	153	40
66	1	148	35
67	1	140	33
68	1	132	26
69	1	169	61

4.7.1 Four Models Involving “sex”

4.7.1.1 Coincidence Model (Age Only)

```
## Ensure sex is a factor (labels will appear in the legend)
sbpdata$sex <- as.factor(sbpdata$sex)

## Fit (you already have this)
fit.age <- lm(sbp ~ age, data = sbpdata)

## Generate predictions over the observed age range
new_age <- seq(min(sbpdata$age, na.rm = TRUE),
```

4 Multiple Linear Regression

```
        max(sbpdata$age, na.rm = TRUE),
        length.out = 200)
pred <- predict(fit.age, newdata = data.frame(age = new_age))


## Simple palette for the sex levels (works for 1-3 levels; expand if needed)
lev <- levels(sbpdata$sex)
cols <- setNames(c("steelblue3", "tomato3", "darkorchid3")[seq_along(lev)], lev)

## Scatter plot with colored points by sex
plot(sbp ~ age, data = sbpdata,
     col = cols[sbpdata$sex], pch = 16,
     xlab = "Age", ylab = "Systolic BP")

## Add predicted line
lines(new_age, pred, lwd = 2)

## Legend
legend("topleft", legend = lev, col = cols[lev], pch = 16, bty = "n", title = "Sex")
```


4.7 Handling Categorical Variables with Dummy Variables



```
data.frame(model.matrix(fit.age))
```

	X.Intercept.	age
1	1	39
2	1	45
3	1	47
4	1	65
5	1	46
6	1	67
7	1	42
8	1	67
9	1	56
10	1	64
11	1	56
12	1	59
13	1	34
14	1	42

4 Multiple Linear Regression

15	1	48
16	1	45
17	1	17
18	1	20
19	1	19
20	1	36
21	1	50
22	1	39
23	1	21
24	1	44
25	1	53
26	1	63
27	1	29
28	1	25
29	1	69
30	1	41
31	1	60
32	1	41
33	1	47
34	1	66
35	1	47
36	1	68
37	1	43
38	1	68
39	1	57
40	1	65
41	1	57
42	1	61
43	1	36
44	1	44
45	1	50
46	1	47
47	1	19
48	1	22
49	1	21
50	1	38
51	1	52
52	1	41
53	1	18
54	1	51
55	1	55
56	1	65
57	1	33
58	1	23
59	1	70

60	1	56
61	1	62
62	1	51
63	1	48
64	1	59
65	1	40
66	1	35
67	1	33
68	1	26
69	1	61

```
print(anova(fit.age))
```

Analysis of Variance Table

Response: sbp

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
age	1	14951.3	14951.3	121.27	< 2.2e-16 ***
Residuals	67	8260.5	123.3		

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

4.7.1.2 Additive Effect Model (Age + Sex)

```
## Parallelism: H0: beta3=0 (Sex has additive effect)
fit.agePLUSsex <- lm(sbp ~ age + sex, data = sbpdata)

## Ensure sex is a factor for labeling/colors
sbpdata$sex <- factor(sbpdata$sex)


## Fit (additive: parallelism)
fit.agePLUSsex <- lm(sbp ~ age + sex, data = sbpdata)

## X-range and palette
ages <- seq(min(sbpdata$age, na.rm = TRUE),
            max(sbpdata$age, na.rm = TRUE),
            length.out = 200)
lev <- levels(sbpdata$sex)
cols <- setNames(c("steelblue3", "tomato3", "darkorchid3")[seq_along(lev)], lev)

## Scatter with colored points by sex
plot(sbp ~ age, data = sbpdata,
```

4 Multiple Linear Regression

```
col = cols[sbpdata$sex], pch = 16,  
xlab = "Age", ylab = "Systolic BP")  
  
## Parallel fitted lines: one per sex (same slope, different intercepts)  
for (sx in lev) {  
  nd <- data.frame(age = ages, sex = factor(sx, levels = lev))  
  yhat <- predict(fit.agePLUSsex, newdata = nd)  
  lines(ages, yhat, col = cols[sx], lwd = 2)  
}  
  
## Legend  
legend("topleft", legend = lev, col = cols[lev], pch = 16, lwd = 2, bty = "n", title = "Sex")
```



```
data.frame(model.matrix(fit.agePLUSsex))
```

X.Intercept. age sex1

4.7 Handling Categorical Variables with Dummy Variables

1	1	39	0
2	1	45	0
3	1	47	0
4	1	65	0
5	1	46	0
6	1	67	0
7	1	42	0
8	1	67	0
9	1	56	0
10	1	64	0
11	1	56	0
12	1	59	0
13	1	34	0
14	1	42	0
15	1	48	0
16	1	45	0
17	1	17	0
18	1	20	0
19	1	19	0
20	1	36	0
21	1	50	0
22	1	39	0
23	1	21	0
24	1	44	0
25	1	53	0
26	1	63	0
27	1	29	0
28	1	25	0
29	1	69	0
30	1	41	1
31	1	60	1
32	1	41	1
33	1	47	1
34	1	66	1
35	1	47	1
36	1	68	1
37	1	43	1
38	1	68	1
39	1	57	1
40	1	65	1
41	1	57	1
42	1	61	1
43	1	36	1
44	1	44	1
45	1	50	1

4 Multiple Linear Regression

46	1	47	1
47	1	19	1
48	1	22	1
49	1	21	1
50	1	38	1
51	1	52	1
52	1	41	1
53	1	18	1
54	1	51	1
55	1	55	1
56	1	65	1
57	1	33	1
58	1	23	1
59	1	70	1
60	1	56	1
61	1	62	1
62	1	51	1
63	1	48	1
64	1	59	1
65	1	40	1
66	1	35	1
67	1	33	1
68	1	26	1
69	1	61	1

```
print(anova(fit.age, fit.agePLUSsex))
```

Analysis of Variance Table

Model 1: sbp ~ age

Model 2: sbp ~ age + sex

	Res.Df	RSS	Df	Sum of Sq	F	Pr(>F)
1	67	8260.5				
2	66	5202.0	1	3058.5	38.805	3.701e-08 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

4.7.1.3 Varying Intercept and Varying Slope Model (Age + Sex + Age:Sex)

```
## Make sure sex is a factor (for colors/legend)
sbpdata$sex <- factor(sbpdata$sex)
```

```

## Fit (interaction: different slopes by sex)
fit.age.TIMES.sex <- lm(sbp ~ age + sex + age:sex, data = sbpdata)

## Age grid and palette
ages <- seq(min(sbpdata$age, na.rm = TRUE),
            max(sbpdata$age, na.rm = TRUE),
            length.out = 200)
lev <- levels(sbpdata$sex)
cols <- setNames(c("steelblue3", "tomato3", "darkorchid3")[seq_along(lev)], lev)


## Scatter: color points by sex
plot(sbp ~ age, data = sbpdata,
     col = cols[sbpdata$sex], pch = 16,
     xlab = "Age", ylab = "Systolic BP")

## Fitted lines: one per sex (different slopes allowed)
for (sx in lev) {
  nd <- data.frame(age = ages, sex = factor(sx, levels = lev))
  yhat <- predict(fit.age.TIMES.sex, newdata = nd)
  lines(ages, yhat, col = cols[sx], lwd = 2)
}

## Legend
legend("topleft", legend = lev, col = cols[lev], pch = 16, lwd = 2, bty = "n", title = "Sex")

```

4 Multiple Linear Regression



Model Matrix and ANOVA

```
data.frame(model.matrix(fit.age.TIMES.sex))
```

	X.Intercept.	age	sex1	age.sex1
1	1	39	0	0
2	1	45	0	0
3	1	47	0	0
4	1	65	0	0
5	1	46	0	0
6	1	67	0	0
7	1	42	0	0
8	1	67	0	0
9	1	56	0	0
10	1	64	0	0
11	1	56	0	0
12	1	59	0	0

4.7 Handling Categorical Variables with Dummy Variables

13	1	34	0	0
14	1	42	0	0
15	1	48	0	0
16	1	45	0	0
17	1	17	0	0
18	1	20	0	0
19	1	19	0	0
20	1	36	0	0
21	1	50	0	0
22	1	39	0	0
23	1	21	0	0
24	1	44	0	0
25	1	53	0	0
26	1	63	0	0
27	1	29	0	0
28	1	25	0	0
29	1	69	0	0
30	1	41	1	41
31	1	60	1	60
32	1	41	1	41
33	1	47	1	47
34	1	66	1	66
35	1	47	1	47
36	1	68	1	68
37	1	43	1	43
38	1	68	1	68
39	1	57	1	57
40	1	65	1	65
41	1	57	1	57
42	1	61	1	61
43	1	36	1	36
44	1	44	1	44
45	1	50	1	50
46	1	47	1	47
47	1	19	1	19
48	1	22	1	22
49	1	21	1	21
50	1	38	1	38
51	1	52	1	52
52	1	41	1	41
53	1	18	1	18
54	1	51	1	51
55	1	55	1	55
56	1	65	1	65
57	1	33	1	33

4 Multiple Linear Regression

58	1	23	1	23
59	1	70	1	70
60	1	56	1	56
61	1	62	1	62
62	1	51	1	51
63	1	48	1	48
64	1	59	1	59
65	1	40	1	40
66	1	35	1	35
67	1	33	1	33
68	1	26	1	26
69	1	61	1	61

```
summary(fit.age.TIMES.sex)
```

Call:

```
lm(formula = sbp ~ age + sex + age:sex, data = sbpdata)
```

Residuals:

Min	1Q	Median	3Q	Max
-20.647	-3.410	1.254	4.314	21.153

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	97.07708	5.17046	18.775	< 2e-16 ***
age	0.94932	0.10864	8.738	1.43e-12 ***
sex1	12.96144	7.01172	1.849	0.0691 .
age:sex1	0.01203	0.14519	0.083	0.9342

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 8.946 on 65 degrees of freedom

Multiple R-squared: 0.7759, Adjusted R-squared: 0.7656

F-statistic: 75.02 on 3 and 65 DF, p-value: < 2.2e-16

```
print(anova(fit.age,fit.agePLUSsex,fit.age.TIMES.sex))
```

Analysis of Variance Table

Model 1: sbp ~ age

Model 2: sbp ~ age + sex

Model 3: sbp ~ age + sex + age:sex

```

      Res.Df    RSS Df Sum of Sq      F    Pr(>F)
1         67 8260.5
2         66 5202.0  1   3058.52 38.2210 4.692e-08 ***
3         65 5201.4  1      0.55  0.0069   0.9342
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```

4.7.1.4 Varying Slope, Equal Intercept Model (Age + Age:Sex)

```

## Make sure sex is a factor (for colors/legend)
sbpdata$sex <- factor(sbpdata$sex)

## Fit (interaction: different slopes by sex)
fit.equal.intercept <- lm(sbp ~ age + age:sex, data = sbpdata)


## Age grid and palette
ages <- seq(min(sbpdata$age, na.rm = TRUE),
            max(sbpdata$age, na.rm = TRUE),
            length.out = 200)
lev <- levels(sbpdata$sex)
cols <- setNames(c("steelblue3", "tomato3", "darkorchid3")[seq_along(lev)], lev)

## Scatter: color points by sex
plot(sbp ~ age, data = sbpdata,
     col = cols[sbpdata$sex], pch = 16,
     xlab = "Age", ylab = "Systolic BP")

## Fitted lines: one per sex (different slopes allowed)
for (sx in lev) {
  nd <- data.frame(age = ages, sex = factor(sx, levels = lev))
  yhat <- predict(fit.equal.intercept, newdata = nd)
  lines(ages, yhat, col = cols[sx], lwd = 2)
}

## Legend
legend("topleft", legend = lev, col = cols[lev], pch = 16, lwd = 2, bty = "n", title = "Sex")

```



4.7.2 Orders of Terms Matters in ANOVA and Warnings in Interpreting t-test Tables

```
fit.int <- lm(sbp ~ 1, data = sbpdata)
fit.sex <- lm(sbp ~ sex, data = sbpdata)

print(anova(fit.int, fit.age, fit.agePLUSsex, fit.age.TIMES.sex))
```

Analysis of Variance Table

```
Model 1: sbp ~ 1
Model 2: sbp ~ age
Model 3: sbp ~ age + sex
Model 4: sbp ~ age + sex + age:sex
```

	Res.Df	RSS	Df	Sum of Sq	F	Pr(>F)
1	68	23211.8				

4.7 Handling Categorical Variables with Dummy Variables

```
2      67  8260.5  1    14951.3 186.8390 < 2.2e-16 ***
3      66  5202.0  1     3058.5  38.2210 4.692e-08 ***
4      65  5201.4  1         0.5   0.0069   0.9342
---
```

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```
print(anova(fit.int,fit.age,fit.equal.intercept, fit.age.TIMES.sex))
```

Analysis of Variance Table

Model 1: sbp ~ 1

Model 2: sbp ~ age

Model 3: sbp ~ age + age:sex

Model 4: sbp ~ age + sex + age:sex

	Res.Df	RSS	Df	Sum of Sq	F	Pr(>F)
1	68	23211.8				
2	67	8260.5	1	14951.3	186.8390	< 2.2e-16 ***
3	66	5474.9	1	2785.6	34.8107	1.437e-07 ***
4	65	5201.4	1	273.4	3.4171	0.06907 .

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```
print(anova(fit.int,fit.sex,fit.agePLUSsex, fit.age.TIMES.sex))
```

Analysis of Variance Table

Model 1: sbp ~ 1

Model 2: sbp ~ sex

Model 3: sbp ~ age + sex

Model 4: sbp ~ age + sex + age:sex

	Res.Df	RSS	Df	Sum of Sq	F	Pr(>F)
1	68	23211.8				
2	67	19282.5	1	3929.2	49.1017	1.684e-09 ***
3	66	5202.0	1	14080.6	175.9583	< 2.2e-16 ***
4	65	5201.4	1	0.5	0.0069	0.9342

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```
summary(fit.age)
```

Call:

4 Multiple Linear Regression

```
lm(formula = sbp ~ age, data = sbpdata)
```

Residuals:

Min	1Q	Median	3Q	Max
-26.782	-7.632	1.968	8.201	22.651

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	103.34905	4.33190	23.86	<2e-16 ***
age	0.98333	0.08929	11.01	<2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 11.1 on 67 degrees of freedom

Multiple R-squared: 0.6441, Adjusted R-squared: 0.6388

F-statistic: 121.3 on 1 and 67 DF, p-value: < 2.2e-16

```
summary(fit.equal.intercept)
```

Call:

```
lm(formula = sbp ~ age + age:sex, data = sbpdata)
```

Residuals:

Min	1Q	Median	3Q	Max
-21.6338	-4.3067	0.9922	4.9819	20.2753

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	104.12501	3.55578	29.283	< 2e-16 ***
age	0.80908	0.07918	10.219	3.14e-15 ***
age:sex1	0.26705	0.04608	5.795	2.09e-07 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 9.108 on 66 degrees of freedom

Multiple R-squared: 0.7641, Adjusted R-squared: 0.757

F-statistic: 106.9 on 2 and 66 DF, p-value: < 2.2e-16

```
summary(fit.agePLUSsex)
```

Call:

4.7 Handling Categorical Variables with Dummy Variables

```
lm(formula = sbp ~ age + sex, data = sbpdata)
```

Residuals:

Min	1Q	Median	3Q	Max
-20.705	-3.299	1.248	4.325	21.160

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	96.77353	3.62085	26.727	< 2e-16 ***
age	0.95606	0.07153	13.366	< 2e-16 ***
sex1	13.51345	2.16932	6.229	3.7e-08 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 8.878 on 66 degrees of freedom

Multiple R-squared: 0.7759, Adjusted R-squared: 0.7691

F-statistic: 114.2 on 2 and 66 DF, p-value: < 2.2e-16

```
summary(fit.age.TIMES.sex)
```

Call:

```
lm(formula = sbp ~ age + sex + age:sex, data = sbpdata)
```

Residuals:

Min	1Q	Median	3Q	Max
-20.647	-3.410	1.254	4.314	21.153

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	97.07708	5.17046	18.775	< 2e-16 ***
age	0.94932	0.10864	8.738	1.43e-12 ***
sex1	12.96144	7.01172	1.849	0.0691 .
age:sex1	0.01203	0.14519	0.083	0.9342

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 8.946 on 65 degrees of freedom

Multiple R-squared: 0.7759, Adjusted R-squared: 0.7656

F-statistic: 75.02 on 3 and 65 DF, p-value: < 2.2e-16

4.8 Model Building

```
library(olsrr)
## Note: Update this path to your local file location
wine <- read.csv("wine.csv")

model.wine <- lm(quality ~ ., data = wine)
```

4.8.1 All Possible Regression

```
ols_step_best_subset(model.wine)
```

Best Subsets Regression	
Model Index	Predictors
1	flavor
2	flavor oakiness
3	aroma flavor oakiness
4	clarity aroma flavor oakiness
5	clarity aroma body flavor oakiness

Subsets Regression Summary								
Model	R-Square	Adj. R-Square	Pred R-Square	C(p)	AIC	SBIC	SBC	MSE
1	0.6242	0.6137	0.5868	9.0436	130.0214	21.6859	134.9341	61.4
2	0.6611	0.6417	0.6058	6.8132	128.0901	20.1242	134.6404	57.0
3	0.7038	0.6776	0.6379	3.9278	124.9781	18.0702	133.1661	51.3
4	0.7147	0.6801	0.6102	4.6747	125.5480	19.2854	135.3736	50.9
5	0.7206	0.6769	0.587	6.0000	126.7552	21.0956	138.2183	51.5

AIC: Akaike Information Criteria

SBIC: Sawa's Bayesian Information Criteria

SBC: Schwarz Bayesian Criteria

MSEP: Estimated error of prediction, assuming multivariate normality

FPE: Final Prediction Error

HSP: Hocking's Sp

APC: Amemiya Prediction Criteria

4.8.2 Automated Stepwise Procedures

```
## Backward Elimination (alpha_out = 0.1)
ols_step_backward_p(model.wine, p_val = 0.1)
```

Stepwise Summary

Step	Variable	AIC	SBC	SBIC	R2	Adj. R2
0	Full Model	126.755	138.218	21.096	0.72060	0.67694
1	body	125.548	135.374	19.285	0.71471	0.68013
2	clarity	124.978	133.166	18.070	0.70377	0.67763

Final Model Output

Model Summary

R	0.839	RMSE	1.098
R-Squared	0.704	MSE	1.207
Adj. R-Squared	0.678	Coef. Var	9.338
Pred R-Squared	0.638	AIC	124.978
MAE	0.868	SBC	133.166

RMSE: Root Mean Square Error

MSE: Mean Square Error

MAE: Mean Absolute Error

AIC: Akaike Information Criteria

SBC: Schwarz Bayesian Criteria

ANOVA

	Sum of Squares	DF	Mean Square	F	Sig.
Regression	108.935	3	36.312	26.925	0.0000
Residual	45.853	34	1.349		
Total	154.788	37			

Parameter Estimates

4 Multiple Linear Regression

model	Beta	Std. Error	Std. Beta	t	Sig	lower	upper
(Intercept)	6.467	1.333		4.852	0.000	3.759	9.176
aroma	0.580	0.262	0.307	2.213	0.034	0.047	1.113
flavor	1.200	0.275	0.603	4.364	0.000	0.641	1.758
oakiness	-0.602	0.264	-0.217	-2.278	0.029	-1.140	-0.065

```
## Forward Selection (alpha_in = 0.1)
ols_step_forward_p(model.wine, p_val = 0.1)
```

Stepwise Summary

Step	Variable	AIC	SBC	SBIC	R2	Adj. R2
0	Base Model	165.209	168.484	55.141	0.00000	0.00000
1	flavor	130.021	134.934	21.686	0.62417	0.61373
2	oakiness	128.090	134.640	20.124	0.66111	0.64175
3	aroma	124.978	133.166	18.070	0.70377	0.67763

Final Model Output

Model Summary

R	0.839	RMSE	1.098
R-Squared	0.704	MSE	1.207
Adj. R-Squared	0.678	Coef. Var	9.338
Pred R-Squared	0.638	AIC	124.978
MAE	0.868	SBC	133.166

RMSE: Root Mean Square Error

MSE: Mean Square Error

MAE: Mean Absolute Error

AIC: Akaike Information Criteria

SBC: Schwarz Bayesian Criteria

ANOVA

Sum of Squares	DF	Mean Square	F	Sig.
-------------------	----	-------------	---	------

Regression	108.935	3	36.312	26.925	0.0000
Residual	45.853	34	1.349		
Total	154.788	37			

Parameter Estimates

model	Beta	Std. Error	Std. Beta	t	Sig	lower	upper
(Intercept)	6.467	1.333		4.852	0.000	3.759	9.176
flavor	1.200	0.275	0.603	4.364	0.000	0.641	1.758
oakiness	-0.602	0.264	-0.217	-2.278	0.029	-1.140	-0.065
aroma	0.580	0.262	0.307	2.213	0.034	0.047	1.113

```
## Stepwise Regression (alpha_in = 0.1, alpha_out = 0.1)
ols_step_both_p(model.wine, p_enter = 0.1, p_remove = 0.1)
```

Stepwise Summary

Step	Variable	AIC	SBC	SBIC	R2	Adj. R2
0	Base Model	165.209	168.484	55.141	0.00000	0.00000
1	flavor (+)	130.021	134.934	21.686	0.62417	0.61373
2	oakiness (+)	128.090	134.640	20.124	0.66111	0.64175
3	aroma (+)	124.978	133.166	18.070	0.70377	0.67763

Final Model Output

Model Summary

R	0.839	RMSE	1.098
R-Squared	0.704	MSE	1.207
Adj. R-Squared	0.678	Coef. Var	9.338
Pred R-Squared	0.638	AIC	124.978
MAE	0.868	SBC	133.166

RMSE: Root Mean Square Error

MSE: Mean Square Error

MAE: Mean Absolute Error

4 Multiple Linear Regression

AIC: Akaike Information Criteria

SBC: Schwarz Bayesian Criteria

ANOVA

	Sum of Squares	DF	Mean Square	F	Sig.
Regression	108.935	3	36.312	26.925	0.0000
Residual	45.853	34	1.349		
Total	154.788	37			

Parameter Estimates

model	Beta	Std. Error	Std. Beta	t	Sig	lower	upper
(Intercept)	6.467	1.333		4.852	0.000	3.759	9.176
flavor	1.200	0.275	0.603	4.364	0.000	0.641	1.758
oakiness	-0.602	0.264	-0.217	-2.278	0.029	-1.140	-0.065
aroma	0.580	0.262	0.307	2.213	0.034	0.047	1.113

4.9 Multicollinearity

4.9.1 A Simple Example

```
y <- c(19, 20, 37, 39, 36, 38)
x1 <- c(4, 4, 7, 7, 7.1, 7.1)
x2 <- c(16, 16, 49, 49, 50.4, 50.4)
cor(data.frame(x1, x2))
```

```
      x1      x2
x1 1.0000000 0.9999713
x2 0.9999713 1.0000000
```

```
fit_multi <- lm(y ~ x1 + x2)
summary(fit_multi)
```

Call:

```
lm(formula = y ~ x1 + x2)
```

Residuals:

```
    1    2    3    4    5    6
-0.5  0.5 -1.0  1.0 -1.0  1.0
```

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-156.056	117.158	-1.332	0.275
x1	65.444	45.890	1.426	0.249
x2	-5.389	4.152	-1.298	0.285

Residual standard error: 1.225 on 3 degrees of freedom

Multiple R-squared: 0.9897, Adjusted R-squared: 0.9829

F-statistic: 144.3 on 2 and 3 DF, p-value: 0.001043

```
fit1_multi <- lm(y ~ x1)
summary(fit1_multi)
```

Call:

```
lm(formula = y ~ x1)
```

Residuals:

```
    1    2    3    4    5    6
-0.5260  0.4740 -0.1925  1.8075 -1.7814  0.2186
```

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-4.0293	2.3332	-1.727	0.159
x1	5.8888	0.3762	15.654	9.73e-05 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.325 on 4 degrees of freedom

Multiple R-squared: 0.9839, Adjusted R-squared: 0.9799

F-statistic: 245.1 on 1 and 4 DF, p-value: 9.725e-05

```
ols_vif_tol(fit_multi)
```

	Variables	Tolerance	VIF
1	x1	5.738191e-05	17427.09
2	x2	5.738191e-05	17427.09

4.9.2 VIFs in the Wine Quality Data

```
wine.x <- wine[, -ncol(wine)] # Assuming quality is the last column  
cor(wine.x)
```

	clarity	aroma	body	flavor	oakiness
clarity	1.00000000	0.0619021	-0.3083783	-0.08515993	0.1832147
aroma	0.06190210	1.0000000	0.5489102	0.73656121	0.2016444
body	-0.30837826	0.5489102	1.0000000	0.64665917	0.1521059
flavor	-0.08515993	0.7365612	0.6466592	1.00000000	0.1797605
oakiness	0.18321471	0.2016444	0.1521059	0.17976051	1.0000000

```
## VIF using olsrr (data frame output)  
ols_vif_tol(model.wine)
```

	Variables	Tolerance	VIF
1	clarity	0.7896462	1.266390
2	aroma	0.4199665	2.381143
3	body	0.4862649	2.056492
4	flavor	0.3728175	2.682277
5	oakiness	0.9118005	1.096731

4.9.3 VIFs in the Children Height Data

```
## Data: Weight, height and age of children  
wgt <- c(64, 71, 53, 67, 55, 58, 77, 57, 56, 51, 76, 68)  
hgt <- c(57, 59, 49, 62, 51, 50, 55, 48, 42, 42, 61, 57)  
age <- c(8, 10, 6, 11, 8, 7, 10, 9, 10, 6, 12, 9)  
  
fit_age_hgt <- lm(wgt ~ hgt + age, data = child.data)  
ols_vif_tol(fit_age_hgt)
```

	Variables	Tolerance	VIF
1	hgt	0.6232021	1.604616
2	age	0.6232021	1.604616

5 Logistic Regression

5.1 Odds as a Function of Probability

For an event with probability p , the odds is

$$\text{odds}(p) = \frac{p}{1-p}$$

and the log-odds (logit) is

$$\text{logit}(p) = \log\left(\frac{p}{1-p}\right)$$

```
## Plot odds(p) with a right-hand axis for log(odds(p)),
## using different line colors for the two curves.
## Defaults: p in [0.01, 0.99].
## Args:
##   p_min, p_max : endpoints for p-grid (0<p_min<p_max<1)
##   n           : number of grid points
##   annotate     : add reference lines/labels if TRUE
##   odds_col     : color for odds(p)
##   logit_col    : color for log(odds(p))
##   lwd1, lwd2   : line widths for the two curves

plot_odds <- function(p_min = 0.01, p_max = 0.99, n = 400,
                      annotate = TRUE,
                      odds_col = "steelblue",
                      logit_col = "firebrick",
                      lwd1 = 2, lwd2 = 2) {
  stopifnot(p_min > 0, p_max < 1, p_min < p_max, n >= 10)
  p <- seq(p_min, p_max, length.out = n)
  odds <- p / (1 - p)
  logit <- log(odds)

  ## Left y-axis: odds(p)
  plot(p, odds, type = "l", lwd = lwd1, col = odds_col,
       xlab = "Probability p",
```

5 Logistic Regression

```
      ylab = "odds(p) = p / (1 - p)")
if (annotate) {
  abline(h = 1, v = 0.5, lty = 2)
  text(0.52, 1.05, "p = 0.5 → odds = 1", adj = 0)
}


## Right y-axis: logit(p) = log(odds)
op <- par(new = TRUE)
on.exit(par(op), add = TRUE)
plot(p, logit, type = "l", lwd = lwd2, col = logit_col,
      axes = FALSE, xlab = "", ylab = "")
axis(4)
mtext("log{odds(p)} = log{p/(1 - p)}", side = 4, line = 3)

if (annotate) {
  abline(v = 0.5, lty = 2)
  # logit(0.5) = 0 reference (horizontal) on the right-axis scale
  usr <- par("usr")
  segments(x0 = usr[1], y0 = 0, x1 = 0.5, y1 = 0, lty = 3)
}

legend("topleft",
      legend = c("odds(p)", "log{odds(p)}"),
      col = c(odds_col, logit_col),
      lwd = c(lwd1, lwd2), bty = "n")

invisible(list(p = p, odds = odds, logit = logit))
}

## Example usage:
## plot_odds() # defaults: steelblue for odds, firebrick for log-odds (right axis)
plot_odds(odds_col = "#1f77b4", logit_col = "#d62728", n = 600)
```

Logistic regression models **log-odds** linearly in predictors, which both keeps fitted probabilities in $(0, 1)$ and turns multiplicative effects on odds into **additive** effects on the linear predictor.

5.2 A Simulated Data

We simulate data from a logistic model where the **logit** is a linear function of x :

$$\text{logit } p(x) = \log \left(\frac{p(x)}{1 - p(x)} \right) = \beta_0 + \beta_1 x,$$

so that

$$p(x) = \text{logit}^{-1}(\beta_0 + \beta_1 x) = \frac{1}{1 + \exp -(\beta_0 + \beta_1 x)}.$$

We then display the observed y_i (binary outcomes) and the true probability curve $p(x)$ in red.

5 Logistic Regression

```
set.seed(123)

## -- Truth (edit as desired) --
n      <- 200
beta0 <- 0
beta1 <- 4

## -- Simulate --
x      <- runif(n, -1, 1)          # predictor
eta    <- beta0 + beta1 * x
p      <- plogis(eta)             # true p(x)
y      <- rbinom(n, size = 1, prob = p) # outcomes

sim.data <- data.frame(x = x, y = y, p = p)
```

5.2.1 Fit a logistic model to the simulated data

```
## -- Optional: fit a model to the simulated data --
sim.fit <- glm(y ~ x, data = sim.data, family = binomial())
p_fit  <- predict(sim.fit, newdata = data.frame(x = x), type = "response")

## -- Plot: points for y_i (jittered), red line for true p(x) --
## Define jitter amount
jit <- 0.05
## jitter to separate 0/1 visually
yj <- jitter(sim.data$y, amount = jit)

plot(sim.data$x, yj,
     pch = 16, col = rgb(0, 0, 0, 0.45),
     xlab = "x",
     ylab = "Observed y (points) & p(x) (curves)",
     ylim = c(-0.1, 1.1))

## True probability curve (red)
xg <- seq(min(x), max(x), length.out = 500)
lines(xg, plogis(beta0 + beta1 * xg), col = "red", lwd = 2)


## Optional: add fitted probability curve (dashed dark red)
lines(xg, predict(sim.fit, newdata = data.frame(x = xg), type = "response"),
     col = "darkred", lwd = 2, lty = 2)

legend("topleft",
```

```

legend = c("y (jittered points)", "true p(x)", "fitted p(x)"),
pch     = c(16, NA, NA),
lty     = c(NA, 1, 2),
col     = c(rgb(0,0,0,0.45), "red", "darkred"),
lwd     = c(NA, 2, 2),
bty     = "n")

```



5.3 Example of Coronary Heart Disease Data

5.3.1 Load a dataset

This dataset is about a follow-up study to determine the development of coronary heart disease (CHD) over 9 years of follow-up of 609 white males from Evans County, Georgia.

Variable meanings (as provided):

- chd: 1 if a person has the disease, 0 otherwise.
- smk: 1 if smoker, 0 if not.
- cat: 1 if catecholamine level is high, 0 if low.
- sbp: systolic blood pressure (continuous).
- age: age in years (continuous).
- chl: cholesterol level (continuous).

5 Logistic Regression

- ecg: 1 if electrocardiogram is abnormal, 0 if normal.
- hpt: 1 if high blood pressure, 0 if normal.

```
## Adjust the path if needed. The default is your original V: drive path.
```

```
data_path <- "evans.dat"
```

```
## Read data (expects a header row)
```

```
CHD.data <- read.table(data_path, header = TRUE)
```

```
CHD.data
```

	id	chd	age	cat	chl	dbp	ecg	sbp	smk	hpt
1	21	0	56	0	270	80	0	138	0	0
2	31	0	43	0	159	74	0	128	1	0
3	51	1	56	1	201	112	1	164	1	1
4	71	0	64	1	179	100	0	200	1	1
5	74	0	49	0	243	82	0	145	1	0
6	91	0	46	0	252	88	0	142	1	0
7	111	1	52	0	179	80	1	128	1	0
8	131	0	63	0	217	92	0	135	0	0
9	141	0	42	0	176	76	0	114	1	0
10	191	0	55	0	250	114	1	182	0	1
11	201	0	74	0	293	100	0	166	0	1
12	241	0	53	0	179	90	0	158	0	0
13	251	0	58	0	201	86	0	142	1	0
14	261	0	56	0	206	85	0	120	1	0
15	271	0	69	0	225	84	0	168	0	1
16	283	1	51	1	259	102	1	135	0	1
17	291	0	43	0	193	78	0	118	1	0
18	311	0	64	1	185	100	1	180	0	1
19	312	0	44	0	150	108	0	160	0	1
20	331	0	42	0	211	86	1	122	0	0
21	351	0	57	0	216	88	0	130	0	0
22	381	1	64	1	247	75	1	130	0	0
23	401	0	49	0	200	82	0	130	0	0
24	411	0	68	1	205	74	0	152	1	0
25	431	0	41	0	225	98	0	135	1	1
26	441	0	64	0	263	98	0	162	1	1
27	451	0	41	0	205	80	0	120	0	0
28	481	0	59	0	253	98	0	154	0	1
29	501	0	50	0	282	90	0	142	1	0
30	521	0	56	0	230	80	0	118	0	0
31	541	0	57	1	203	112	0	182	0	1
32	561	0	42	0	211	86	0	144	0	0
33	571	0	59	0	234	84	0	164	1	1

5.3 Example of Coronary Heart Disease Data

34	581	0	44	0	202	94	1	174	1	1
35	611	0	52	0	162	78	0	134	1	0
36	621	0	45	0	191	85	0	135	0	0
37	641	0	41	0	220	110	0	178	0	1
38	651	0	59	0	240	80	0	130	0	0
39	671	0	52	0	189	110	0	168	0	1
40	681	0	64	0	247	102	0	170	0	1
41	731	0	46	0	181	122	1	176	1	1
42	741	0	42	0	168	75	0	104	1	0
43	751	0	54	0	187	86	0	146	1	0
44	761	0	48	0	196	98	0	130	0	1
45	811	0	45	0	155	70	0	142	1	0
46	851	0	66	1	173	100	0	160	1	1
47	861	0	41	0	138	70	0	115	1	0
48	871	0	76	0	269	94	0	175	1	1
49	881	1	49	0	266	102	0	152	1	1
50	921	0	57	1	200	100	0	160	1	1
51	941	0	51	0	188	84	0	124	1	0
52	961	1	43	0	218	108	1	136	1	1
53	971	0	43	0	212	80	1	108	1	0
54	981	0	45	0	212	102	0	150	1	1
55	991	0	45	0	180	80	0	122	1	0
56	1061	1	46	1	166	76	1	162	0	1
57	1071	0	40	0	257	84	0	130	0	0
58	1081	0	48	0	243	82	1	154	1	0
59	1091	0	64	1	179	100	1	148	1	1
60	1111	0	70	0	167	64	0	112	1	0
61	1151	0	52	0	178	84	1	112	1	0
62	1171	0	55	0	178	94	0	152	0	0
63	1181	0	49	0	211	68	0	114	1	0
64	1191	1	56	0	171	85	0	125	1	0
65	1201	1	66	1	205	80	0	150	1	0
66	1221	0	48	0	229	130	0	195	1	1
67	1231	0	47	0	238	120	1	160	1	1
68	1471	0	54	1	195	112	0	174	1	1
69	1501	0	44	0	162	82	0	120	0	0
70	1561	0	51	0	240	84	1	126	1	0
71	1691	0	43	0	177	102	1	138	1	1
72	1701	0	68	0	252	88	1	112	1	0
73	1741	0	49	0	217	105	0	148	0	1
74	1751	0	55	0	263	84	0	114	0	0
75	1761	0	51	0	229	100	0	162	1	1
76	1791	0	50	0	245	96	0	144	0	1
77	1811	0	65	0	177	74	0	122	0	0
78	1821	0	42	0	203	78	0	134	1	0

5 Logistic Regression

79	1851	0	57	0	194	75	1	114	0	0
80	1881	0	42	0	288	110	0	142	0	1
81	1891	0	53	0	217	70	0	120	1	0
82	1901	0	57	1	163	94	0	184	0	1
83	1911	0	61	0	180	84	0	136	0	0
84	1951	0	53	0	209	98	0	142	1	1
85	1961	0	45	0	200	80	0	135	0	0
86	1971	0	44	0	194	80	0	120	1	0
87	2241	0	63	0	227	90	1	135	0	0
88	2252	0	42	0	158	92	0	135	1	0
89	2273	0	73	1	183	120	1	220	0	1
90	2281	0	47	0	253	110	0	140	1	1
91	2311	0	56	0	198	88	0	122	1	0
92	2371	1	41	0	228	132	0	162	1	1
93	2381	0	58	0	217	86	0	140	0	0
94	2391	0	55	0	163	70	0	110	1	0
95	2401	0	46	0	212	124	0	184	1	1
96	2461	0	57	0	144	95	0	130	0	1
97	2481	0	44	0	134	74	0	114	1	0
98	2501	0	52	1	183	96	0	158	1	1
99	2511	0	56	0	212	108	0	144	0	1
100	2531	0	64	0	214	82	0	128	1	0
101	2541	0	54	0	249	92	0	120	1	0
102	2571	0	52	0	180	78	1	104	1	0
103	2591	0	42	0	212	92	0	125	1	0
104	2611	0	46	0	167	82	0	120	1	0
105	2621	0	46	0	273	94	0	152	0	0
106	2631	0	42	0	210	96	0	134	1	1
107	2641	0	54	1	173	110	0	170	1	1
108	2671	0	43	0	256	72	0	114	1	0
109	2681	0	53	0	234	80	0	122	0	0
110	2691	1	40	0	221	100	0	140	1	1
111	2711	0	46	0	261	86	0	128	1	0
112	2731	0	43	0	299	80	0	116	0	0
113	2851	0	43	0	192	75	0	115	1	0
114	2861	0	47	0	185	80	1	146	1	0
115	2871	0	44	0	283	70	0	108	1	0
116	2881	0	49	0	176	92	0	134	1	0
117	2891	1	56	1	331	110	0	190	1	1
118	2901	1	56	0	203	82	0	120	1	0
119	2911	0	64	1	217	92	0	166	1	1
120	2921	0	54	0	164	72	0	122	1	0
121	2931	0	54	0	256	98	0	148	0	1
122	2991	0	51	0	184	98	0	170	0	1
123	3001	0	49	0	165	80	0	114	1	0

5.3 Example of Coronary Heart Disease Data

124	3011	0	47	0	189	92	0	145	0	0
125	3031	0	58	0	221	88	0	140	1	0
126	3061	0	70	1	126	66	1	164	1	1
127	3601	0	42	0	169	80	1	122	1	0
128	3611	0	59	0	266	92	0	138	0	0
129	3621	0	57	1	153	92	0	148	1	0
130	3651	0	76	1	211	114	1	228	1	1
131	3661	0	43	0	113	76	0	114	1	0
132	3701	0	46	0	200	85	0	145	1	0
133	3721	0	75	1	172	114	1	162	1	1
134	3751	0	42	0	131	84	0	130	0	0
135	3761	0	64	0	214	84	0	120	0	0
136	3771	0	63	1	236	94	1	190	0	1
137	3791	0	54	0	213	90	0	142	0	0
138	3811	0	66	0	226	90	0	166	0	1
139	3813	0	44	0	200	110	0	160	1	1
140	3841	0	72	0	188	78	0	130	0	0
141	3861	0	50	0	268	102	0	138	0	1
142	3871	0	59	1	195	114	1	208	0	1
143	3881	1	59	0	216	95	0	140	1	1
144	3891	0	53	0	182	92	0	130	1	0
145	3901	0	48	0	178	95	0	135	1	1
146	3911	0	40	0	191	76	0	152	1	0
147	3941	0	61	0	255	80	0	120	0	0
148	3951	0	42	0	225	80	0	126	1	0
149	4161	0	42	0	166	90	0	145	0	0
150	4191	0	49	0	278	84	0	126	1	0
151	4202	0	40	0	235	72	0	116	0	0
152	4221	0	51	0	251	86	0	128	1	0
153	4242	0	44	0	217	90	0	146	0	0
154	4261	0	44	0	181	94	0	144	1	0
155	4271	0	47	0	208	108	0	178	0	1
156	4291	0	51	0	182	112	0	182	0	1
157	4301	0	69	0	228	75	0	115	1	0
158	4321	0	58	1	170	88	1	152	1	0
159	4331	0	74	1	147	80	1	200	0	1
160	4341	0	48	0	190	78	0	114	1	0
161	4381	0	64	0	205	98	1	140	0	1
162	4401	0	53	0	216	78	0	124	1	0
163	4411	0	71	0	170	90	0	140	1	0
164	4421	0	47	0	127	74	0	110	1	0
165	4451	0	56	0	235	92	0	128	1	0
166	4461	0	40	0	200	72	0	118	0	0
167	4491	0	46	0	283	100	0	148	1	1
168	4531	0	68	1	157	94	0	162	0	1

5 Logistic Regression

169	4551	1	54	0	206	76	1	142	0	0
170	4581	0	54	0	197	88	0	125	1	0
171	4591	0	45	0	163	75	0	115	1	0
172	4601	0	66	0	176	60	1	124	0	0
173	4641	0	58	0	211	88	0	146	1	0
174	4681	0	49	0	161	75	0	115	0	0
175	4711	0	51	0	244	90	0	128	0	0
176	4731	0	44	0	172	100	0	138	0	1
177	4751	0	61	1	166	86	0	156	1	0
178	4771	0	48	0	184	76	0	116	1	0
179	4781	0	63	0	143	92	0	122	1	0
180	4791	0	54	0	196	84	0	138	1	0
181	4801	0	52	0	189	88	0	142	1	0
182	4811	0	45	0	227	98	1	140	1	1
183	4821	0	62	0	236	94	0	160	0	1
184	4831	0	41	0	240	86	0	144	0	0
185	4851	0	41	0	256	90	0	145	1	0
186	4861	0	61	0	200	84	0	148	1	0
187	4871	0	42	0	199	104	0	166	1	1
188	4901	0	42	0	161	88	0	124	0	0
189	4911	0	72	0	211	80	1	104	0	0
190	4951	0	43	0	180	64	0	92	0	0
191	4961	1	72	0	200	86	1	138	0	0
192	4971	0	51	0	206	80	0	132	1	0
193	4981	0	58	0	254	94	0	152	1	0
194	5011	0	41	0	215	90	0	142	1	0
195	5061	0	71	1	162	98	1	184	1	1
196	5071	1	63	0	145	96	0	162	1	1
197	5091	0	44	0	220	90	1	130	1	0
198	5101	0	45	0	298	108	0	170	1	1
199	5111	0	54	0	300	94	0	148	1	0
200	5131	1	52	1	306	108	0	178	1	1
201	5141	0	55	0	302	134	1	206	1	1
202	5181	1	41	0	158	80	0	140	1	0
203	5191	0	54	0	194	130	1	170	1	1
204	5211	0	64	1	229	94	1	156	1	0
205	5251	0	61	0	259	82	0	118	0	0
206	5281	0	40	0	214	94	0	130	0	0
207	5301	0	51	0	168	106	0	156	1	1
208	5361	0	51	0	265	90	0	158	1	0
209	5391	0	75	0	225	80	0	125	0	0
210	5421	1	40	0	219	80	0	115	1	0
211	5451	1	63	0	202	110	0	160	0	1
212	5461	0	42	1	217	94	1	138	0	0
213	5471	1	64	0	231	85	0	120	1	0

5.3 Example of Coronary Heart Disease Data

214	5521	1	50	0	215	114	0	170	1	1
215	5601	0	49	0	146	98	1	145	1	1
216	5621	0	48	0	198	75	0	120	1	0
217	5631	0	58	0	206	92	0	154	0	0
218	5641	0	46	0	227	98	0	168	1	1
219	5671	0	46	0	214	92	1	166	1	1
220	6341	0	42	0	225	100	1	162	1	1
221	6351	0	57	0	193	86	0	124	0	0
222	6371	0	50	0	186	102	0	160	0	1
223	6391	0	46	0	147	85	0	122	1	0
224	6411	0	45	0	205	100	0	166	0	1
225	6421	0	57	1	196	98	1	196	1	1
226	6441	0	46	0	195	96	0	138	0	1
227	6451	0	45	1	153	108	1	212	1	1
228	6461	0	58	1	172	96	1	168	0	1
229	6482	0	42	0	293	110	0	176	1	1
230	6491	0	53	0	274	106	0	158	1	1
231	6501	0	55	0	221	106	0	162	0	1
232	6511	0	53	0	197	70	0	112	0	0
233	6531	0	69	1	194	100	1	150	0	1
234	6551	0	58	0	204	74	0	122	1	0
235	6561	0	46	0	203	84	0	114	1	0
236	6591	0	62	0	293	90	1	142	1	0
237	6631	0	61	0	197	72	1	110	0	0
238	6641	0	49	0	195	82	0	138	1	0
239	6651	0	48	0	184	96	0	144	1	1
240	6661	1	55	0	209	85	0	130	1	0
241	6681	0	52	1	209	98	1	170	0	1
242	6691	0	61	0	214	100	0	158	0	1
243	6721	0	68	1	130	106	1	200	0	1
244	6731	0	55	0	196	70	0	125	0	0
245	6741	0	52	1	237	126	1	224	0	1
246	6751	0	43	0	185	85	1	140	1	0
247	6761	1	47	0	248	104	1	132	1	1
248	6781	0	57	0	252	106	0	166	0	1
249	6791	0	55	0	198	96	0	144	1	1
250	6801	0	71	0	176	62	0	138	0	0
251	6811	0	74	1	193	98	1	202	0	1
252	6821	1	65	0	185	105	0	156	0	1
253	6831	0	65	0	241	102	1	146	0	1
254	6871	0	44	0	231	70	0	108	0	0
255	6881	0	40	0	157	78	1	122	0	0
256	6891	0	45	0	152	106	1	148	1	1
257	6911	0	50	0	237	102	0	156	1	1
258	6921	0	64	1	175	110	1	142	0	1

5 Logistic Regression

259	6931	1	56	0	195	94	1	150	0	0
260	6941	0	62	1	151	88	0	165	0	1
261	6961	0	44	0	205	80	0	128	1	0
262	6981	0	73	0	190	75	0	115	0	0
263	7001	0	46	0	239	100	0	160	1	1
264	7021	0	51	0	232	80	0	120	0	0
265	7031	0	59	1	170	100	0	180	1	1
266	7051	1	67	1	319	104	0	182	0	1
267	7091	0	54	0	225	86	0	122	0	0
268	7101	0	49	0	252	90	0	128	1	0
269	7121	0	46	0	224	84	0	130	1	0
270	7131	0	42	1	229	90	1	145	0	0
271	8641	0	68	0	195	76	1	116	1	0
272	8651	0	43	0	230	85	1	135	1	0
273	8671	0	56	1	186	98	1	154	0	1
274	8682	0	68	1	192	94	0	154	1	0
275	8711	0	46	0	184	78	0	110	1	0
276	8721	1	64	1	233	94	0	140	1	0
277	8731	0	54	0	175	96	0	156	1	1
278	8751	0	48	0	188	106	0	148	1	1
279	8771	0	41	0	232	82	0	126	1	0
280	8811	0	65	1	178	106	1	194	0	1
281	8841	0	41	0	187	108	0	154	0	1
282	8851	1	42	0	207	86	1	128	1	0
283	8971	0	66	0	94	86	0	134	0	0
284	8981	0	44	0	211	90	0	142	1	0
285	9011	0	42	0	275	100	1	150	1	1
286	9021	0	51	0	165	85	0	130	1	0
287	9031	0	56	0	282	94	0	134	1	0
288	9051	1	64	1	239	94	0	162	1	1
289	9061	0	44	0	256	106	0	162	1	1
290	9071	1	55	0	175	108	0	160	1	1
291	9091	0	55	0	306	82	0	160	1	1
292	9101	1	67	0	188	102	1	168	0	1
293	9191	1	56	1	221	78	1	154	1	0
294	9201	1	63	1	213	156	1	256	0	1
295	9261	1	67	0	250	100	0	158	0	1
296	9471	0	48	0	268	120	0	172	1	1
297	9601	1	45	0	263	86	0	132	0	0
298	9631	0	49	0	150	98	1	120	1	1
299	9651	1	70	1	251	108	1	174	1	1
300	9671	0	45	0	180	102	0	156	1	1
301	9681	0	48	0	336	110	0	174	1	1
302	9711	1	42	0	210	70	0	124	1	0
303	9721	0	69	1	179	110	0	175	1	1

5.3 Example of Coronary Heart Disease Data

304	9731	0	44	0	177	75	0	120	0	0
305	9751	0	48	0	227	92	0	158	1	0
306	9791	0	46	0	195	72	0	120	1	0
307	9801	0	52	0	227	76	0	116	1	0
308	9811	0	73	0	250	84	0	154	0	0
309	9831	0	67	0	218	96	1	148	0	1
310	9841	0	63	0	229	100	0	168	1	1
311	9871	0	45	1	197	80	1	134	0	0
312	9881	0	46	0	190	86	0	122	1	0
313	9891	0	68	1	189	104	1	202	1	1
314	9901	0	49	0	185	80	0	120	1	0
315	9911	1	63	0	194	90	0	190	1	1
316	9931	0	59	0	192	66	0	134	0	0
317	9941	0	67	1	261	80	1	160	1	1
318	9951	0	49	0	174	78	1	108	0	0
319	9961	0	65	1	189	114	1	168	1	1
320	9981	0	44	0	248	100	0	145	1	1
321	10011	0	45	0	214	94	0	122	0	0
322	10041	0	47	0	275	76	0	114	1	0
323	10051	0	46	0	259	92	0	130	1	0
324	10071	0	52	0	230	68	0	100	0	0
325	10091	0	60	0	206	84	0	138	1	0
326	10121	0	45	0	275	95	0	125	1	1
327	10151	1	67	1	237	100	1	170	1	1
328	10181	0	60	0	289	80	1	118	0	0
329	10201	0	65	1	176	82	0	200	1	1
330	10221	0	72	1	232	80	1	210	1	1
331	10231	1	71	0	184	90	0	160	1	1
332	10241	0	55	0	283	108	1	178	1	1
333	10271	0	54	0	214	110	0	170	1	1
334	10401	0	52	1	161	76	0	162	1	1
335	10402	0	48	0	232	98	0	154	1	1
336	10921	0	66	0	228	72	0	120	1	0
337	10951	0	52	1	206	120	1	206	0	1
338	10971	0	64	0	218	80	0	110	1	0
339	11011	0	42	0	262	92	0	142	1	0
340	11081	0	52	0	227	66	0	98	0	0
341	11101	0	51	0	215	60	0	100	0	0
342	11141	0	54	0	146	70	0	115	0	0
343	11151	0	51	0	268	85	0	140	1	0
344	11161	0	60	0	211	94	0	166	0	1
345	11221	0	48	0	213	90	1	145	1	0
346	11281	0	73	0	249	108	0	206	1	1
347	11291	0	50	0	218	92	0	130	1	0
348	11321	0	45	0	221	92	0	128	0	0

5 Logistic Regression

349	11341	1	56	1	228	92	0	152	1	0
350	11351	1	46	0	240	104	0	142	1	1
351	11361	1	76	1	279	96	0	136	1	1
352	11391	0	52	0	186	70	0	118	0	0
353	11441	0	54	0	160	110	1	200	1	1
354	11461	0	53	1	222	104	1	154	0	1
355	11481	0	43	0	211	65	0	112	1	0
356	11491	0	46	0	195	132	1	230	1	1
357	11501	0	63	0	290	90	0	150	0	0
358	11511	0	44	0	220	95	0	138	0	1
359	11531	0	42	0	161	80	0	124	1	0
360	11553	0	74	1	212	98	0	164	1	1
361	11611	0	53	0	182	86	0	136	1	0
362	11651	0	56	1	223	110	1	208	1	1
363	11661	0	47	0	290	92	0	136	1	0
364	11711	0	43	0	249	90	1	162	1	1
365	11721	0	51	0	174	92	0	124	1	0
366	11731	0	63	1	204	92	1	190	1	1
367	11781	0	49	0	245	62	0	124	1	0
368	11791	0	57	1	216	114	0	174	1	1
369	11811	0	43	0	245	120	1	145	0	1
370	11831	0	58	0	151	98	0	138	1	1
371	11851	0	49	1	178	102	0	166	1	1
372	11891	0	47	0	227	88	0	132	1	0
373	11911	0	45	0	253	104	0	152	1	1
374	11941	1	65	1	222	88	1	162	0	1
375	11971	0	51	0	258	94	1	178	1	1
376	11981	0	49	0	182	84	1	124	1	0
377	11991	0	51	0	184	96	0	150	1	1
378	12051	1	67	0	357	90	0	129	0	0
379	12111	0	47	0	193	90	0	135	1	0
380	12121	0	50	0	198	82	1	136	1	0
381	12141	0	48	0	263	76	0	102	0	0
382	12151	0	48	0	254	74	0	124	0	0
383	12181	0	64	0	248	74	0	126	1	0
384	12221	0	43	0	197	84	0	122	1	0
385	12231	0	41	0	282	98	0	132	0	1
386	12241	0	48	0	238	106	0	144	1	1
387	12251	0	50	0	156	74	0	122	1	0
388	12271	0	46	0	234	70	0	120	1	0
389	12281	0	44	0	203	82	0	110	1	0
390	12291	1	65	0	200	90	0	160	1	1
391	12293	0	44	0	209	84	0	132	1	0
392	12311	0	40	0	245	94	0	142	0	0
393	12351	0	56	0	124	86	0	142	0	0

5.3 Example of Coronary Heart Disease Data

394	12371	0	56	1	199	86	1	154	1	0
395	12381	1	47	0	148	85	1	145	1	0
396	12391	0	48	0	246	92	0	122	1	0
397	12401	0	46	0	233	96	0	138	0	1
398	12431	0	48	0	265	100	1	142	1	1
399	12461	0	50	0	207	86	1	142	1	0
400	12471	0	69	0	227	72	1	108	1	0
401	12481	0	45	0	205	130	1	182	1	1
402	12641	0	57	0	189	102	1	128	1	1
403	12681	1	69	0	191	102	0	164	1	1
404	12741	0	45	0	171	91	0	145	1	0
405	12742	0	52	0	178	91	1	145	1	0
406	12751	0	63	0	229	94	0	148	1	0
407	12761	0	61	1	169	90	0	140	1	0
408	12801	0	48	0	238	88	0	134	1	0
409	12831	1	45	0	216	94	0	138	1	0
410	12861	0	66	1	178	110	0	198	0	1
411	12891	0	54	0	173	92	0	162	0	1
412	12901	0	45	0	173	64	0	120	1	0
413	12911	1	66	0	180	104	1	162	1	1
414	12921	0	53	0	168	110	0	154	1	1
415	12941	0	40	0	277	80	0	120	0	0
416	13021	0	55	0	181	78	0	132	1	0
417	13041	0	48	0	272	98	1	156	1	1
418	13051	0	49	0	307	88	0	130	0	0
419	13101	0	61	0	203	94	1	146	0	0
420	13111	0	41	0	212	90	0	120	1	0
421	13121	0	43	0	248	118	0	142	1	1
422	13131	0	47	0	208	110	0	160	1	1
423	13321	0	46	0	218	86	0	126	1	0
424	13351	0	63	1	163	76	0	175	1	1
425	13391	0	62	0	261	88	0	130	1	0
426	13421	0	72	0	224	100	1	190	0	1
427	13431	0	50	0	292	80	1	128	1	0
428	13451	0	46	0	202	100	1	172	1	1
429	13461	0	44	0	145	72	0	114	1	0
430	13471	0	46	1	183	88	1	162	1	1
431	13481	0	47	0	188	88	0	126	1	0
432	13511	0	51	1	209	106	0	180	1	1
433	13521	0	46	0	217	84	0	144	1	0
434	13531	0	47	0	180	78	0	126	1	0
435	13541	0	44	0	190	90	0	140	0	0
436	13551	0	55	0	211	80	0	115	1	0
437	13571	0	56	0	204	76	0	124	1	0
438	13591	0	54	0	185	98	0	170	0	1

5 Logistic Regression

439	13611	1	50	0	206	70	0	108	1	0
440	13641	0	59	0	265	96	0	150	1	1
441	13651	0	47	0	246	80	0	130	0	0
442	13661	0	65	1	171	102	0	166	0	1
443	13662	0	41	0	211	91	0	145	0	0
444	13671	0	47	0	139	96	1	192	1	1
445	13691	0	49	0	155	84	0	124	0	0
446	13721	0	50	0	229	90	0	134	0	0
447	13731	0	56	1	148	110	0	168	1	1
448	13751	0	50	0	198	86	0	134	1	0
449	13761	0	55	1	186	120	0	172	1	1
450	13771	0	52	0	211	70	0	112	1	0
451	13811	0	57	0	210	80	0	120	1	0
452	13841	1	47	1	212	122	1	220	1	1
453	13861	0	59	0	227	70	0	122	0	0
454	13901	0	47	0	232	90	0	142	0	0
455	13911	0	42	0	176	88	0	122	1	0
456	13931	0	56	0	166	86	0	126	0	0
457	13941	1	43	0	268	88	1	132	1	0
458	13951	0	55	0	178	80	0	104	1	0
459	13961	0	49	1	147	134	1	300	1	1
460	13971	0	71	1	164	94	0	174	0	1
461	14041	0	71	1	187	114	1	172	0	1
462	14071	0	69	1	165	96	1	140	0	1
463	14691	0	47	0	250	88	1	122	1	0
464	14701	0	44	0	199	70	0	120	1	0
465	14711	1	65	1	233	116	1	180	1	1
466	14731	0	65	0	182	74	0	124	0	0
467	14741	0	40	0	210	94	0	128	0	0
468	14751	0	47	0	235	86	0	128	1	0
469	14761	0	54	0	172	92	0	144	1	0
470	14771	0	64	1	198	100	0	178	0	1
471	14781	0	59	1	212	90	1	198	0	1
472	14801	0	72	0	285	90	1	150	1	0
473	14811	0	52	0	194	82	0	132	0	0
474	14861	0	56	0	237	70	0	106	0	0
475	14871	0	56	0	153	66	1	96	0	0
476	14881	0	54	0	219	92	0	152	1	0
477	14901	0	60	1	188	114	1	210	1	1
478	14911	0	63	0	276	95	1	145	0	1
479	14931	0	73	0	203	68	0	130	1	0
480	14941	1	49	0	228	98	0	140	1	1
481	14981	0	57	0	199	80	0	134	0	0
482	15001	0	53	1	162	90	0	158	1	0
483	15011	0	43	0	221	94	0	125	1	0

5.3 Example of Coronary Heart Disease Data

484	15021	0	68	0	150	78	0	145	0	0
485	15031	0	53	0	140	80	0	120	1	0
486	15041	0	50	0	162	65	1	110	0	0
487	15061	0	49	0	171	86	1	125	1	0
488	15091	0	62	0	206	94	0	144	0	0
489	15111	0	52	0	226	76	0	130	1	0
490	15121	0	57	1	113	94	1	146	1	0
491	15141	0	45	0	197	78	0	118	0	0
492	15161	0	50	0	180	88	0	132	1	0
493	15171	0	50	1	180	118	1	214	1	1
494	15191	0	53	0	196	86	0	144	1	0
495	15201	0	51	0	211	98	0	135	0	1
496	15211	0	58	1	164	96	1	155	1	1
497	15221	0	64	0	218	85	0	154	0	0
498	15251	1	49	0	191	76	0	132	1	0
499	15271	0	55	0	189	64	0	110	0	0
500	15311	0	50	0	156	82	0	114	1	0
501	15321	0	50	0	223	80	0	130	1	0
502	15361	0	57	0	165	76	0	132	1	0
503	15401	0	55	1	200	94	1	188	1	1
504	15421	0	48	1	162	135	1	250	1	1
505	15431	0	56	1	207	110	0	172	1	1
506	15441	0	72	0	262	84	0	172	1	1
507	15511	1	67	1	236	106	1	200	0	1
508	15541	0	70	1	192	90	1	162	0	1
509	15562	0	57	1	203	100	0	170	1	1
510	15611	0	67	1	200	160	1	224	0	1
511	15641	0	62	0	280	86	0	124	1	0
512	15651	0	72	1	229	140	1	270	1	1
513	15661	0	70	0	290	84	0	138	0	0
514	15671	0	65	0	222	88	0	146	0	0
515	15691	0	58	0	259	100	1	154	0	1
516	15711	0	64	0	205	80	0	140	0	0
517	15761	0	44	0	276	74	0	112	1	0
518	15791	0	55	0	171	68	0	110	0	0
519	15831	0	71	0	287	90	0	130	0	0
520	15851	1	72	0	174	78	1	192	1	1
521	15882	0	71	0	277	110	1	200	0	1
522	15891	0	54	0	192	85	0	130	0	0
523	15911	0	51	0	196	90	0	128	1	0
524	15921	0	68	0	203	74	0	138	1	0
525	15931	0	47	0	271	85	0	145	1	0
526	15941	0	49	1	169	85	1	145	0	0
527	15951	0	48	0	201	98	0	150	1	1
528	15981	0	74	0	244	94	0	164	0	1

5 Logistic Regression

529	15991	0	49	0	161	92	0	120	0	0
530	16321	1	53	0	192	106	0	164	1	1
531	16431	0	46	0	192	86	1	116	1	0
532	16441	0	59	0	230	84	0	158	1	0
533	16461	0	50	0	312	98	1	138	0	1
534	16481	1	69	1	230	100	1	170	0	1
535	16501	1	75	1	233	90	1	222	1	1
536	16531	0	42	0	207	72	0	106	1	0
537	16541	0	50	0	317	90	1	138	0	0
538	16571	0	44	0	213	84	0	118	1	0
539	16581	0	44	0	220	98	0	140	0	1
540	16591	0	42	0	225	95	0	140	0	1
541	16622	0	42	0	288	104	0	150	1	1
542	16691	0	44	0	168	94	1	134	1	0
543	16701	0	57	1	182	96	1	138	1	1
544	16711	1	68	1	242	84	0	128	1	0
545	16752	0	69	0	258	82	0	145	1	0
546	16761	0	74	1	172	100	0	190	1	1
547	16841	0	56	1	239	140	1	220	1	1
548	16871	1	58	1	209	94	1	140	1	0
549	16891	0	46	0	181	84	0	124	0	0
550	16911	0	60	1	199	100	0	162	0	1
551	16931	0	62	0	217	90	0	144	1	0
552	16971	0	74	0	200	78	0	118	1	0
553	17071	0	44	0	268	80	0	126	1	0
554	17111	0	54	0	202	86	0	134	0	0
555	17121	0	49	0	224	86	1	134	1	0
556	17131	0	46	0	302	102	0	160	1	1
557	17151	0	45	0	239	90	0	128	0	0
558	17161	0	57	0	205	88	0	140	0	0
559	17171	0	56	1	192	170	1	270	0	1
560	17181	0	42	0	282	114	0	170	1	1
561	17191	0	52	0	232	94	0	144	1	0
562	17211	0	49	0	229	92	0	162	1	1
563	17231	0	51	0	336	86	0	130	1	0
564	17251	0	40	0	146	84	0	125	0	0
565	17271	0	43	0	224	72	0	115	0	0
566	17291	0	45	0	228	76	1	136	1	0
567	17361	0	63	0	211	108	0	144	1	1
568	17401	0	52	0	212	76	0	118	0	0
569	17481	1	67	1	243	118	1	220	1	1
570	17961	0	72	1	208	94	1	174	1	1
571	17991	0	54	0	284	98	0	146	1	1
572	18061	0	52	0	190	88	1	130	1	0
573	18071	0	49	0	264	92	0	162	0	1

5.3 Example of Coronary Heart Disease Data

574	18101	0	42	0	288	108	0	146	0	1
575	18121	0	41	0	181	94	0	136	1	0
576	18131	1	56	1	283	100	0	188	1	1
577	18141	0	46	0	217	66	0	120	1	0
578	18151	1	52	0	250	80	0	132	1	0
579	18161	0	44	0	209	70	0	116	1	0
580	18171	1	43	0	189	106	0	154	1	1
581	18201	0	73	0	190	78	0	138	0	0
582	18401	0	54	0	223	82	0	122	1	0
583	18411	0	46	0	241	84	0	120	1	0
584	18421	0	44	0	214	96	1	142	1	1
585	18441	0	43	0	207	86	0	122	0	0
586	18481	0	46	0	186	86	0	130	0	0
587	18491	1	74	1	212	70	1	144	1	0
588	18511	0	54	1	211	94	1	152	0	0
589	18521	0	63	0	223	86	0	158	1	0
590	18551	0	58	1	206	108	1	192	0	1
591	18581	0	50	0	194	92	0	134	1	0
592	18631	0	71	0	193	82	0	115	1	0
593	18661	0	52	0	213	90	0	140	1	0
594	18681	0	63	0	318	82	1	126	0	0
595	18711	0	66	0	216	104	1	154	0	1
596	18731	0	60	0	211	66	0	128	1	0
597	18752	0	47	0	219	88	0	128	1	0
598	18771	0	57	0	322	98	0	144	1	1
599	18801	0	66	1	239	100	1	184	1	1
600	18841	0	66	1	195	104	0	158	1	1
601	18871	0	47	0	243	78	0	118	1	0
602	18921	0	60	0	223	92	1	122	1	0
603	18971	0	48	0	174	102	0	160	1	1
604	19003	0	61	1	163	86	1	144	1	0
605	19011	0	64	0	225	90	0	160	1	1
606	19061	1	46	0	252	122	0	158	1	1
607	19091	0	49	0	261	102	0	166	1	1
608	19121	0	51	0	184	88	0	118	1	0
609	19161	0	64	0	206	82	0	152	0	0

```
colnames(CHD.data)
```

```
[1] "id" "chd" "age" "cat" "chl" "dbp" "ecg" "sbp" "smk" "hpt"
```

5.3.2 Fit Logistic Regression Model for a Single Variable

```

vars <- c("smk", "sbp", "age", "chl")
#jit <- 0.01 # global jitter amount for y

## par(mfrow = c(2, 2), mar = c(4, 4, 2, 4) + 0.1) # extra right margin for axis(4)

for (v in vars) {
  ## Univariate logistic regression using ORIGINAL variable name in the formula
  fit <- glm(
    formula = reformulate(v, response = "chd"),
    data     = CHD.data,
    family   = binomial()
  )
  print(summary(fit))

  ## Base scatter of chd with small jitter (left axis: probability scale)
  plot(
    CHD.data[[v]],
    jitter(CHD.data$chd, amount = jit),
    pch = 16, col = rgb(0, 0, 0, 0.45),
    xlab = v, ylab = "chd (jittered)",
    main = paste("chd vs", v),
    ylim = c(-0.1, 1.1)
  )

  ## Fitted (x) in red (left axis)
  if (length(unique(CHD.data[[v]])) == 2) {
    # binary predictor
    xcat <- sort(unique(CHD.data[[v]]))
    nd   <- setNames(data.frame(xcat), v)
    pcat <- predict(fit, newdata = nd, type = "response")
    points(xcat, pcat, pch = 19, col = "red")
    lines(xcat, pcat, col = "red", lwd = 2)

    # Right-axis: logit{ (x) } with fixed y-limits
    logit_p <- log(pcat / (1 - pcat))
    par(new = TRUE)
    plot(
      xcat, logit_p, type = "l", lwd = 2, col = "blue",
      axes = FALSE, xlab = "", ylab = "",
      xlim = range(CHD.data[[v]]), ylim = c(-2.5, 0)
    )
    axis(4)
  }
}

```

```

mtext("logit(p(x))", side = 4, line = 3)
par(new = FALSE)

} else {
  # continuous predictor
  xg <- seq(min(CHD.data[[v]]), max(CHD.data[[v]]), length.out = 400)
  nd <- setNames(data.frame(xg), v)
  pg <- predict(fit, newdata = nd, type = "response")
  lines(xg, pg, col = "red", lwd = 2)

  # Right-axis: logit{ (x)} with fixed y-limits
  logit_pg <- log(pg / (1 - pg))
  par(new = TRUE)
  plot(
    xg, logit_pg, type = "l", lwd = 2, col = "blue",
    axes = FALSE, xlab = "", ylab = "",
    xlim = range(xg), ylim = c(-2.5, 0)
  )
  axis(4)
  mtext("logit(p(x))", side = 4, line = 3)
  par(new = FALSE)
}
}

```

Call:

```
glm(formula = reformulate(v, response = "chd"), family = binomial(),
    data = CHD.data)
```

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-2.4898	0.2524	-9.865	<2e-16 ***
smk	0.6706	0.2919	2.297	0.0216 *

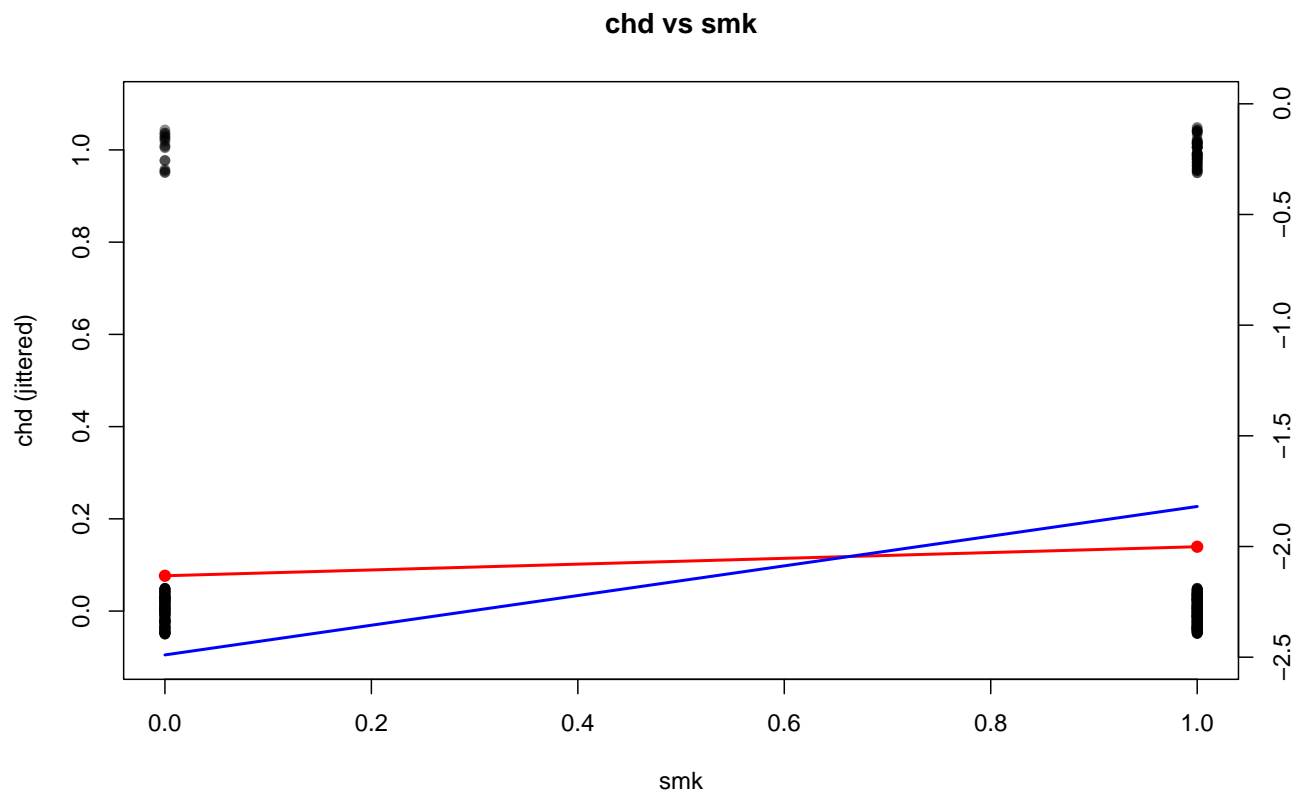
Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 438.56 on 608 degrees of freedom
 Residual deviance: 432.81 on 607 degrees of freedom
 AIC: 436.81

Number of Fisher Scoring iterations: 5

5 Logistic Regression



Call:

```
glm(formula = reformulate(v, response = "chd"), family = binomial(),
     data = CHD.data)
```

Coefficients:

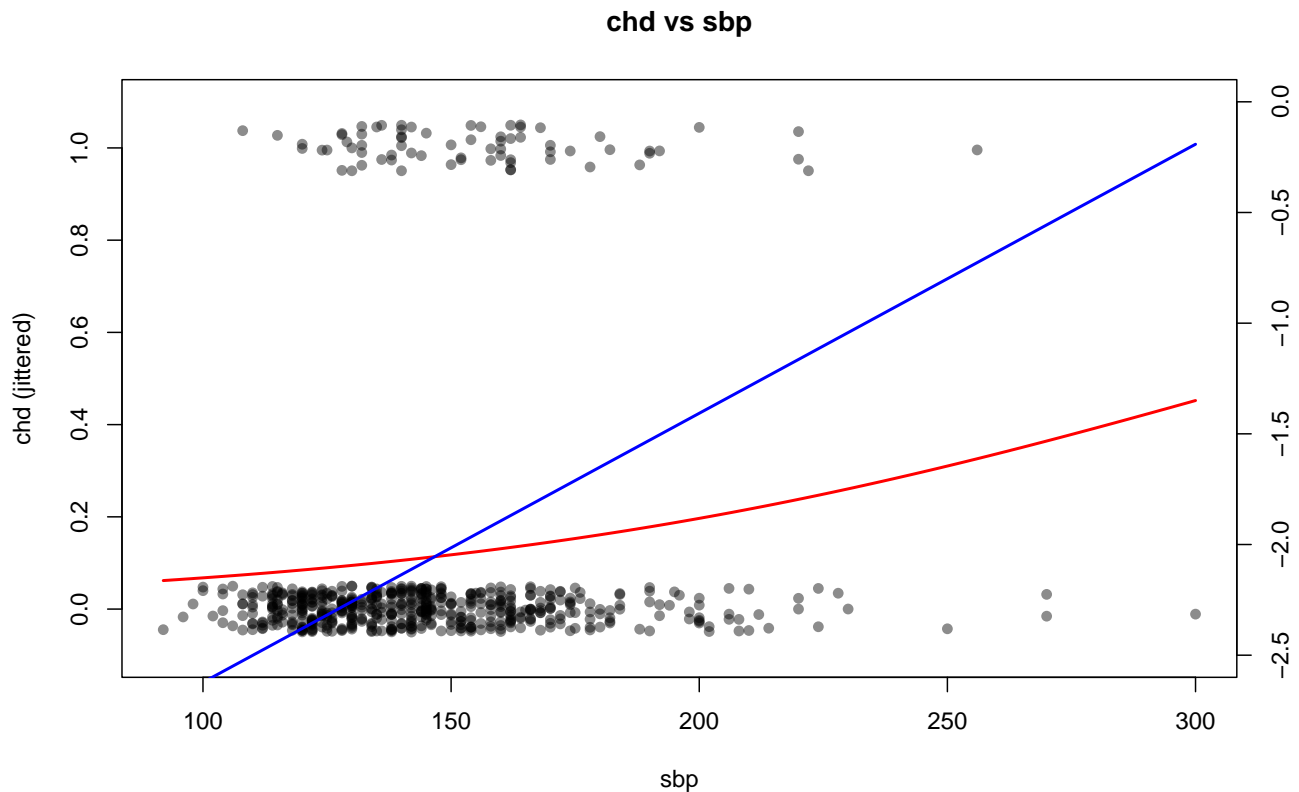
	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-3.837912	0.629805	-6.094	1.1e-09 ***
sbp	0.012154	0.004036	3.011	0.0026 **

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 438.56 on 608 degrees of freedom
Residual deviance: 430.06 on 607 degrees of freedom
AIC: 434.06

Number of Fisher Scoring iterations: 4



```
Call:
glm(formula = reformulate(v, response = "chd"), family = binomial(),
    data = CHD.data)
```

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-4.47833	0.75610	-5.923	3.16e-09 ***
age	0.04445	0.01315	3.381	0.000723 ***

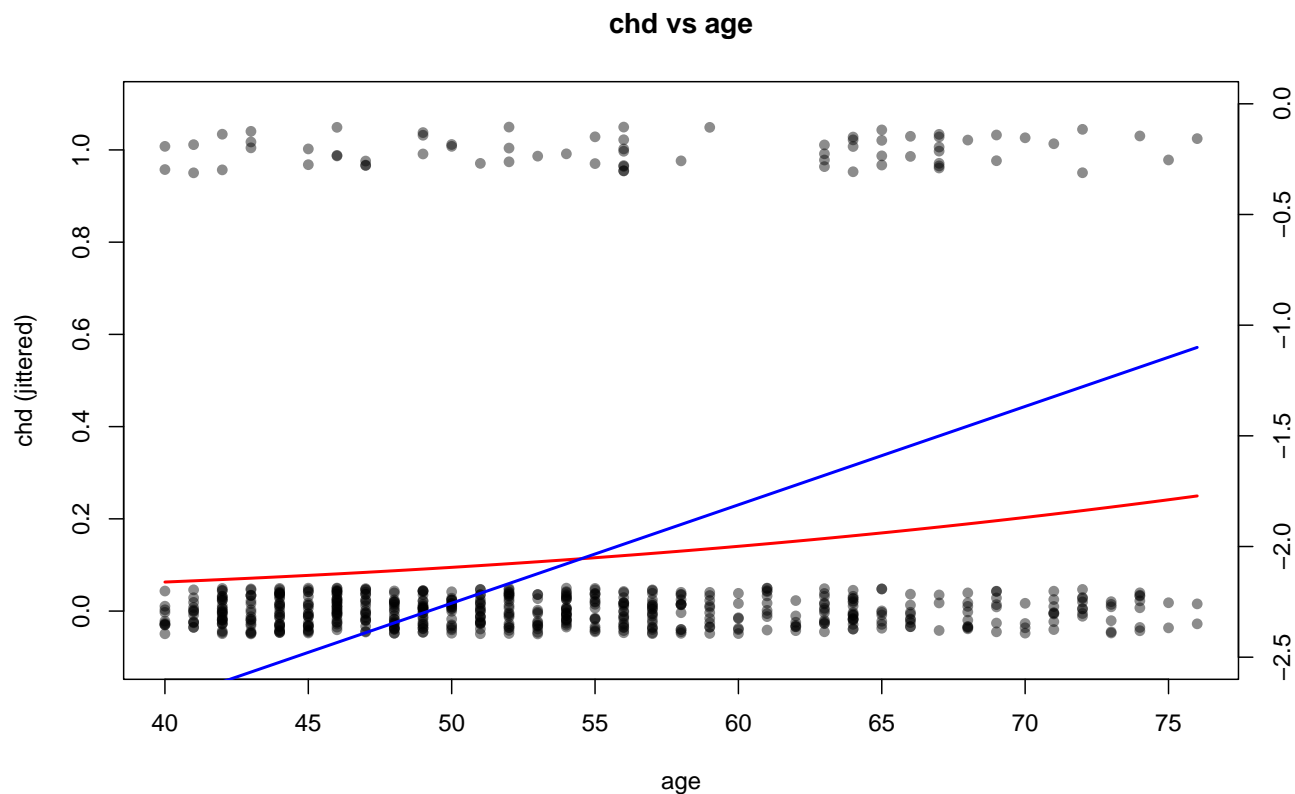
Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 438.56 on 608 degrees of freedom
 Residual deviance: 427.22 on 607 degrees of freedom
 AIC: 431.22

Number of Fisher Scoring iterations: 5

5 Logistic Regression



```
Call:
glm(formula = reformulate(v, response = "chd"), family = binomial(),
     data = CHD.data)
```

Coefficients:

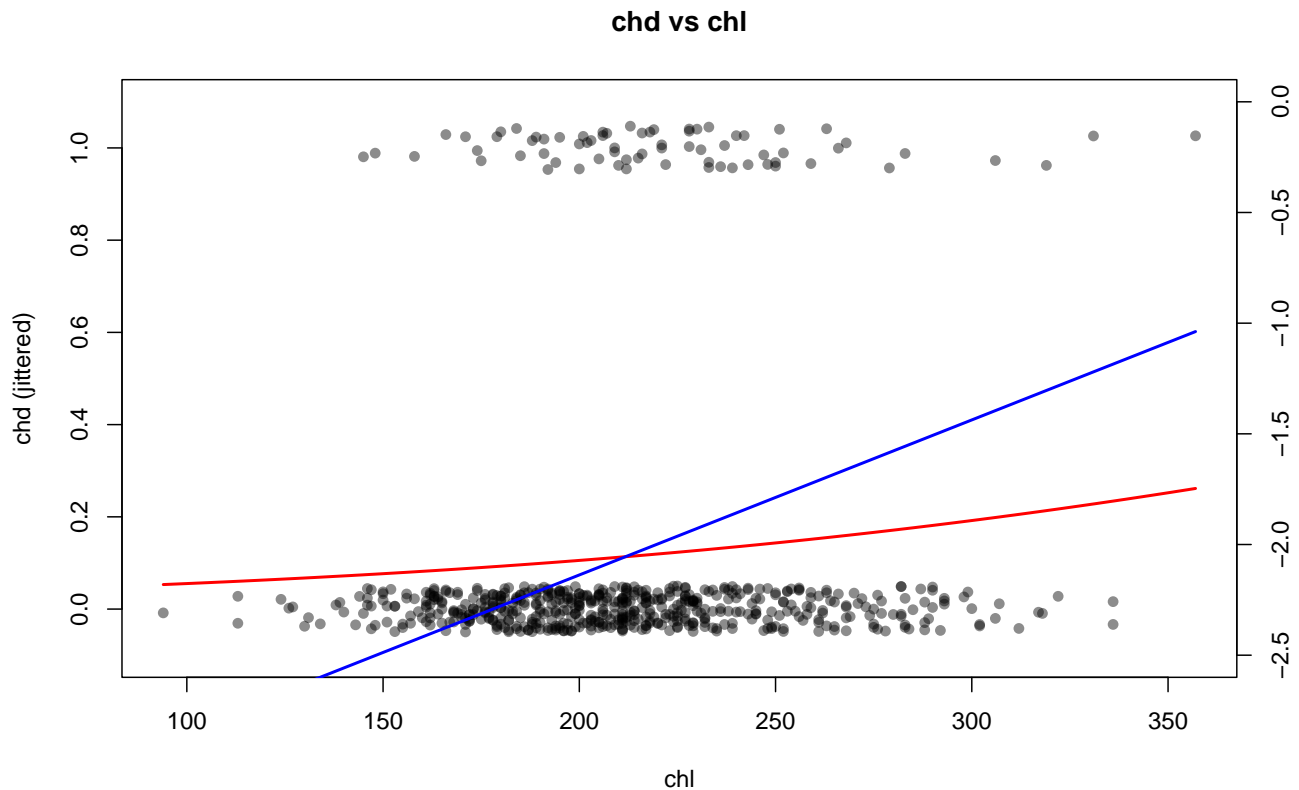
	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-3.538260	0.686879	-5.151	2.59e-07 ***
chl	0.007004	0.003064	2.286	0.0223 *

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 438.56 on 608 degrees of freedom
 Residual deviance: 433.42 on 607 degrees of freedom
 AIC: 437.42

Number of Fisher Scoring iterations: 4



5.3.3 Fit Logistic Regression Model with all variables

We fit a logistic regression with a logit link:

```
fit1_chd <- glm(
  chd ~ smk + cat + sbp + age + chl + ecg + hpt,
  data = CHD.data,
  family = binomial(link = "logit")
)
summary(fit1_chd)
```

Call:

```
glm(formula = chd ~ smk + cat + sbp + age + chl + ecg + hpt,
     family = binomial(link = "logit"), data = CHD.data)
```

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)	
(Intercept)	-6.048892	1.345165	-4.497	6.9e-06	***
smk	0.855951	0.306505	2.793	0.00523	**
cat	0.732763	0.376129	1.948	0.05139	.

5 Logistic Regression

```
sbp      -0.006995   0.006976  -1.003   0.31600
age       0.033956   0.015344   2.213   0.02690 *
chl       0.008970   0.003274   2.740   0.00615 **
ecg       0.417776   0.295553   1.414   0.15750
hpt       0.655498   0.359976   1.821   0.06861 .
---
```

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

```
Null deviance: 438.56  on 608  degrees of freedom
Residual deviance: 399.35  on 601  degrees of freedom
AIC: 415.35
```

Number of Fisher Scoring iterations: 5

Notes for interpretation:

- Positive coefficients increase the log-odds of CHD; negative coefficients decrease it.
- For indicator variables (e.g., smk), $\exp(\beta)$ is the adjusted odds ratio comparing the group with value 1 versus 0, holding others fixed.
- For continuous predictors (e.g., sbp, age), $\exp(\beta)$ is the multiplicative change in the odds for a one-unit increase. For a d -unit increase, the OR is $\exp(d * \beta)$.

5.4 Inference for Coefficients: Confidence Intervals and Covariance Matrix

We extract profile-likelihood CIs and the covariance matrix to confirm standard errors.

```
ci_95 <- confint(fit1_chd, level = 0.95)      # profile-likelihood CI
vcov_mat <- vcov(fit1_chd)                   # covariance matrix of coefficients
se_vec  <- sqrt(diag(vcov_mat))              # standard errors

ci_95
```

	2.5 %	97.5 %
(Intercept)	-8.718003347	-3.427904298
smk	0.275699158	1.483333169
cat	-0.006873216	1.471885644
sbp	-0.021166144	0.006266328
age	0.003687290	0.064005215
chl	0.002533226	0.015404292
ecg	-0.171584621	0.990632546


```
hpt          -0.050184520  1.364993401
```

```
vcov_mat
```

```

              (Intercept)          smk          cat          sbp
(Intercept)  1.809468553 -1.014526e-01  0.1391440386 -4.908229e-03
smk          -0.101452600  9.394560e-02 -0.0032961000 -1.653230e-04
cat          0.139144039 -3.296100e-03  0.1414730484 -9.299960e-04
sbp          -0.004908229 -1.653230e-04 -0.0009299960  4.866901e-05
age          -0.011142995  7.738971e-04 -0.0017879998 -1.311191e-05
chl          -0.002111134  3.161443e-05  0.0003146354 -1.821907e-06
ecg          0.003442546  9.255483e-03 -0.0204455233 -2.982539e-04
hpt          0.139817180  6.954592e-03 -0.0044220690 -1.486400e-03
              age          chl          ecg          hpt
(Intercept) -1.114300e-02 -2.111134e-03  3.442546e-03  1.398172e-01
smk          7.738971e-04  3.161443e-05  9.255483e-03  6.954592e-03
cat          -1.788000e-03  3.146354e-04 -2.044552e-02 -4.422069e-03
sbp          -1.311191e-05 -1.821907e-06 -2.982539e-04 -1.486400e-03
age          2.354442e-04 -1.480501e-06 -4.972374e-05  4.044434e-04
chl          -1.480501e-06  1.071766e-05  5.040548e-05 -6.046197e-05
ecg          -4.972374e-05  5.040548e-05  8.735130e-02  9.506863e-04
hpt          4.044434e-04 -6.046197e-05  9.506863e-04  1.295828e-01
```

```
se_vec # should match the SE column in summary(fit1_chd)
```

```

(Intercept)          smk          cat          sbp          age          chl
1.345164879 0.306505459 0.376129032 0.006976318 0.015344190 0.003273784
          ecg          hpt
0.295552531 0.359976081
```

5.5 Inference for Odds Ratios

5.5.1 Interpretation of Odds Ratios in Logistic Regression

A multiple logistic regression model expresses the log-odds (logit) of an event as a linear function of predictors:

$$\log\left(\frac{p}{1-p}\right) = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \cdots + \beta_k x_k.$$

Here,

5 Logistic Regression

- $p = \Pr(Y = 1 \mid x_1, x_2, \dots, x_k)$ is the probability of the event,
- β_0 is the intercept, and
- each β_j represents the **change in the log-odds** of the event per one-unit increase in x_j , *holding all other variables constant*.

Exponentiating both sides gives the model in odds form:

$$\frac{p}{1-p} = \exp(\beta_0) \times \exp(\beta_1 x_1) \times \exp(\beta_2 x_2) \times \dots \times \exp(\beta_k x_k).$$

An R function for OR at two Profiles

The `or_from_predict` R function is a utility designed to calculate the Odds Ratio (OR) and its 95% confidence interval (CI) between two specific covariate profiles (`new1` and `new0`) for a given logistic regression model (`fit`). The calculation is performed on the link (logit) scale. For a logistic model $\text{logit}(p) = \eta = \mathbf{X}\beta$, the log-Odds Ratio (`logOR`) is the difference between the linear predictors (η_1, η_0) for the two profiles:

$$\widehat{\text{logOR}} = \eta_1 - \eta_0 = (\mathbf{x}_1^T - \mathbf{x}_0^T)\beta = \mathbf{c}^T\beta \quad (5.1)$$

Here, $\mathbf{c} = \mathbf{x}_1 - \mathbf{x}_0$ is the linear contrast vector derived from the model matrices of the two profiles. The function estimates the variance of this contrast as $\text{Var}(\widehat{\text{logOR}}) = \mathbf{c}^T \mathbf{V} \mathbf{c}$, where \mathbf{V} is the model's variance-covariance matrix (`vcov(fit)`). The standard error $SE = \sqrt{\mathbf{c}^T \mathbf{V} \mathbf{c}}$ is used to compute the $100(1 - \alpha)\%$ confidence interval for the `logOR`:

$$\widehat{\text{logOR}} \pm z_{1-\alpha/2} \times SE.$$

These values (estimate and CI bounds) are then exponentiated to produce the final $\widehat{\text{OR}} = \exp(\widehat{\text{logOR}})$ and its 95% CI. The function also prints two helpful summaries to the console: a data frame showing only the variables that differ between the `new0` and `new1` profiles, and a 2x3 table presenting the estimates and CIs for both the OR and the `logOR`.

The R function to find ORs

```
## Compute OR and 95% CI via predict() on the LINK scale
## OR = exp( eta(new1) - eta(new0) ), where eta(.) = logit{(.)}
## Compute OR via predict() contrast on the LINK scale, also:
## (ii) print a 2-row data.frame of only variables that differ between new0 and new1
## (iii) print a 2x3 table (rows: OR, logOR; cols: Estimate, CI_low, CI_up)
or_from_predict <- function(fit, new1, new0, level = 0.95, digits = 4, tol = 1e-12) {
  stopifnot(is.data.frame(new1), is.data.frame(new0))

  ## --- REFACTORED SECTION START ---
  ## ---- (ii) Two-row data.frame with only changed variables ----

  ## Helper function to find differing variables between two profiles
  ## This is defined *inside* or_from_predict for encapsulation
  get_changed_vars <- function(d0, d1, tolerance) {
```

```

common <- intersect(names(d0), names(d1))
diffv <- vapply(common, function(nm) {
  x0 <- d0[[nm]]; x1 <- d1[[nm]]
  if (is.numeric(x0) && is.numeric(x1)) {
    !isTRUE(all.equal(as.numeric(x0), as.numeric(x1), tolerance = tolerance))
  } else {
    !identical(x0, x1)
  }
}, logical(1))

keep <- common[diffv]
if (length(keep) == 0L) {
  out <- data.frame(`_no_changes_` = "no differences")
  rownames(out) <- c("new0", "new1")
  return(out)
}
out <- rbind(d0[keep], d1[keep])
rownames(out) <- c("new0", "new1")
out
}

## Call the helper function
changes_df <- get_changed_vars(new0, new1, tol)
## --- REFACTORED SECTION END ---

## ---- Linear contrast for log-OR and its variance ----

## (i) Calculate logOR estimate using predict(type="link")
## eta(.) = logit{p(.)}
eta1 <- predict(fit, newdata = new1, type = "link")
eta0 <- predict(fit, newdata = new0, type = "link")
logOR_hat <- as.numeric(eta1 - eta0) # logOR = eta1 - eta0

## (ii) Calculate standard error using the contrast vector 'cvec'
X1 <- model.matrix(delete.response(terms(fit)), data = new1)
X0 <- model.matrix(delete.response(terms(fit)), data = new0)
cvec <- as.numeric(X1 - X0)
V <- vcov(fit)
se_logOR <- sqrt(as.numeric(t(cvec) %*% V %*% cvec))

alpha <- 1 - level
z <- qnorm(1 - alpha / 2)
ci_log <- c(logOR_hat - z * se_logOR, logOR_hat + z * se_logOR)

```

5 Logistic Regression

```
## ---- 2x3 table: rows OR and logOR; columns Estimate, CI_low, CI_up ----
res_tab <- data.frame(
  Estimate = c(exp(logOR_hat),      logOR_hat),
  CI_low   = c(exp(ci_log[1L]),     ci_log[1L]),
  CI_up    = c(exp(ci_log[2L]),     ci_log[2L]),
  row.names = c("OR", "logOR")
)

## ---- Print requested items ----
cat("\nVariables that differ between new0 and new1:\n")
print(changes_df)
cat("\nOdds Ratio summary:\n")
print(round(res_tab, digits = digits)) # Added rounding for neatness

## ---- Return (invisibly) ----
invisible(list(
  OR      = exp(logOR_hat),
  CI_OR   = exp(ci_log),
  logOR   = logOR_hat,
  CI_logOR = ci_log,
  se_logOR = se_logOR,
  changes  = changes_df,
  table    = res_tab
))
}

## --- Example usage ---
## Suppose 'fit1_chd' is your fitted model and 'CHD.data' is your data
## base_prof <- as.data.frame(lapply(CHD.data, function(col) if (is.numeric(col)) mean(col) el
## new0 <- base_prof; new0$smk <- 0
## new1 <- base_prof; new1$smk <- 1
## or_from_predict(fit1_chd, new1 = new1, new0 = new0)

mean_profile <- function(data, vars_binary_as = c(0,1)) {
  ## Build a single-row data.frame of typical values:
  out <- lapply(data, function(col) {
    if (is.numeric(col)) {
      # If strictly 0/1, keep mean (works fine for GLM prediction),
      # or switch to mode if you prefer.
      if (all(col %in% c(0,1))) mean(col) else mean(col, na.rm = TRUE)
    } else {
      # Fallback to first level for factors/characters
      if (is.factor(col)) levels(col)[1] else unique(col)[1]
    }
  })
}
```

```

    }
  })
  as.data.frame(out)
}

```

5.5.2 Examples of Finding ORs and Their CIs for the CHD Dataset

5.5.2.1 OR Smoking (smk) (1 vs 0)

```

### 1) Smoking OR: smk = 1 vs 0 (other vars at their means)
## Example profiles at sample means (adjust as you like)
base_prof <- mean_profile(CHD.data)
new0 <- base_prof; new0$smk <- 0
new1 <- base_prof; new1$smk <- 1

res_smk <- or_from_predict(fit1_chd, new1 = new1, new0 = new0)

```

Variables that differ between new0 and new1:

```

      smk
new0    0
new1    1

```

Odds Ratio summary:

	Estimate	CI_low	CI_up
OR	2.3536	1.2907	4.2917
logOR	0.8560	0.2552	1.4567

How to read this:

- $OR_{smk} > 1$ suggests higher odds of CHD among smokers (adjusted for other variables). If the 95% CI excludes 1, the association is statistically significant at the 5% level.

5.5.2.2 OR for Systolic Blood Pressure (sbp): from 120 to 160

We compute the adjusted OR for a 40-unit increase in sbp (from 120 to 160):

```

### 2) SBP OR: 160 vs 120 (other vars at their means)
new0 <- base_prof; new0$sbp <- 120
new1 <- base_prof; new1$sbp <- 160

res_sbp <- or_from_predict(fit1_chd, new1 = new1, new0 = new0)

```

5 Logistic Regression

Variables that differ between new0 and new1:

```
      sbp
new0 120
new1 160
```

Odds Ratio summary:

	Estimate	CI_low	CI_up
OR	0.7559	0.4375	1.3062
logOR	-0.2798	-0.8267	0.2671

5.5.2.3 OR for Combined Effects of Two Variables: Smoking with an Age Difference

Suppose we compare two groups that differ in **smoking status** and **age**:

- **Group A:** smk = 1, age = 50 (all other covariates equal)
- **Group B:** smk = 0, age = 30

The log-odds contrast is ($A = \beta_{smk} + (50 - 20)\beta_{age}$), so the OR is ($\exp(A)$).

```
new0 <- base_prof; new0$age <- 30; new0$smk <- 0
new1 <- base_prof; new1$age <- 50; new1$smk <- 1

res_ageAsmk <- or_from_predict(fit1_chd, new1 = new1, new0 = new0)
```

Variables that differ between new0 and new1:

```
      age smk
new0  30   0
new1  50   1
```

Odds Ratio summary:

	Estimate	CI_low	CI_up
OR	4.6417	1.8546	11.6168
logOR	1.5351	0.6177	2.4525

5.6 Assessing Statistical Significance with Wilks' Theorem (Analogue of F-test for OLS)

In the context of logistic regression, Wilks' theorem provides the basis for the Likelihood Ratio Test (LRT) used to assess the significance of predictor variables. The theorem states that when comparing a full model (M_1) to a nested null model (M_0), the test statistic, Λ , asymptotically follows a chi-squared (χ^2) distribution under the null hypothesis (i.e., that the simpler model M_0 is correct).

5.6 Assessing Statistical Significance with Wilks' Theorem (Analogue of F-test for OLS)

The statistic Λ is calculated as the difference in the maximized log-likelihoods:

$$\Lambda = -2(\log L_0 - \log L_1) \quad (5.2)$$

where $\log L_0$ and $\log L_1$ are the log-likelihoods of the null and full models, respectively. In logistic regression, this is equivalent to the difference in the deviances: $\Lambda = \text{Deviance}_0 - \text{Deviance}_1$. This test statistic Λ represents the reduction in deviance (a measure of badness-of-fit) achieved by adding the extra predictors to the model.

The following R code chunk generates a conceptual plot of this relationship:

```
#| label: plot-lrt-concept
#| echo: false
#| fig-cap: "Conceptual plot of Deviance versus Number of Parameters, illustrating the Likelihood

library(ggplot2)

## 1. Create conceptual data for the plot
## These are just for illustration
n_obs <- 60 # Number of observations
p0 <- 1      # Parameters in null model (intercept)
p1 <- 21     # Parameters in full model (e.g., intercept + 7 predictors)
psat <- n_obs # Parameters in saturated model (1 per observation)

D0 <- 41 # Null deviance
D1 <- 20 # Full model deviance (residual deviance of M1)
D_sat <- 0 # Saturated model deviance

## Data frame for the three points
plot_data <- data.frame(
  model = c("M_0 (Null)", "M_1 (Full)", "M_Sat (Saturated)"),
  params = c(p0, p1, psat),
  deviance = c(D0, D1, D_sat),
  ## Add custom justification and nudges for labels
  hjust_val = c(0.5, 0.5, 1.1), # Right-align the last label
  nudge_x_val = c(0, 0, 0)
)

## 2. Create the ggplot
ggplot(plot_data, aes(x = params, y = deviance)) +
  ## Draw dashed guide lines for D0 and D1
  geom_segment(aes(x = p0, y = D0, xend = p1, yend = D0), linetype = "dashed", color = "grey70")
  geom_segment(aes(x = p1, y = D1, xend = psat, yend = D1), linetype = "dashed", color = "grey70")

  ## Connect the points with lines
  geom_line(color = "black", linetype = "solid", linewidth = 0.5) +
```

5 Logistic Regression

```
## Draw the main points
geom_point(size = 4, aes(color = model)) +

## --- MODIFIED LABEL PLACEMENT ---
## Label the points using custom nudge/justification
geom_text(
  aes(label = model, hjust = hjust_val, nudge_x = nudge_x_val),
  nudge_y = 2.5, # Use a much smaller vertical nudge
  size = 4
) +

## --- ADDED BACK D0 and D1 ANNOTATIONS ---
## D0 (Null Deviance)
geom_segment(
  aes(x = p0 - 2, y = D0, xend = p0 - 2, yend = D_sat), # Nudged left
  arrow = arrow(ends = "both", length = unit(0.1, "inches")),
  color = "darkgreen",
  linewidth = 1
) +
annotate(
  "text",
  x = p0 - 3, y = D0 / 2, # Nudged left
  label = "D[0]", parse = TRUE,
  color = "darkgreen", hjust = 0.5, size = 5
) +

## D1 (Residual Deviance of Full Model)
geom_segment(
  aes(x = psat + 2, y = D1, xend = psat + 2, yend = D_sat), # Nudged right
  arrow = arrow(ends = "both", length = unit(0.1, "inches")),
  color = "darkblue",
  linewidth = 1
) +
annotate(
  "text",
  x = psat + 3, y = D1 / 2, # Nudged right
  label = "D[1]", parse = TRUE,
  color = "darkblue", hjust = 0.5, size = 5
) +

## LRT statistic  $\Lambda = D0 - D1$ 
geom_segment(
  aes(x = p1 + 2, y = D0, xend = p1 + 2, yend = D1), # Nudged right
  arrow = arrow(ends = "both", length = unit(0.1, "inches")),
```



5.6 Assessing Statistical Significance with Wilks' Theorem (Analogue of F-test for OLS)

```
    color = "red",
    linewidth = 1
) +
annotate(
  "text",
  x = p1 + 3, y = D1 + (D0 - D1) / 2, # Nudged right
  label = "Lambda == D[0] - D[1]",
  parse = TRUE,
  color = "red", hjust = 0, size = 5
) +

## Customize axes to show the symbolic labels
scale_x_continuous(
  breaks = c(p0, p1, psat),
  labels = c(expression(p[0]), expression(p[1]), expression(n)),
  expand = expansion(mult = 0.1) # Add some padding
) +
scale_y_continuous(
  breaks = c(D_sat, D1, D0),
  labels = c(expression(0), expression(D[1]), expression(D[0]))
) +

## Labels and Title
labs(
  title = "Relationship between Deviance and Model Complexity",
  x = "Number of Parameters",
  y = "Model Deviance"
) +

## Clean theme
theme_minimal(base_size = 14) +
theme(
  plot.title = element_text(hjust = 0.5),
  legend.position = "none" # Remove legend, as points are labeled
) +
scale_color_manual(values = c("M_0 (Null)" = "blue", "M_1 (Full)" = "blue", "M_Sat (Saturated)"
```



As the diagram illustrates, the null model (M_0) has fewer parameters (p_0) and a higher deviance (D_0 , or worse fit), while the full model (M_1) has more parameters (p_1) and a lower deviance (D_1). The Likelihood Ratio Test statistic D is the magnitude of this drop in deviance.

For assessing the overall significance of a regression model (`fit1_chd`), this involves comparing it to its corresponding intercept-only (null) model. The degrees of freedom for the χ^2 test is the difference in the number of parameters, $df = p_1 - p_0$, which equals the number of predictors in the full model.

Here is an R code chunk demonstrating how to compute this p-value directly from a `glm` fit object, assuming it is named `fit1_chd`.

```
## Calculate the Likelihood Ratio Test statistic (D) and degrees of freedom (df)
## by comparing the model's deviance to the null (intercept-only) deviance,
## both of which are stored in the 'fit1_chd' object.
summary(fit1_chd)
```

Call:

```
glm(formula = chd ~ smk + cat + sbp + age + chl + ecg + hpt,
     family = binomial(link = "logit"), data = CHD.data)
```

5.6 Assessing Statistical Significance with Wilks' Theorem (Analogue of F-test for OLS)

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)	
(Intercept)	-6.048892	1.345165	-4.497	6.9e-06	***
smk	0.855951	0.306505	2.793	0.00523	**
cat	0.732763	0.376129	1.948	0.05139	.
sbp	-0.006995	0.006976	-1.003	0.31600	
age	0.033956	0.015344	2.213	0.02690	*
chl	0.008970	0.003274	2.740	0.00615	**
ecg	0.417776	0.295553	1.414	0.15750	
hpt	0.655498	0.359976	1.821	0.06861	.

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 438.56 on 608 degrees of freedom
Residual deviance: 399.35 on 601 degrees of freedom
AIC: 415.35

Number of Fisher Scoring iterations: 5

```
lrt_statistic <- fit1_chd$null.deviance - fit1_chd$deviance
lrt_df <- fit1_chd$df.null - fit1_chd$df.residual

## Compute the p-value from the chi-squared distribution
## We use lower.tail = FALSE to get P(ChiSq > D)
p_value <- pchisq(lrt_statistic, lrt_df, lower.tail = FALSE)

## Create and print the result in an ANOVA-like table
## Row 1: Null model
## Row 2: Full model (fit1_chd), showing the test against the null
lrt_table <- data.frame(
  "Resid. Df" = c(fit1_chd$df.null, fit1_chd$df.residual),
  "Resid. Dev" = c(round(fit1_chd$null.deviance, 4), round(fit1_chd$deviance, 4)),
  "Test Df" = c(NA, lrt_df),
  "Test Statistic (D)" = c(NA, round(lrt_statistic, 4)),
  "p-value" = c(NA, format.pval(p_value, digits = 4)),
  row.names = c("Null Model", "Full Model (fit1_chd)"),
  check.names = FALSE # Prevent R from changing 'p-value' to 'p.value'
)

cat("Likelihood Ratio Test for Model Significance:\n")
```

Likelihood Ratio Test for Model Significance:

5 Logistic Regression

```
lrt_table
```

	Resid. Df	Resid. Dev	Test Df	Test Statistic (D)	p-value
Null Model	608	438.5583	NA	NA	<NA>
Full Model (fit1_chd)	601	399.3539	7	39.2044	1.787e-06

Using built-in anova() function

```
fit0_chd <- glm (chd~1, data = CHD.data, family = binomial())  
anova(fit0_chd, fit1_chd)
```

Analysis of Deviance Table

Model 1: chd ~ 1

Model 2: chd ~ smk + cat + sbp + age + chl + ecg + hpt

	Resid. Df	Resid. Dev	Df	Deviance	Pr(>Chi)
1	608	438.56			
2	601	399.35	7	39.204	1.787e-06 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```
anova(fit1_chd, test="LRT")
```

Analysis of Deviance Table

Model: binomial, link: logit

Response: chd

Terms added sequentially (first to last)

	Df	Deviance	Resid. Df	Resid. Dev	Pr(>Chi)
NULL			608	438.56	
smk 1	1	5.7453	607	432.81	0.0165324 *
cat 1	1	14.3716	606	418.44	0.0001501 ***
sbp 1	1	0.7574	605	417.68	0.3841353
age 1	1	5.2821	604	412.40	0.0215455 *
chl 1	1	7.8619	603	404.54	0.0050489 **
ecg 1	1	1.8701	602	402.67	0.1714609
hpt 1	1	3.3159	601	399.35	0.0686113 .

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

5.7 Assessing Predictive Effect-Size (Analogue to R_{adj}^2)

While the LRT assesses overall model significance (in-sample fit), it's also crucial to evaluate how well the model predicts new, unseen data (out-of-sample performance). A common method is to split the data into a training set (e.g., 2/3 of the data) and a test set (e.g., 1/3). The model is fit using only the training data and then used to make predictions for the test data. We can then compare these predictions to the actual outcomes in the test set.

5.7.1 Understanding the Confusion Matrix and Metrics

To evaluate a model's predictive performance, we classify its probabilistic predictions using a threshold (typically 0.5) and compare them to the true outcomes in a **Confusion Matrix**:

	Predicted: 0	Predicted: 1
Actual: 0	True Negative (TN)	False Positive (FP)
Actual: 1	False Negative (FN)	True Positive (TP)

From this matrix, we derive several key performance metrics:

- **Misclassification Error Rate (ER):** The proportion of all predictions that were incorrect.

$$\text{Error Rate} = \frac{FP + FN}{TP + TN + FP + FN}$$

- **Precision (Positive Predictive Value):** Answers: "Of all the times the model predicted positive, how often was it correct?" This is crucial when the cost of a **False Positive** is high.

$$\text{Precision} = \frac{TP}{TP + FP}$$

- **Recall (Sensitivity or True Positive Rate):** Answers: "Of all the actual positive cases, how many did the model find?" This is crucial when the cost of a **False Negative** is high.

$$\text{Recall (TPR)} = \frac{TP}{TP + FN}$$

- **ROC Curve and AUC:** An **ROC (Receiver Operating Characteristic) Curve** is a graph that shows a model's diagnostic ability across *all possible classification thresholds*. It plots the **True Positive Rate (Recall)** on the y-axis against the **False Positive Rate** ($FPR = \frac{FP}{FP + TN}$) on the x-axis.

- **Interpretation:** The curve shows the trade-off between sensitivity (finding all the positives) and specificity (not mislabeling negatives). A random "no-skill" classifier is represented by a diagonal line from (0,0) to (1,1). A perfect classifier would hug the **top-left corner** (TPR = 1, FPR = 0).

- **AUC (Area Under the Curve):** The AUC summarizes the entire curve into a single number from 0 to 1. An AUC of 0.5 corresponds to a random guess, while an AUC of 1.0 represents a perfect model.
- **Precision-Recall (PR) Curve:** A PR Curve plots **Precision** (y-axis) against **Recall** (x-axis) at all possible thresholds.
 - **Interpretation:** This curve shows the trade-off between how *reliable* a positive prediction is (Precision) and how *complete* the model is at finding all positives (Recall).
 - **When to Use:** The PR curve is particularly informative when the dataset is **imbalanced** (i.e., one class, like “fraud” or “disease,” is much rarer than the other). Unlike the ROC curve, the PR curve’s baseline (the “no-skill” line) is a horizontal line at the proportion of positive cases, which makes it easier to see if the model is performing significantly better than chance in a low-positive-rate scenario. A perfect classifier would hug the **top-right corner** (Precision = 1, Recall = 1).

5.7.2 Illustration with the Simulated Dataset

This section applies the train/test split and model evaluation workflow to the `sim.data` created in the previous step.

```
## Load the pROC library for AUC calculation
## install.packages("pROC") # Uncomment to install if needed
library(pROC)

## --- 1. Split the data ---
## We use 'sim.data' which has 200 rows
set.seed(123) # for reproducibility
n_sim <- nrow(sim.data)
train_size_sim <- floor(2/3 * n_sim)
train_indices_sim <- sample(1:n_sim, size = train_size_sim)
train_data_sim <- sim.data[train_indices_sim, ]
test_data_sim <- sim.data[-train_indices_sim, ]

## --- 2. Refit the model on the training data ---
## We fit the model y ~ x on the training data
fit_train_sim <- glm(
  y ~ x,
  data = train_data_sim,
  family = binomial(link = "logit")
)

## --- 3. Make predictions on the test data ---
## Note: The true probabilities 'p' are also in test_data_sim
## We predict from the *fitted* model
pred_probs_sim <- predict(fit_train_sim, newdata = test_data_sim, type = "response")
```

Plotting the Predictive Probabilities with True Labels

```
## --- 5. Plot sorted predicted probabilities ---


## Create a data frame for plotting
plot_data_sim <- data.frame(
  Prob = pred_probs_sim,
  Actual = as.factor(test_data_sim$y),
  TrueProb = test_data_sim$p # Include true probs for comparison
)

## Sort by predicted probability
plot_data_sim <- plot_data_sim[order(plot_data_sim$Prob), ]
plot_data_sim$Rank <- 1:nrow(plot_data_sim)

## Create the plot
plot(
  plot_data_sim$Rank,
  plot_data_sim$Prob,
  pch = ifelse(plot_data_sim$Actual == 0, 1, 4),
  col = ifelse(plot_data_sim$Actual == 0, "blue", "red"),
  xlab = "Index (Sorted by Predicted Probability)",
  ylab = "Predicted Probability",
  main = "Predicted Probabilities vs. Actual Class (Simulated Data)",
  ylim = c(0, 1)
)
abline(h = 0.5, lty = 2, col = "black")
abline(h = 0.1, lty = 3, col = "grey")

## Add the true probability curve (sorted by predicted prob)
## This shows how well the fitted model's predictions align with the true probs
#lines(plot_data_sim$Rank, plot_data_sim$TrueProb[order(plot_data_sim$Prob)], col = "darkgreen",

## Add a legend
legend(
  "topleft",
  legend = c("Actual 0 (o)", "Actual 1 (x)"),
  pch = c(1, 4),
  lty = c(NA, NA),
  lwd = c(NA, NA),
  col = c("blue", "red")
)
```

**Confusion Matrix with threshold=0.5**

```
## --- 4. Assess accuracy ---

## 4a. Misclassification Error Rate (using 0.5 threshold)
threshold <- 0.5
pred_class_sim <- ifelse(pred_probs_sim > threshold, 1, 0)
conf_matrix_sim <- table(Actual = test_data_sim$y, Predicted = pred_class_sim)

## --- MODIFIED LINES START ---
cat("Confusion Matrix (Counts, threshold = 0.5):\n")
```

Confusion Matrix (Counts, threshold = 0.5):

```
print(conf_matrix_sim)
```

```
      Predicted
Actual 0 1
0 26 5
1 9 27
```



```
cat("\nRow Proportions (Given Actual, % Predicted -- Relates to TPR/FPR):\n")
```

Row Proportions (Given Actual, % Predicted -- Relates to TPR/FPR):

```
## margin = 1 calculates proportions across rows
print(round(prop.table(conf_matrix_sim, margin = 1), 3))
```

	Predicted	
Actual	0	1
0	0.839	0.161
1	0.250	0.750

```
cat("\nColumn Proportions (Given Predicted, % Actual -- Relates to Precision):\n")
```

Column Proportions (Given Predicted, % Actual -- Relates to Precision):

```
## margin = 2 calculates proportions across columns
print(round(prop.table(conf_matrix_sim, margin = 2), 3))
```

	Predicted	
Actual	0	1
0	0.743	0.156
1	0.257	0.844

```
## --- MODIFIED LINES END ---
```

```
## Check if matrix has 2x2 dimensions, otherwise metrics will fail
if (all(dim(conf_matrix_sim) == c(2, 2))) {
  TN <- conf_matrix_sim[1, 1]
  FP <- conf_matrix_sim[1, 2]
  FN <- conf_matrix_sim[2, 1]
  TP <- conf_matrix_sim[2, 2]

  ## Calculate metrics
  error_rate <- (FP + FN) / (TP + TN + FP + FN)
  TPR_Recall <- TP / (TP + FN) # True Positive Rate (Recall / Sensitivity)
  FPR <- FP / (FP + TN)      # False Positive Rate (1 - Specificity)
  Precision <- TP / (TP + FP) # Positive Predictive Value
```

5 Logistic Regression

```
cat(paste("\nMisclassification Error Rate:", round(error_rate, 4), "\n"))
cat(paste("True Positive Rate (Recall):", round(TPR_Recall, 4), "\n"))
cat(paste("False Positive Rate:", round(FPR, 4), "\n"))
cat(paste("Precision:", round(Precision, 4), "\n"))
} else {
  cat("\nCannot calculate full metrics: model predicted only one class.\n")
}
```

Misclassification Error Rate: 0.209
True Positive Rate (Recall): 0.75
False Positive Rate: 0.1613
Precision: 0.8438

Confusion Matrix with threshold=0.1

```
threshold <- 0.1
pred_class_sim <- ifelse(pred_probs_sim > threshold, 1, 0)
conf_matrix_sim <- table(Actual = test_data_sim$y, Predicted = pred_class_sim)

## --- MODIFIED LINES START ---
cat("Confusion Matrix (Counts, threshold = 0.1):\n")
```

Confusion Matrix (Counts, threshold = 0.1):

```
print(conf_matrix_sim)
```

	Predicted	
Actual	0	1
0	15	16
1	1	35

```
cat("\nRow Proportions (Given Actual, % Predicted -- Relates to TPR/FPR):\n")
```

Row Proportions (Given Actual, % Predicted -- Relates to TPR/FPR):

```
## margin = 1 calculates proportions across rows
print(round(prop.table(conf_matrix_sim, margin = 1), 3))
```

	Predicted	
Actual	0	1
0	0.484	0.516
1	0.028	0.972

```
cat("\nColumn Proportions (Given Predicted, % Actual -- Relates to Precision):\n")
```

```
Column Proportions (Given Predicted, % Actual -- Relates to Precision):
```

```
## margin = 2 calculates proportions across columns
print(round(prop.table(conf_matrix_sim, margin = 2), 3))
```

	Predicted	
Actual	0	1
0	0.938	0.314
1	0.062	0.686

```
## --- MODIFIED LINES END ---
```

```
## Check if matrix has 2x2 dimensions
if (all(dim(conf_matrix_sim) == c(2, 2))) {
  TN <- conf_matrix_sim[1, 1]
  FP <- conf_matrix_sim[1, 2]
  FN <- conf_matrix_sim[2, 1]
  TP <- conf_matrix_sim[2, 2]

  ## Calculate metrics
  error_rate <- (FP + FN) / (TP + TN + FP + FN)
  TPR_Recall <- TP / (TP + FN) # True Positive Rate (Recall / Sensitivity)
  FPR <- FP / (FP + TN)      # False Positive Rate (1 - Specificity)
  Precision <- TP / (TP + FP) # Positive Predictive Value

  cat(paste("\nMisclassification Error Rate:", round(error_rate, 4), "\n"))
  cat(paste("True Positive Rate (Recall):", round(TPR_Recall, 4), "\n"))
  cat(paste("False Positive Rate:", round(FPR, 4), "\n"))
  cat(paste("Precision:", round(Precision, 4), "\n"))
} else {
  cat("\nCannot calculate full metrics: model predicted only one class.\n")
}
```


5 Logistic Regression

Misclassification Error Rate: 0.2537
True Positive Rate (Recall): 0.9722
False Positive Rate: 0.5161
Precision: 0.6863

ROC curve and Area Under the ROC (AUC)

```
## 4b. Area Under the Curve (AUC)
roc_curve_sim <- roc(test_data_sim$y, pred_probs_sim, quiet = TRUE)

## Plot the ROC curve
plot(roc_curve_sim, main = "ROC Curve (Simulated Test Data)", print.auc = TRUE)
```



```
auc_value_sim <- auc(roc_curve_sim)
cat(paste("Area Under the Curve (AUC):", round(auc_value_sim, 4), "\n\n"))
```

Area Under the Curve (AUC): 0.8996

PR curve and Area Under PR Curve (AUPR)

```
## Load the ROCR library
## install.packages("ROCR") # Uncomment to install if needed
library(ROCR)

## --- 1. Create a 'prediction' object ---
## 'prediction' takes all predictions and all true labels
pred_obj <- prediction(pred_probs_sim, test_data_sim$y)

## --- 2. Create a 'performance' object for PR ---
## "prec" is for precision, "rec" is for recall
perf_pr <- performance(pred_obj, measure = "prec", x.measure = "rec")


## --- 3. Calculate Area Under the PR Curve (AUPR) ---
perf_auc <- performance(pred_obj, measure = "aucpr") # "aucpr" = Area Under PR Curve
aupr_value <- perf_auc@y.values[[1]]
cat(paste("Area Under PR Curve (AUPR):", round(aupr_value, 4), "\n"))
```

Area Under PR Curve (AUPR): 0.9157

```
## --- 4. Plot the performance object ---
plot(perf_pr,
      main = "Precision-Recall Curve (Simulated Test Data)",
      xlim = c(0, 1),
      ylim = c(0, 1),
      col = "black")

## --- 5. Calculate and add the 'no-skill' baseline ---
baseline_precision_sim <- sum(test_data_sim$y == 1) / length(test_data_sim$y)
abline(h = baseline_precision_sim, col = "blue", lty = 2)

## --- 6. Add a legend with AUPR ---
legend("bottomleft",
      legend = c(
        paste("Model (AUPR =", round(aupr_value, 4), ")"), # <-- MODIFIED LINE
        paste("Baseline (", round(baseline_precision_sim, 3), ")")
      ),
      col = c("black", "blue"),
      lty = c(1, 2),
      bty = "n") # bty="n" removes the box
```



5.7.3 Application to the CHD Dataset

```
## Load the pROC library for AUC calculation
## install.packages("pROC") # Uncomment to install if needed
library(pROC)

## --- 1. Split the data ---
set.seed(123) # for reproducibility
n <- nrow(CHD.data)
train_size <- floor(2/3 * n)
train_indices <- sample(1:n, size = train_size)
train_data <- CHD.data[train_indices, ]
test_data <- CHD.data[-train_indices, ]

## --- 2. Refit the model on the training data ---
fit_train <- glm(
  chd ~ smk + cat + sbp + age + chl + ecg + hpt,
  data = train_data,
  family = binomial(link = "logit")
)
```

```
)

## --- 3. Make predictions on the test data ---
pred_probs <- predict(fit_train, newdata = test_data, type = "response")
```

Plot Predictive Probabilities


```
## --- 5. Plot sorted predicted probabilities ---

## Create a data frame for plotting
plot_data <- data.frame(
  Prob = pred_probs,
  Actual = as.factor(test_data$chd)
)

## Sort by predicted probability
plot_data <- plot_data[order(plot_data$Prob), ]
plot_data$Rank <- 1:nrow(plot_data)

## Create the plot
## We use 'pch' (plot character) to set different symbols
## 'pch = 1' is 'o' (default)
## 'pch = 4' is 'x'
plot(
  plot_data$Rank,
  plot_data$Prob,
  pch = ifelse(plot_data$Actual == 0, 1, 4),
  col = ifelse(plot_data$Actual == 0, "blue", "red"),
  xlab = "Index (Sorted by Predicted Probability)",
  ylab = "Log-odds of Predicted Probability",
  main = "Predicted Probabilities vs. Actual Class",
  ylim = c(0,1)
)
abline(h=0.5)
abline(h=0.1, col="grey")


## Add a legend
legend(
  "topleft",
  legend = c("Actual 0 (o)", "Actual 1 (x)"),
  pch = c(1, 4),
  col = c("blue", "red")
)
```



ROC curve and Area Under the ROC (AUC)

```
## 4b. Area Under the Curve (AUC)
roc_curve <- roc(test_data$chd, pred_probs, quiet = TRUE)

## Plot the ROC curve
plot(roc_curve, main = "ROC Curve (Test Data)", print.auc = TRUE)
```

```
auc_value <- auc(roc_curve)
cat(paste("Area Under the Curve (AUC):", round(auc_value, 4), "\n\n"))
```

Area Under the Curve (AUC): 0.6872

PR curve and Area Under PR Curve (AUPR)

```
## Load the ROCR library
## install.packages("ROCR") # Uncomment to install if needed
library(ROCR)

## --- 1. Create a 'prediction' object ---
## 'prediction' takes all predictions and all true labels
## We use 'pred_probs' and 'test_data$chd' from the CHD data split
pred_obj <- prediction(pred_probs, test_data$chd)

## --- 2. Create a 'performance' object for PR ---
## "prec" is for precision, "rec" is for recall
```

5 Logistic Regression

```
perf_pr <- performance(pred_obj, measure = "prec", x.measure = "rec")

## --- 3. Calculate Area Under the PR Curve (AUPR) ---
perf_auc <- performance(pred_obj, measure = "aucpr") # "aucpr" = Area Under PR Curve
aupr_value <- perf_auc@y.values[[1]]
cat(paste("Area Under PR Curve (AUPR):", round(aupr_value, 4), "\n"))
```


Area Under PR Curve (AUPR): 0.2826

```
## --- 4. Plot the performance object ---
plot(perf_pr,
      main = "Precision-Recall Curve (Test Data)",
      xlim = c(0, 1),
      ylim = c(0, 1),
      col = "black")

## --- 5. Calculate and add the 'no-skill' baseline ---
baseline_precision <- sum(test_data$chd == 1) / length(test_data$chd)
abline(h = baseline_precision, col = "blue", lty = 2)

## --- 6. Add a legend with AUPR ---
legend("bottomleft",
      legend = c(
        paste("Model (AUPR =", round(aupr_value, 4), ")"), # <-- MODIFIED LINE
        paste("Baseline (", round(baseline_precision, 3), ")")
      ),
      col = c("black", "blue"),
      lty = c(1, 2),
      bty = "n") # bty="n" removes the box
```

Precision–Recall Curve (Test Data)



6 One-factor Design

6.1 Completely Randomized Design

6.1.1 Plasma Etching Experiment

This section analyzes data from a **Completely Randomized Design (CRD)**. In a CRD, experimental units (in this case, the silicon wafers being etched) are assigned to treatments (the RF Power levels) completely at random. The primary goal is to determine if changing the RF Power level has a statistically significant effect on the mean etch rate.

6.1.1.1 Data and Visualization

We begin by loading the data into a single, tidy `data.frame`. The response variable, `rate`, contains all the etch rate observations. The predictor variable, `power`, is a **factor**, which is R's way of representing a categorical variable. This tells R to treat the different power levels as distinct groups.

```
## Define the data vectors
rate <- c(575, 542, 530, 539, 570, 565, 593, 590, 579, 610,
          600, 651, 610, 637, 629, 725, 700, 715, 685, 710)
power_levels <- c(160, 180, 200, 220)

## Create the data frame
etching_df <- data.frame(
  rate = rate,
  power = factor(rep(power_levels, each = 5))
)

## Display the first few rows
etching_df
```


	rate	power
1	575	160
2	542	160
3	530	160
4	539	160
5	570	160

6 One-factor Design

6	565	180
7	593	180
8	590	180
9	579	180
10	610	180
11	600	200
12	651	200
13	610	200
14	637	200
15	629	200
16	725	220
17	700	220
18	715	220
19	685	220
20	710	220

Grouped Boxplots

```
boxplot(rate~power, data=etching_df)
```



Using ggplot to visualize grouped data

```

library(ggplot2)
library(dplyr) # Using dplyr for easier data manipulation

## Calculate group means and their start/end indices
mean_rates <- etching_df %>%
  mutate(obs_index = row_number()) %>%
  group_by(power) %>%
  summarise(
    mean_rate = mean(rate),
    x_start = min(obs_index) - 0.5,
    x_end = max(obs_index) + 0.5
  )

ggplot(etching_df, aes(x = 1:nrow(etching_df), y = rate, color = power)) +
  geom_point(size = 3, alpha = 0.7) + # Plot individual data points
  geom_segment(
    data = mean_rates,
    aes(x = x_start, xend = x_end, y = mean_rate, yend = mean_rate),
    linetype = "dashed",
    size = 1.2
  ) + # Add line segments for group means
  labs(
    title = "Etch Rate Observations by RF Power Level",
    x = "Observation Index",
    y = "Etch Rate",
    color = "RF Power (W)"
  ) +
  scale_color_brewer(palette = "Set1") +
  theme_minimal() +
  theme(plot.title = element_text(hjust = 0.5))

```




Figure 6.1: Index Plot of Etch Rate with Group-Specific Mean Lines

6.1.1.2 Model Fitting with Sum-to-Zero Constraint

We fit a linear model using the `lm()` function to perform an **Analysis of Variance (ANOVA)**. The model is specified as `rate ~ power`, and we now include the `data = etching_df` argument.

To get interpretable estimates for the treatment effects (τ_i), we use a **sum-to-zero constraint** (`contr.sum`), which forces the sum of the treatment effects to be zero ($\sum \tau_i = 0$).

```
fit <- lm(rate ~ power, data = etching_df, contrasts = list(power = contr.sum))
cat ("Model Matrix:\n")
```

Model Matrix:

```
model.matrix(fit)
```

```
(Intercept) power1 power2 power3
1           1       1       0       0
```



```

2          1      1      0      0
3          1      1      0      0
4          1      1      0      0
5          1      1      0      0
6          1      0      1      0
7          1      0      1      0
8          1      0      1      0
9          1      0      1      0
10         1      0      1      0
11         1      0      0      1
12         1      0      0      1
13         1      0      0      1
14         1      0      0      1
15         1      0      0      1
16         1     -1     -1     -1
17         1     -1     -1     -1
18         1     -1     -1     -1
19         1     -1     -1     -1
20         1     -1     -1     -1

```

```

attr("assign")
[1] 0 1 1 1
attr("contrasts")
attr("contrasts")$power
      [,1] [,2] [,3]
160     1     0     0
180     0     1     0
200     0     0     1
220    -1    -1    -1

```

```

summary.fit <- summary(fit)
cat ("Summary of lm fitting results:\n")

```

Summary of lm fitting results:

```
summary.fit
```

Call:

```
lm(formula = rate ~ power, data = etching_df, contrasts = list(power = contr.sum))
```

Residuals:

Min	1Q	Median	3Q	Max
-25.4	-13.0	2.8	13.2	25.6

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	617.750	4.085	151.234	< 2e-16 ***
power1	-66.550	7.075	-9.406	6.39e-08 ***
power2	-30.350	7.075	-4.290	0.000563 ***
power3	7.650	7.075	1.081	0.295602

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 18.27 on 16 degrees of freedom

Multiple R-squared: 0.9261, Adjusted R-squared: 0.9122

F-statistic: 66.8 on 3 and 16 DF, p-value: 2.883e-09

6.1.1.3 Point Estimation of Parameters

The output of the model provides estimates for the overall mean ($\hat{\mu}$) and the treatment effects for the first k-1 levels ($\hat{\tau}_1, \hat{\tau}_2, \hat{\tau}_3$).

- $\hat{\mu}$ (the Intercept) is the estimate of the grand mean etch rate across all power levels.
- $\hat{\tau}_i$ is the estimated effect of the i-th power level, representing how much that level's mean deviates from the grand mean.

Using the sum-to-zero constraint, we can manually calculate the effect for the final level, $\hat{\tau}_4$.

```
## Extract coefficients
est <- coef(fit)
tau4.hat <- -sum(est[-1])
taui.hat <- c(est[-1], tau4.hat)
print("Estimated Treatment Effects (tau_i):")
```

```
[1] "Estimated Treatment Effects (tau_i):"
```

```
print(taui.hat)
```

```
power1 power2 power3
-66.55 -30.35  7.65  89.25
```

```
## Estimates of treatment means (mu_i)
mu.hat <- est[1]
mui.hat <- mu.hat + taui.hat
print("Estimated Treatment Means (mu_i):")
```

```
[1] "Estimated Treatment Means (mu_i):"
```

```
print(mui.hat)
```

```
power1 power2 power3
551.2  587.4  625.4  707.0
```

6.1.1.4 ANOVA Table

The ANOVA table partitions the total variation into variation **between** treatment groups (power) and variation **within** treatment groups (random error). The **p-value** ($\Pr(>F)$) indicates if the treatment has a significant effect.

```
anova(fit)
```

Analysis of Variance Table

Response: rate

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
power	3	66871	22290.2	66.797	2.883e-09 ***
Residuals	16	5339	333.7		

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

6.1.1.5 95% Confidence Intervals for Treatment Means

A **confidence interval** provides a range of plausible values for the true mean etch rate at each power level.

```
## Number of replicates
n <- 5
## Extract sqrt(MSE) and error df
sqrt.MSE <- summary.fittedvalues$sigma
DF <- fit$df.residual
## Find t-value
t.value <- qt(0.975, DF)
## Calculate CIs
CI.lower <- mui.hat - t.value * sqrt.MSE / sqrt(n)
CI.upper <- mui.hat + t.value * sqrt.MSE / sqrt(n)

## Display CIs
data.frame(Power_Level = power_levels, Mean = mui.hat, Lower_CI = CI.lower, Upper_CI = CI.upper)
```

6 One-factor Design

	Power_Level	Mean	Lower_CI	Upper_CI
power1	160	551.2	533.8815	568.5185
power2	180	587.4	570.0815	604.7185
power3	200	625.4	608.0815	642.7185
	220	707.0	689.6815	724.3185

Alternatively, one can use a model without intercept

```
fit_nointercpt <- lm(rate ~ 0+power, data = etching_df)
summary(fit_nointercpt)
```

Call:

```
lm(formula = rate ~ 0 + power, data = etching_df)
```

Residuals:

Min	1Q	Median	3Q	Max
-25.4	-13.0	2.8	13.2	25.6

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
power160	551.200	8.169	67.47	<2e-16 ***
power180	587.400	8.169	71.90	<2e-16 ***
power200	625.400	8.169	76.55	<2e-16 ***
power220	707.000	8.169	86.54	<2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 18.27 on 16 degrees of freedom

Multiple R-squared: 0.9993, Adjusted R-squared: 0.9991

F-statistic: 5768 on 4 and 16 DF, p-value: < 2.2e-16

```
confint(fit_nointercpt)
```

	2.5 %	97.5 %
power160	533.8815	568.5185
power180	570.0815	604.7185
power200	608.0815	642.7185
power220	689.6815	724.3185

6.1.1.6 Comparison with Default “Treatment” Contrast

Fitting the model without specifying contrasts uses R’s default (“treatment” contrast), which sets $\tau_1 = 0$. The fundamental results (ANOVA, treatment means) remain unchanged.

```
fit1 <- lm(rate ~ power, data = etching_df)
summary(fit1)
```

Call:

```
lm(formula = rate ~ power, data = etching_df)
```

Residuals:

Min	1Q	Median	3Q	Max
-25.4	-13.0	2.8	13.2	25.6

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	551.200	8.169	67.471	< 2e-16 ***
power180	36.200	11.553	3.133	0.00642 **
power200	74.200	11.553	6.422	8.44e-06 ***
power220	155.800	11.553	13.485	3.73e-10 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 18.27 on 16 degrees of freedom

Multiple R-squared: 0.9261, Adjusted R-squared: 0.9122

F-statistic: 66.8 on 3 and 16 DF, p-value: 2.883e-09

6.1.1.7 Pairwise Comparisons

Since our ANOVA result was significant, we perform **post-hoc tests** to determine exactly which pairs of power levels have different means.

6.1.1.7.1 Tukey’s HSD Test

Tukey’s Honest Significant Difference (HSD) controls the **family-wise error rate**, adjusting p-values to account for multiple comparisons.

```
fit.aov <- aov(rate ~ power, data = etching_df)
TukeyHSD(fit.aov)
```

6 One-factor Design

Tukey multiple comparisons of means
95% family-wise confidence level

```
Fit: aov(formula = rate ~ power, data = etching_df)
```

```
$power
      diff      lwr      upr    p adj
180-160  36.2    3.145624 69.25438 0.0294279
200-160  74.2   41.145624 107.25438 0.0000455
220-160 155.8  122.745624 188.85438 0.0000000
200-180  38.0    4.945624  71.05438 0.0215995
220-180 119.6   86.545624 152.65438 0.0000001
220-200  81.6   48.545624 114.65438 0.0000146
```

6.1.1.7.2 Fisher's LSD Test

The **Fisher's Least Significant Difference (LSD)** test does not control the family-wise error rate but is more powerful.

```
with(etching_df, pairwise.t.test(rate, power, p.adj = "none"))
```

Pairwise comparisons using t tests with pooled SD

data: rate and power

	160	180	200
180	0.0064	-	-
200	8.4e-06	0.0046	-
220	3.7e-10	1.7e-08	2.7e-06

P value adjustment method: none

6.1.1.8 Checking Model Assumptions

The validity of our ANOVA results depends on three key assumptions about the model's residuals. We use diagnostic plots to check them.

```
r <- rstudent(fit)
fitted <- fitted.values(fit)
```

6.1.1.8.1 Normality of Residuals

A **Normal Q-Q plot** is used to check if the residuals are normally distributed. The points should fall closely along the straight diagonal line.

```
qqnorm(r)
qqline(r)
```




Figure 6.2: Normal Q-Q plot of standardized residuals.

6.1.1.8.2 Independence of Residuals

A plot of **residuals versus run order** helps check for independence. We look for random scatter around the zero line.

```
plot(r, ylab = "Standardized residuals", xlab = "Run order",
     main = "Plot of residuals vs. run order")
abline(h = 0)
```




Figure 6.3: Standardized residuals vs. run order.

6.1.1.8.3 Constant Variance (Homoscedasticity)

A plot of **residuals versus fitted values** helps check for constant variance. The spread of residuals should be roughly constant across all fitted values.

```
plot(fitted, r, ylab = "Standardized residuals",  
     xlab = "Fitted values", main = "Plot of residuals vs. fitted values")  
abline(h = 0)
```





Figure 6.4: Standardized residuals vs. fitted values.

6.2 Unbalanced Designs with nequal Sample Sizes

The ANOVA framework also handles **unbalanced designs**. We again start by creating a data frame.

```
## Create the data frame
bricks_df <- data.frame(
  density = c(21.8, 21.9, 21.7, 21.6, 21.7,
              21.7, 21.4, 21.5, 21.4,
              21.9, 21.8, 21.8, 21.6, 21.5,
              21.9, 21.7, 21.8, 21.4),
  temperature = factor(c(rep(100, 5), rep(125, 4), rep(150, 5), rep(175, 4)))
)

bricks_df
```

```
  density temperature
1    21.8         100
2    21.9         100
```

6 One-factor Design

3	21.7	100
4	21.6	100
5	21.7	100
6	21.7	125
7	21.4	125
8	21.5	125
9	21.4	125
10	21.9	150
11	21.8	150
12	21.8	150
13	21.6	150
14	21.5	150
15	21.9	175
16	21.7	175
17	21.8	175
18	21.4	175

```
## Fit the model and get the ANOVA table
fit2 <- lm(density ~ temperature, data = bricks_df)
summary(fit2)
```

Call:

```
lm(formula = density ~ temperature, data = bricks_df)
```

Residuals:

Min	1Q	Median	3Q	Max
-0.300	-0.100	0.000	0.095	0.200

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	21.74000	0.07171	303.150	<2e-16 ***
temperature125	-0.24000	0.10757	-2.231	0.0425 *
temperature150	-0.02000	0.10142	-0.197	0.8465
temperature175	-0.04000	0.10757	-0.372	0.7156

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.1604 on 14 degrees of freedom

Multiple R-squared: 0.3025, Adjusted R-squared: 0.153

F-statistic: 2.024 on 3 and 14 DF, p-value: 0.1569

```
anova(fit2)
```

Analysis of Variance Table

Response: density

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
temperature	3	0.15611	0.052037	2.0237	0.1569
Residuals	14	0.36000	0.025714		

In this case, the large p-value (0.133) indicates that there is no statistically significant evidence that firing temperature affects brick density.

6.3 Randomized Complete Block Design

6.3.1 Vascular Graft Experiment

This section analyzes a **Randomized Complete Block Design (RCBD)**, used to control for a known source of variability (here, “batches of resin,” treated as **blocks**).

6.3.1.1 Data and Visualization

We structure the data in a `data.frame` to identify the response, treatment (pressure), and block (batch) for each observation.

```
## Define data vectors
strength <- c(90.3, 89.2, 98.2, 93.9, 87.4, 97.9,
              92.5, 89.5, 90.6, 94.7, 87.0, 95.8,
              85.5, 90.8, 89.6, 86.2, 88.0, 93.4,
              82.5, 89.5, 85.6, 87.4, 78.9, 90.7)
pressure_levels <- rep(c(8500, 8700, 8900, 9100), each = 6)
batch_levels <- rep(1:6, 4)

## Create the data frame
graft_df <- data.frame(
  strength = strength,
  pressure = factor(pressure_levels),
  batch = factor(batch_levels)
)


graft_df
```

6 One-factor Design

	strength	pressure	batch
1	90.3	8500	1
2	89.2	8500	2
3	98.2	8500	3
4	93.9	8500	4
5	87.4	8500	5
6	97.9	8500	6
7	92.5	8700	1
8	89.5	8700	2
9	90.6	8700	3
10	94.7	8700	4
11	87.0	8700	5
12	95.8	8700	6
13	85.5	8900	1
14	90.8	8900	2
15	89.6	8900	3
16	86.2	8900	4
17	88.0	8900	5
18	93.4	8900	6
19	82.5	9100	1
20	89.5	9100	2
21	85.6	9100	3
22	87.4	9100	4
23	78.9	9100	5
24	90.7	9100	6

Visualize the Block and Treatment Effects

```
par (mfrow = c(1,2))
#boxplot
plot(strength ~ batch, data=graft_df , main = "Block")
plot(strength ~ pressure, data=graft_df , main = "Pressure")
```



Interaction Plots

```
ggplot(graft_df, aes(x = pressure, y = strength, group = batch, color = batch)) +
  stat_summary(fun = mean, geom = "line", size = 1) +
  stat_summary(fun = mean, geom = "point", size = 3) +
  labs(
    title = "Interaction Plot: Batch and Pressure",
    x = "Pressue",
    y = "Strength",
    color = "Batch"
  ) +
  theme_minimal() +
  theme(plot.title = element_text(hjust = 0.5))
```




Figure 6.5: Interaction between Material Type and Temperature.

6.3.1.2 Model Fitting and ANOVA

The model `strength ~ pressure + batch` partitions the total variance into treatment, block, and error components. Our primary interest is in the significance of the pressure factor.

```
rcbd.fit1 <- aov(strength ~ pressure + batch, data = graft_df)
anova(rcbd.fit1)
```

Analysis of Variance Table

Response: strength

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
pressure	3	178.17	59.390	8.1071	0.001916 **
batch	5	192.25	38.450	5.2487	0.005532 **
Residuals	15	109.89	7.326		

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

The small p-value for pressure (0.0019) provides strong evidence that extrusion pressure significantly affects graft strength after accounting for batch differences.

6.3.1.3 Model Adequacy Checks

The assumptions for an RCBD are the same as for a CRD. We perform the same diagnostic checks.

```
rcbd.r1 <- rstudent(rcbd.fit1)
rcbd.fitted1 <- fitted.values(rcbd.fit1)

qqnorm(rcbd.r1, main = "Normal Q-Q Plot")
qqline(rcbd.r1)
plot(rcbd.fitted1, rcbd.r1, ylab = "Standardized residuals",
     xlab = "Fitted values", main = "Residuals vs. Fitted")
abline(h = 0)
plot(graft_df$pressure, rcbd.r1, ylab = "Standardized residuals",
     xlab = "Extrusion pressure", main = "Residuals vs. Treatment")
abline(h = 0)
plot(graft_df$batch, rcbd.r1, ylab = "Standardized residuals",
     xlab = "Batches of raw material", main = "Residuals vs. Block")
abline(h = 0)
```

6.3.1.4 Pairwise Comparisons

Again, since the treatment factor (pressure) is significant, we perform post-hoc tests.

6.3.1.4.1 Tukey's HSD Test

Tukey's HSD compares all pairs of treatment levels while controlling the family-wise error rate.

```
TukeyHSD(rcbd.fit1, which = "pressure")
```


Tukey multiple comparisons of means
95% family-wise confidence level

```
Fit: aov(formula = strength ~ pressure + batch, data = graft_df)
```


```
$pressure
```

	diff	lwr	upr	p adj
8700-8500	-1.133333	-5.637161	3.370495	0.8854831
8900-8500	-3.900000	-8.403828	0.603828	0.1013084
9100-8500	-7.050000	-11.553828	-2.546172	0.0020883


6 One-factor Design




(a) Normal Q-Q Plot



(a) Residuals vs. Fitted Values



(a) Residuals vs. Treatment (Pressure)



(a) Residuals vs. Block (Batch)


```
8900-8700 -2.766667 -7.270495 1.737161 0.3245644
9100-8700 -5.916667 -10.420495 -1.412839 0.0086667
9100-8900 -3.150000 -7.653828 1.353828 0.2257674
```

6.3.1.4.2 Fisher's LSD Test

The `LSD.test()` function from the `agricolae` package correctly handles the error structure of an RCBD.

```
## install.packages("agricolae")
library(agricolae)

out <- LSD.test(rcbd.fit1, trt = "pressure", p.adj = "none", group = FALSE)
print(out$comparison)
```

	difference	pvalue	signif.	LCL	UCL
8500 - 8700	1.133333	0.4795		-2.1974047	4.464071
8500 - 8900	3.900000	0.0247	*	0.5692620	7.230738
8500 - 9100	7.050000	0.0004	***	3.7192620	10.380738
8700 - 8900	2.766667	0.0970	.	-0.5640714	6.097405
8700 - 9100	5.916667	0.0018	**	2.5859286	9.247405
8900 - 9100	3.150000	0.0621	.	-0.1807380	6.480738

7 Two-Factor Factorial Design

7.1 Battery Design Experiment

This analysis explores data from a **two-factor factorial experiment** designed to assess the lifespan of a battery. The experiment investigates two factors: **material type** (with 3 levels) and **operating temperature** (with 3 levels: 15°C, 70°C, and 125°C). The primary goal is to understand not only how each factor individually affects battery life but, more importantly, whether the effect of temperature depends on the material type used. This combined effect is known as an **interaction**.

7.2 Data Setup and Preparation

First, we organize the raw data into a structured `data.frame`. This is a best practice in R that makes the data easier to manage and the code more readable. We create columns for the response variable `life` and the two factors, `material` and `temperature`, ensuring they are treated as categorical variables (factors) for the analysis.

```
## Response variable: battery life
life <- c(130,155,74,180, 34,40,80,75, 20,70,82,58,
         150,188,159,126, 136,122,106,115, 25,70,58,45,
         138,110,168,160, 174,120,150,139, 96,104,82,60)

## Create the data frame
battery_df <- data.frame(
  life = life,
  material = factor(rep(1:3, each = 12)),
  temperature = factor(rep(rep(c(15, 70, 125), each = 4), 3))
)

## Preview the data
battery_df
```

	life	material	temperature
1	130	1	15
2	155	1	15
3	74	1	15

7 Two-Factor Factorial Design

4	180	1	15
5	34	1	70
6	40	1	70
7	80	1	70
8	75	1	70
9	20	1	125
10	70	1	125
11	82	1	125
12	58	1	125
13	150	2	15
14	188	2	15
15	159	2	15
16	126	2	15
17	136	2	70
18	122	2	70
19	106	2	70
20	115	2	70
21	25	2	125
22	70	2	125
23	58	2	125
24	45	2	125
25	138	3	15
26	110	3	15
27	168	3	15
28	160	3	15
29	174	3	70
30	120	3	70
31	150	3	70
32	139	3	70
33	96	3	125
34	104	3	125
35	82	3	125
36	60	3	125

7.3 Exploratory Data Analysis and Visualization

Before fitting a formal model, we visualize the data to get an intuition for the relationships between the factors and the response.


7.4 Boxplots of Main Effects

Boxplots are excellent for examining the distribution of battery life for each level of our factors independently. This gives us a preliminary look at the **main effects**—the individual impact of material type and temperature.

```
library(ggplot2)

## Boxplot for Material Type
ggplot(battery_df, aes(x = material, y = life, fill = material)) +
  geom_boxplot() +
  labs(title = "Battery Life by Material Type", x = "Material Type", y = "Life (hours)") +
  theme_minimal() +
  theme(legend.position = "none")

## Boxplot for Temperature
ggplot(battery_df, aes(x = temperature, y = life, fill = temperature)) +
  geom_boxplot() +
  labs(title = "Battery Life by Temperature", x = "Temperature (°C)", y = "Life (hours)") +
  theme_minimal() +
  theme(legend.position = "none")
```



(a) Distribution of Battery Life by Material and Temperature. (a) Distribution of Battery Life by Material and Temperature.

7.5 Interaction Plot

The most crucial plot for a factorial experiment is the **interaction plot**. It displays the mean battery life for each combination of material and temperature. If the lines are parallel, it suggests there is no interaction. If the lines are not parallel (i.e., they cross or diverge), it indicates that the effect of temperature on battery life is different for each material type, signaling a likely interaction.

7 Two-Factor Factorial Design

```
ggplot(battery_df, aes(x = temperature, y = life, group = material, color = material)) +  
  stat_summary(fun = mean, geom = "line", size = 1) +  
  stat_summary(fun = mean, geom = "point", size = 3) +  
  labs(  
    title = "Interaction Plot: Material Type and Temperature",  
    x = "Temperature (°C)",  
    y = "Average Battery Life (hours)",  
    color = "Material Type"  
  ) +  
  theme_minimal() +  
  theme(plot.title = element_text(hjust = 0.5))
```




Figure 7.3: Interaction between Material Type and Temperature.

The non-parallel lines in the plot strongly suggest that a significant interaction effect is present. Specifically, the performance of Material 3 drops less dramatically with increasing temperature compared to Materials 1 and 2.

7.6 Model Fitting and Analysis of Variance (ANOVA)

We now fit a linear model to formally test the significance of the main effects and the interaction term. The model `life ~ material * temperature` is shorthand for `life ~ material + temperature + material:temperature`. We use a sum-to-zero contrast (`contr.sum`) for balanced interpretation of the effects. The **ANOVA table** will tell us if the variation caused by our factors is statistically significant compared to the random variation in the data.

```
## Fit the full factorial model
battery_fit <- lm(life ~ material * temperature,
                 data = battery_df,
                 contrasts = list(material = contr.sum, temperature = contr.sum))

summary(battery_fit)
```

Call:

```
lm(formula = life ~ material * temperature, data = battery_df,
   contrasts = list(material = contr.sum, temperature = contr.sum))
```

Residuals:

Min	1Q	Median	3Q	Max
-60.750	-14.625	1.375	17.938	45.250

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	105.528	4.331	24.367	< 2e-16 ***
material1	-22.361	6.125	-3.651	0.00111 **
material2	2.806	6.125	0.458	0.65057
temperature1	39.306	6.125	6.418	7.1e-07 ***
temperature2	2.056	6.125	0.336	0.73975
material1:temperature1	12.278	8.662	1.417	0.16778
material2:temperature1	8.111	8.662	0.936	0.35735
material1:temperature2	-27.972	8.662	-3.229	0.00325 **
material2:temperature2	9.361	8.662	1.081	0.28936

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 25.98 on 27 degrees of freedom

Multiple R-squared: 0.7652, Adjusted R-squared: 0.6956

F-statistic: 11 on 8 and 27 DF, p-value: 9.426e-07

```
## Generate the ANOVA table
anova(battery_fit)
```

Analysis of Variance Table

Response: life

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
material	2	10684	5341.9	7.9114	0.001976 **
temperature	2	39119	19559.4	28.9677	1.909e-07 ***
material:temperature	4	9614	2403.4	3.5595	0.018611 *
Residuals	27	18231	675.2		

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

The ANOVA table shows very small p-values ($\text{Pr}(>F)$) for material, temperature, and, most importantly, the material:temperature interaction. This confirms our visual inspection: all effects are statistically significant. **Because the interaction is significant, our interpretation should focus on the interaction itself rather than the main effects in isolation.**

7.7 Model Adequacy Checks

The validity of our ANOVA results depends on the model's residuals meeting certain assumptions (normality, constant variance, independence). We check these with diagnostic plots.

```
## Extract standardized residuals and fitted values
battery_fit_diag <- data.frame(
  residuals = rstandard(battery_fit),
  fitted = fitted.values(battery_fit)
)

## Normal Q-Q Plot
p1 <- ggplot(battery_fit_diag, aes(sample = residuals)) +
  stat_qq() +
  stat_qq_line() +
  labs(title = "Normal Q-Q Plot", x = "Theoretical Quantiles", y = "Standardized Residuals") +
  theme_minimal()


## Residuals vs. Fitted Plot
p2 <- ggplot(battery_fit_diag, aes(x = fitted, y = residuals)) +
  geom_point() +
  geom_hline(yintercept = 0, linetype = "dashed", color = "red") +
  labs(title = "Residuals vs. Fitted Values", x = "Fitted Values", y = "Standardized Residuals")
```




```
theme_minimal()
```

```
p1
```

```
p2
```



(a) Diagnostic plots for the battery life model.



(a) Diagnostic plots for the battery life model.

The Normal Q-Q plot shows the points falling roughly along the line, suggesting the normality assumption is met. The Residuals vs. Fitted plot shows a random scatter of points around the zero line, indicating that the variance is reasonably constant. The model assumptions appear to be satisfied.

7.8 Post-Hoc Analysis: Pairwise Comparisons

Since the interaction is significant, we must compare the means of the nine specific treatment combinations (3 materials \times 3 temperatures). Simply comparing the average effect of Material 1 vs. Material 2 would be misleading, as that difference depends on the temperature.

7.9 Tukey's HSD Test

Tukey's Honest Significant Difference (HSD) test is a post-hoc test that compares all possible pairs of means while controlling the family-wise error rate. We apply it to an aov model object. The output for the material:temperature interaction shows which specific combinations are significantly different from one another.

```
## Fit the model using aov() for Tukey's test
battery_aov <- aov(life ~ material * temperature, data = battery_df)
```

7 Two-Factor Factorial Design

```
## Perform Tukey's HSD test
TukeyHSD(battery_aov)
```

Tukey multiple comparisons of means
95% family-wise confidence level

```
Fit: aov(formula = life ~ material * temperature, data = battery_df)
```

\$material

	diff	lwr	upr	p adj
2-1	25.16667	-1.135677	51.46901	0.0627571
3-1	41.91667	15.614323	68.21901	0.0014162
3-2	16.75000	-9.552344	43.05234	0.2717815

\$temperature

	diff	lwr	upr	p adj
70-15	-37.25000	-63.55234	-10.94766	0.0043788
125-15	-80.66667	-106.96901	-54.36432	0.0000001
125-70	-43.41667	-69.71901	-17.11432	0.0009787

\$`material:temperature`

	diff	lwr	upr	p adj
2:15-1:15	21.00	-40.823184	82.823184	0.9616404
3:15-1:15	9.25	-52.573184	71.073184	0.9998527
1:70-1:15	-77.50	-139.323184	-15.676816	0.0065212
2:70-1:15	-15.00	-76.823184	46.823184	0.9953182
3:70-1:15	11.00	-50.823184	72.823184	0.9994703
1:125-1:15	-77.25	-139.073184	-15.426816	0.0067471
2:125-1:15	-85.25	-147.073184	-23.426816	0.0022351
3:125-1:15	-49.25	-111.073184	12.573184	0.2016535
3:15-2:15	-11.75	-73.573184	50.073184	0.9991463
1:70-2:15	-98.50	-160.323184	-36.676816	0.0003449
2:70-2:15	-36.00	-97.823184	25.823184	0.5819453
3:70-2:15	-10.00	-71.823184	51.823184	0.9997369
1:125-2:15	-98.25	-160.073184	-36.426816	0.0003574
2:125-2:15	-106.25	-168.073184	-44.426816	0.0001152
3:125-2:15	-70.25	-132.073184	-8.426816	0.0172076
1:70-3:15	-86.75	-148.573184	-24.926816	0.0018119
2:70-3:15	-24.25	-86.073184	37.573184	0.9165175
3:70-3:15	1.75	-60.073184	63.573184	1.0000000
1:125-3:15	-86.50	-148.323184	-24.676816	0.0018765
2:125-3:15	-94.50	-156.323184	-32.676816	0.0006078
3:125-3:15	-58.50	-120.323184	3.323184	0.0742711
2:70-1:70	62.50	0.676816	124.323184	0.0460388

3:70-1:70	88.50	26.676816	150.323184	0.0014173
1:125-1:70	0.25	-61.573184	62.073184	1.0000000
2:125-1:70	-7.75	-69.573184	54.073184	0.9999614
3:125-1:70	28.25	-33.573184	90.073184	0.8281938
3:70-2:70	26.00	-35.823184	87.823184	0.8822881
1:125-2:70	-62.25	-124.073184	-0.426816	0.0474675
2:125-2:70	-70.25	-132.073184	-8.426816	0.0172076
3:125-2:70	-34.25	-96.073184	27.573184	0.6420441
1:125-3:70	-88.25	-150.073184	-26.426816	0.0014679
2:125-3:70	-96.25	-158.073184	-34.426816	0.0004744
3:125-3:70	-60.25	-122.073184	1.573184	0.0604247
2:125-1:125	-8.00	-69.823184	53.823184	0.9999508
3:125-1:125	28.00	-33.823184	89.823184	0.8347331
3:125-2:125	36.00	-25.823184	97.823184	0.5819453

7.10 Fisher's LSD Method

The **Fisher's Least Significant Difference (LSD)** method is another option for pairwise comparisons. To test the interaction means, we must specify both factors in the `trt` argument.

```
library(agricolae)

## Perform LSD test on the interaction term
lsd_results <- LSD.test(battery_aov, trt = c("material", "temperature"),
                        p.adj = "none", group = FALSE)

## Print the comparison table
print(lsd_results$comparison)
```

	difference	pvalue	signif.	LCL	UCL
1:125 - 1:15	-77.25	0.0003	***	-114.950479	-39.549521
1:125 - 1:70	0.25	0.9892		-37.450479	37.950479
1:125 - 2:125	8.00	0.6667		-29.700479	45.700479
1:125 - 2:15	-98.25	0.0000	***	-135.950479	-60.549521
1:125 - 2:70	-62.25	0.0022	**	-99.950479	-24.549521
1:125 - 3:125	-28.00	0.1392		-65.700479	9.700479
1:125 - 3:15	-86.50	0.0001	***	-124.200479	-48.799521
1:125 - 3:70	-88.25	0.0001	***	-125.950479	-50.549521
1:15 - 1:70	77.50	0.0002	***	39.799521	115.200479
1:15 - 2:125	85.25	0.0001	***	47.549521	122.950479
1:15 - 2:15	-21.00	0.2631		-58.700479	16.700479
1:15 - 2:70	15.00	0.4214		-22.700479	52.700479
1:15 - 3:125	49.25	0.0124	*	11.549521	86.950479

7 Two-Factor Factorial Design

1:15 - 3:15	-9.25	0.6187		-46.950479	28.450479
1:15 - 3:70	-11.00	0.5544		-48.700479	26.700479
1:70 - 2:125	7.75	0.6765		-29.950479	45.450479
1:70 - 2:15	-98.50	0.0000	***	-136.200479	-60.799521
1:70 - 2:70	-62.50	0.0021	**	-100.200479	-24.799521
1:70 - 3:125	-28.25	0.1358		-65.950479	9.450479
1:70 - 3:15	-86.75	0.0001	***	-124.450479	-49.049521
1:70 - 3:70	-88.50	0.0000	***	-126.200479	-50.799521
2:125 - 2:15	-106.25	0.0000	***	-143.950479	-68.549521
2:125 - 2:70	-70.25	0.0007	***	-107.950479	-32.549521
2:125 - 3:125	-36.00	0.0605	.	-73.700479	1.700479
2:125 - 3:15	-94.50	0.0000	***	-132.200479	-56.799521
2:125 - 3:70	-96.25	0.0000	***	-133.950479	-58.549521
2:15 - 2:70	36.00	0.0605	.	-1.700479	73.700479
2:15 - 3:125	70.25	0.0007	***	32.549521	107.950479
2:15 - 3:15	11.75	0.5279		-25.950479	49.450479
2:15 - 3:70	10.00	0.5907		-27.700479	47.700479
2:70 - 3:125	34.25	0.0732	.	-3.450479	71.950479
2:70 - 3:15	-24.25	0.1980		-61.950479	13.450479
2:70 - 3:70	-26.00	0.1685		-63.700479	11.700479
3:125 - 3:15	-58.50	0.0036	**	-96.200479	-20.799521
3:125 - 3:70	-60.25	0.0029	**	-97.950479	-22.549521
3:15 - 3:70	-1.75	0.9248		-39.450479	35.950479

The results from both Tukey's HSD and Fisher's LSD provide detailed p-values for comparing pairs of treatment combinations, allowing us to make specific conclusions, such as "at 125°C, Material 3 has a significantly longer life than Materials 1 and 2."